

CS4100: 計算機結構

Instruction Set Architecture



國立清華大學資訊工程學系
一零零零學年度第二學期

Outline

- ◆ **Instruction set architecture (Sec 2.1)**
- ◆ **Operands**
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ **Signed and unsigned numbers**
- ◆ **Representing instructions**
- ◆ **Operations**
 - Logical
 - Decision making and branches
- ◆ **Supporting procedures in hardware**
- ◆ **Communicating with people**
- ◆ **Addressing for 32-bit immediate and addresses**
- ◆ **Translating and starting a program**
- ◆ **A sort example**
- ◆ **Arrays versus pointers**
- ◆ **ARM and x86 instruction sets**

What Is Computer Architecture?

Computer Architecture =
Instruction Set Architecture
+ Machine Organization

- ◆ “... the attributes of a [computing] system as seen by the [assembly language] programmer, *i.e.* the conceptual structure and functional behavior ...”

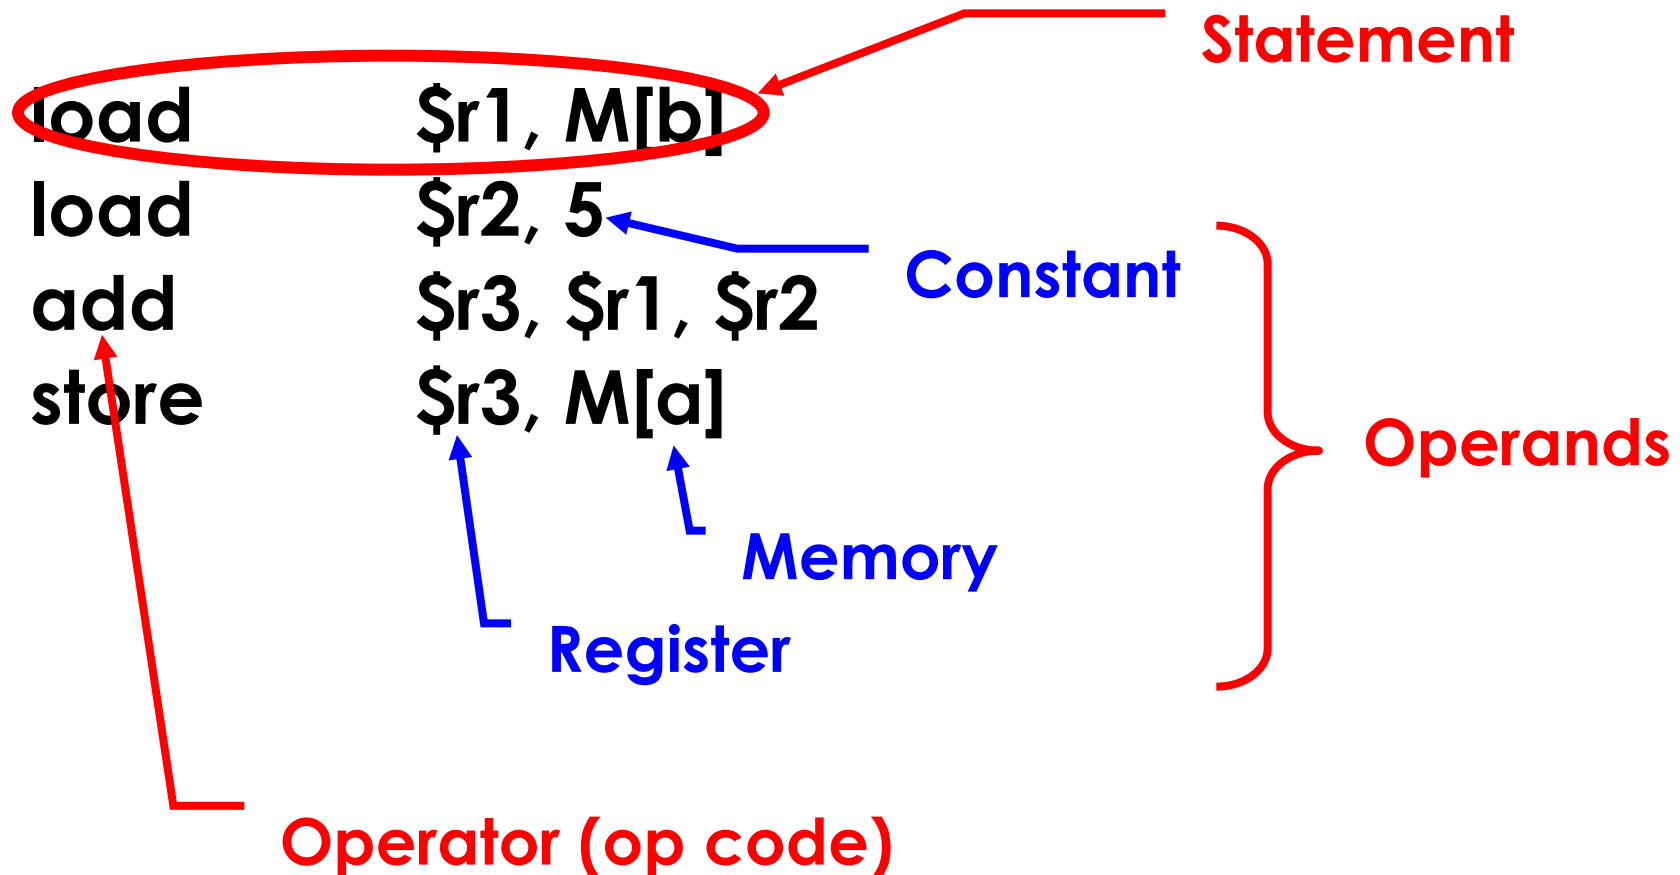
What are specified?

Recall in C Language

- ◆ Operators: +, -, *, /, % (mod), ...
 - $7/4==1$, $7\%4==3$
- ◆ Operands:
 - Variables: `lower`, `upper`, `fahr`, `celsius`
 - Constants: `0`, `1000`, `-17`, `15.4`
- ◆ Assignment statement:
 variable = expression
 - Expressions consist of operators operating on operands, e.g.,
 `celsius = 5*(fahr-32)/9;`
 `a = b+c+d-e;`

When Translating to Assembly ...

a = b + 5;



Components of an ISA


- ◆ Organization of programmable storage
 - registers
 - memory: flat, segmented
 - modes of addressing and accessing data items and instructions
- ◆ Data types and data structures
 - encoding and representation
- ◆ Instruction formats
- ◆ Instruction set (or operation code)
 - ALU, control transfer, exceptional handling

MIPS ISA as an Example

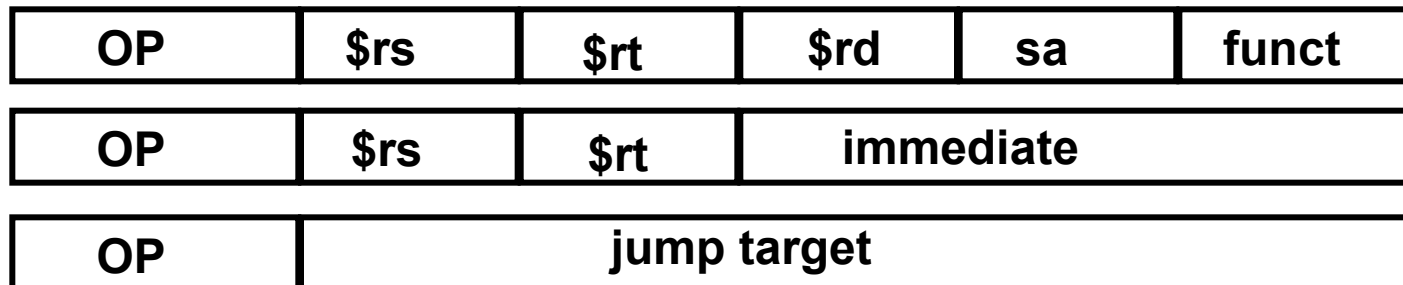
◆ Instruction categories:

- Load/Store
- Computational
- Jump and Branch
- Floating Point
- Memory Management
- Special

Registers



3 Instruction Formats: all 32 bits wide



Outline

- ◆ Instruction set architecture
- ◆ Operands (Sec 2.2, 2.3)
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions
- ◆ Operations
 - Logical
 - Decision making and branches
- ◆ Supporting procedures in hardware
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program
- ◆ A sort example
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets

Operations of Hardware

- ◆ Syntax of basic MIPS arithmetic/logic instructions:

1 2 3 4
`add $s0, $s1, $s2 # f = g + h`

1) operation by name

2) operand getting result (“destination”)

3) 1st operand for operation (“source1”)

4) 2nd operand for operation (“source2”)

- ◆ Each instruction is **32** bits
- ◆ Syntax is rigid: 1 operator, 3 operands
 - Why? Keep hardware simple via regularity
- ◆ *Design Principle 1: Simplicity favors regularity*
 - Regularity makes implementation simpler
 - Simplicity enables higher performance at lower cost

Example

- ◆ How to do the following C statement?

`f = (g + h) - (i + j);`

Compiled MIPS code:

```
add t0, g, h    # temp t0 = g + h
add t1, i, j    # temp t1 = i + j
sub f, t0, t1   # f = t0 - t1
```

Operands and Registers

- ◆ Unlike high-level language, assembly don't use variables
 - => assembly operands are registers
 - Limited number of special locations built directly into the hardware
 - Operations are performed on these
- ◆ Benefits:
 - Registers in hardware => faster than memory
 - Registers are easier for a compiler to use
 - e.g., as a place for temporary storage
 - Registers can hold variables to reduce memory traffic and improve code density (since register named with fewer bits than memory location)

MIPS Registers

- ◆ 32 registers, each is 32 bits wide
 - Why 32? *Design Principle 2: smaller is faster*
 - Groups of 32 bits called a **word** in MIPS
 - Registers are numbered from 0 to 31
 - Each can be referred to by number or name
 - Number references:
\$0, \$1, \$2, ... \$30, \$31
 - By convention, each register also has a name to make it easier to code, e.g.,
\$16 - \$22 → \$s0 - \$s7 (C variables)
\$8 - \$15 → \$t0 - \$t7 (temporary)
- ◆ 32 x 32-bit FP registers (paired DP)
- ◆ Others: HI, LO, PC

Registers Conventions for MIPS

0	zero	constant 0
1	at	reserved for assembler
2	v0	expression evaluation &
3	v1	function results
4	a0	arguments
5	a1	
6	a2	
7	a3	
8	t0	temporary: caller saves
...		(callee can clobber)
15	t7	
16	s0	callee saves
...		(caller can clobber)
23	s7	
24	t8	temporary (cont'd)
25	t9	
26	k0	reserved for OS kernel
27	k1	
28	gp	pointer to global area
29	sp	stack pointer
30	fp	frame pointer
31	ra	return address (HW)

Fig. 2.18

MIPS R2000 Organization

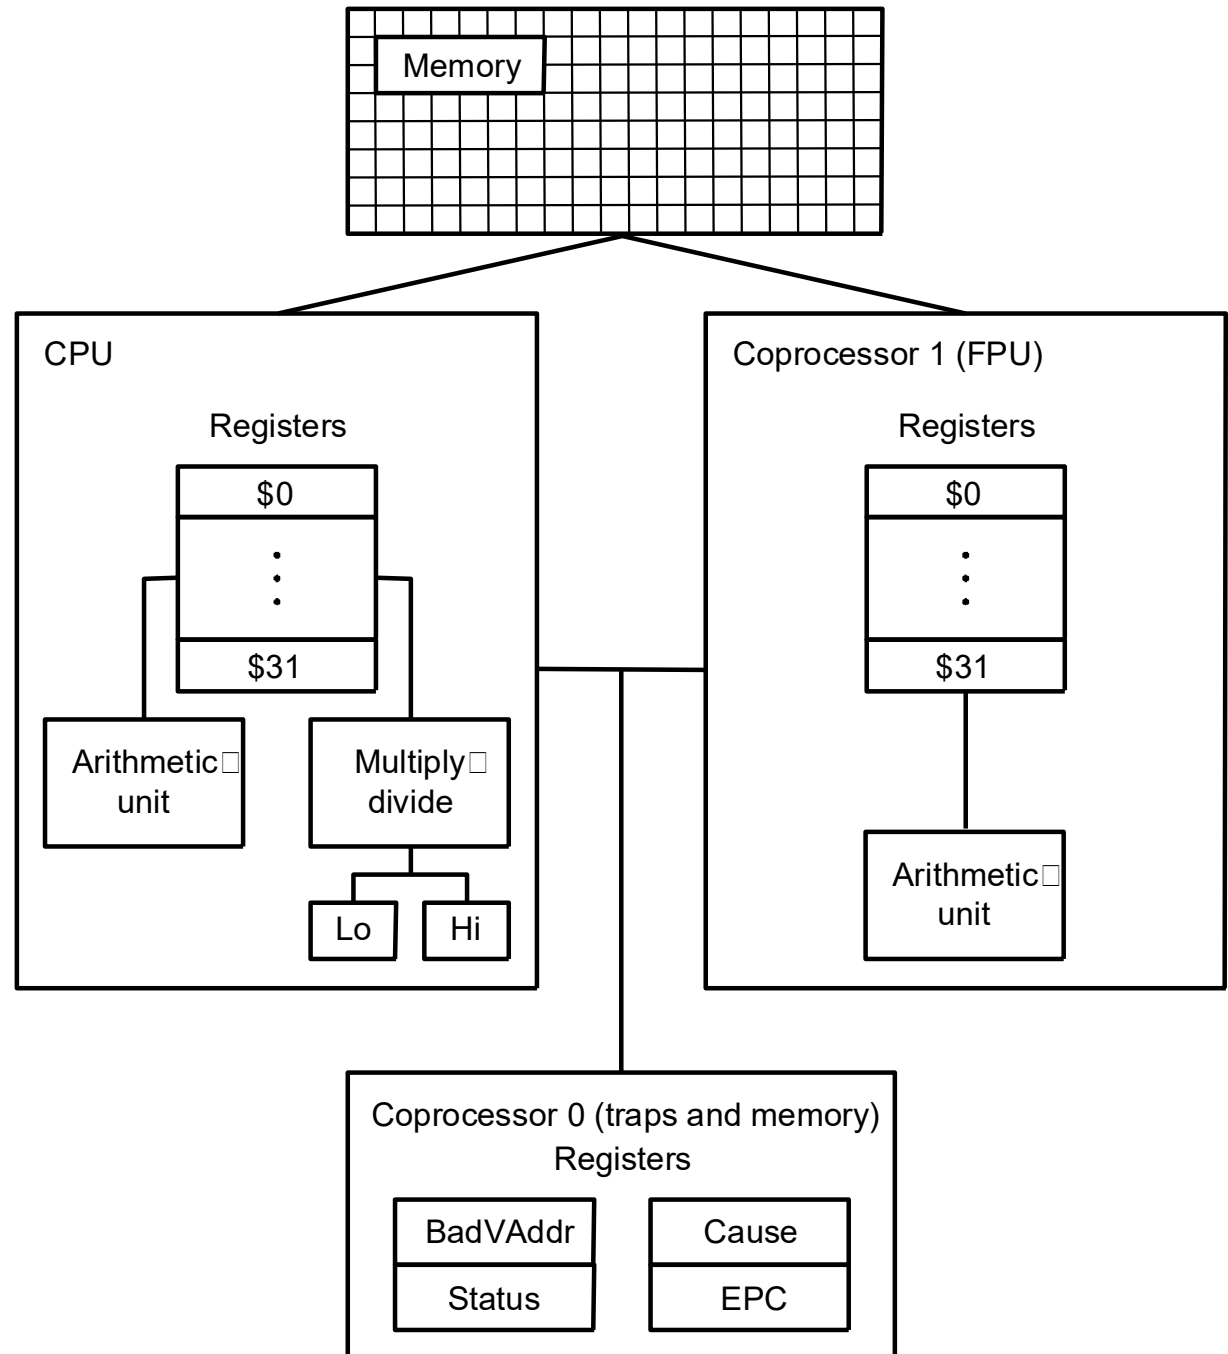


Fig. A.10.1

Example

- ◆ How to do the following C statement?

`f = (g + h) - (i + j);`

- ◆ `f, ..., j` in `$s0, ..., $s4`
- ◆ use intermediate temporary register `t0, t1`

```
add $t0,$s1,$s2    # t0 = g + h
add $t1,$s3,$s4    # t1 = i + j
sub $s0,$t0,$t1    # f=(g+h)-(i+j)
```

Operations of Hardware

- ◆ Syntax of basic MIPS arithmetic/logic instructions:

1 2 3 4
`add $s0, $s1, $s2` # $f = g + h$

1) operation by name

2) operand getting result (“destination”)

3) 1st operand for operation (“source1”)

4) 2nd operand for operation (“source2”)

- ◆ Each instruction is **32** bits
- ◆ Syntax is rigid: 1 operator, 3 operands
 - Why? Keep hardware simple via regularity
- ◆ *Design Principle 1: Simplicity favors regularity*
 - Regularity makes implementation simpler
 - Simplicity enables higher performance at lower cost

Register Architecture

◆ Accumulator (1 register):

1 address: **add A** $\text{acc} \leftarrow \text{acc} + \text{mem}[A]$

1+x address: **addx A** $\text{acc} \leftarrow \text{acc} + \text{mem}[A+x]$

◆ Stack:

0 address: **add** $\text{tos} \leftarrow \text{tos} + \text{next}$

◆ General Purpose Register:

2 address: **add A,B** $\text{EA}(A) \leftarrow \text{EA}(A) + \text{EA}(B)$

3 address: **add A,B,C** $\text{EA}(A) \leftarrow \text{EA}(B) + \text{EA}(C)$

◆ Load/Store: (a special case of GPR)

3 address: **add** $\$ra, \$rb, \$rc$ $\$ra \leftarrow \$rb + \$rc$

load $\$ra, \rb $\$ra \leftarrow \text{mem}[\$rb]$

store $\$ra, \rb $\text{mem}[\$rb] \leftarrow \ra

Register Organization Affects Programming

Code for $C = A + B$ for four register organizations:

Stack	Accumulator	Register (reg-mem)	Register (load-store)
Push A	Load A	Load \$r1,A	Load \$r1,A
Push B	Add B	Add \$r1,B	Load \$r2,B
Add	Store C	Store C,\$r1	Add \$r3,\$r1,\$r2
Pop C			Store C,\$r3

=> Register organization is an attribute of ISA!

Comparison: Byte per instruction? Number of instructions? Cycles per instruction?

Since 1975 all machines use GPRs

Outline

- ◆ Instruction set architecture
- ◆ Operands (Sec 2.2, 2.3)
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions
- ◆ Operations
 - Logical
 - Decision making and branches
- ◆ Supporting procedures in hardware
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program
- ◆ A sort example
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets

Memory Operands

- ◆ C variables map onto registers; what about large data structures like arrays?
 - Memory contains such data structures
- ◆ But MIPS arithmetic instructions operate on registers, not directly on memory
 - Data transfer instructions (lw, sw, ...) to transfer between memory and register
 - A way to address memory operands

Data Transfer: Memory to Register (1/2)

- ◆ To transfer a word of data, need to specify two things:
 - Register: specify this by number (0 - 31)
 - Memory address: more difficult
 - Think of memory as a 1D array
 - Address it by supplying a pointer to a memory address
 - Offset (in bytes) from this pointer
 - The desired memory address is the sum of these two values, e.g., 8 (\$t0)
 - Specifies the memory address pointed to by the value in \$t0, plus 8 bytes (why “bytes”, not “words”?)
 - Each address is **32** bits

Data Transfer: Memory to Register (2/2)

◆ Load Instruction Syntax:

1 2 3 4
`lw $t0, 12($s0)`

- 1) operation name
- 2) register that will receive value
- 3) numerical offset in bytes
- 4) register containing pointer to memory

◆ Example: `lw $t0, 12($s0)`

- `lw` (Load Word, so a word (32 bits) is loaded at a time)
- Take the pointer in `$s0`, add 12 bytes to it, and then load the value from the memory pointed to by this calculated sum into register `$t0`

◆ Notes:


- `$s0` is called the *base register*, 12 is called the *offset*
- Offset is generally used in accessing elements of array:
base register points to the beginning of the array

Example

◆ $\$s0 = 1000$

◆ $\text{lw } \$t0, 12(\$s0)$

◆ $\$t0 = ?$



Data Transfer: Register to Memory

- ◆ Also want to store value from a register into memory
- ◆ Store instruction syntax is identical to Load instruction syntax
- ◆ Example: `sw $t0, 12($s0)`
 - `sw` (meaning Store Word, so 32 bits or one word are **stored** at a time)
 - This instruction will take the pointer in `$s0`, add 12 bytes to it, and then store the value from register `$t0` into the memory address pointed to by the calculated sum


Example

◆ $\$s0 = 1000$

◆ $\$t0 = 25$

◆ $\text{sw } \$t0, 12(\$s0)$

◆ $M[?] = 25$



Compilation with Memory

- ◆ Compile by hand using registers:

\$s1:g, \$s2:h, \$s3:base address of A

$g = h + A[8];$

- ◆ What offset in `lw` to select an array element `A[8]` in a C program?

- $4 \times 8 = 32$ bytes to select `A[8]`

- 1st transfer from memory to register:

`lw $t0, 32($s3) # $t0 gets A[8]`

- Add 32 to `$s3` to select `A[8]`, put into `$t0`

- ◆ Next add it to `h` and place in `g`

`add $s1, $s2, $t0 # $s1 = h+A[8]`

Memory Operand Example 2

◆ C code:

`A[12] = h + A[8];`

- `h` in `$s2`, base address of `A` in `$s3`

◆ Compiled MIPS code:

- Index 8 requires offset of 32

```
lw  $t0, 32($s3)    # load word
add $t0, $s2, $t0
sw  $t0, 48($s3)    # store word
```


Addressing: Byte versus Word

- ◆ Every word in memory has an address, similar to an index in an array
- ◆ Early computers numbered words like C numbers elements of an array:
 - Memory[0], Memory[1], Memory[2], ...

Called the “address” of a word
- ◆ Computers need to access 8-bit bytes as well as words (4 bytes/word)
- ◆ Today, machines address memory as bytes, hence word addresses differ by 4
 - Memory[0], Memory[4], Memory[8], ...
 - This is also why lw and sw use bytes in offset

A Note about Memory: Alignment


- ◆ MIPS requires that all words start at addresses that are multiples of 4 bytes



- ◆ Called Alignment: objects must fall on address that is multiple of their size

Another Note: Endianess

- ◆ Byte order: numbering of bytes within a word
- ◆ Big Endian: address of most significant byte at least address of a word
 - IBM 360/370, Motorola 68k, MIPS, Sparc, HP PA
- ◆ Little Endian: address of least significant byte at least address
 - Intel 80x86, DEC Vax, DEC Alpha (Windows NT)



Role of Registers vs. Memory

- ◆ What if more variables than registers?
 - Compiler tries to keep most frequently used variables in registers
 - Writes less common variables to memory: spilling
- ◆ Why not keep all variables in memory?
 - Smaller is faster:
registers are faster than memory
 - Registers more versatile:
 - MIPS arithmetic instructions can read 2 registers, operate on them, and write 1 per instruction
 - MIPS data transfers only read or write 1 operand per instruction, and no operation

Outline

- ◆ Instruction set architecture
- ◆ Operands (Sec 2.2, 2.3)
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions
- ◆ Operations
 - Logical
 - Decision making and branches
- ◆ Supporting procedures in hardware
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program
- ◆ A sort example
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets

Constants

- ◆ Small constants used frequently (50% of operands)

e.g., $A = A + 5;$

$B = B + 1;$

$C = C - 18;$

- ◆ Put 'typical constants' in memory and load them

- ◆ Constant data specified in an instruction:

`addi $29, $29, 4`

`slti $8, $18, 10`

`andi $29, $29, 6`

`ori $29, $29, 4`

- ◆ *Design Principle 3: Make the common case fast*

Immediate Operands

◆ Immediate: numerical constants

- Often appear in code, so there are special instructions for them
- **Add Immediate:**

$f = g + 10$ (in C)

`addi $s0, $s1, 10` (in MIPS)

where `$s0, $s1` are associated with `f, g`

- Syntax similar to `add` instruction, except that last argument is a number instead of a register
- No subtract immediate instruction
 - Just use a negative constant

`addi $s2, $s1, -1`

The Constant Zero

- ◆ The number zero (0), appears very often in code; so we define register zero
- ◆ MIPS register 0 (\$zero) is the constant 0
 - Cannot be overwritten
 - This is defined in hardware, so an instruction like
`addi $0, $0, 5` will not do anything
- ◆ Useful for common operations
 - E.g., move between registers
`add $t2, $s1, $zero`

Outline

- ◆ Instruction set architecture
- ◆ Operands
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers (Sec 2.4, read by students)
- ◆ Representing instructions
- ◆ Operations
 - Logical
 - Decision making and branches
- ◆ Supporting procedures in hardware
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program
- ◆ A sort example
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets

Outline

- ◆ Instruction set architecture
- ◆ Operands
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions (Sec 2.5)
- ◆ Operations
 - Logical
 - Decision making and branches
- ◆ Supporting procedures in hardware
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program
- ◆ A sort example
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets

Instructions as Numbers

- ◆ Currently we only work with words (32-bit blocks):
 - Each register is a word
 - `lw` and `sw` both access memory one word at a time
- ◆ So how do we represent instructions?
 - Remember: Computer only understands 1s and 0s, so “`add $t0, $0, $0`” is meaningless to hardware
 - MIPS wants simplicity: since data is in words, make instructions be words...

MIPS Instruction Format

- ◆ One instruction is 32 bits
=> divide instruction word into “fields”
 - Each field tells computer something about instruction
- ◆ We could define different fields for each instruction, but MIPS is based on simplicity, so define 3 basic types of instruction formats:
 - *R-format*: for register
 - *I-format*: for immediate, and `lw` and `sw` (since the offset counts as an immediate)
 - *J-format*: for jump

R-Format Instructions (1/2)

- ◆ Define the following “fields”:

6	5	5	5	5	6
opcode	rs	rt	rd	shamt	funct

- opcode: partially specifies what instruction it is (Note: 0 for all R-Format instructions)
- funct: combined with opcode to specify the instruction
Question: Why aren't opcode and funct a single 12-bit field?
- rs (Source Register): *generally* used to specify register containing first operand
- rt (Target Register): *generally* used to specify register containing second operand
- rd (Destination Register): *generally* used to specify register which will receive result of computation

R-Format Instructions (2/2)

◆ Notes about register fields:

- Each register field is exactly 5 bits, which means that it can specify any unsigned integer in the range 0-31. Each of these fields specifies one of the 32 registers by number.

◆ Final field:

- **shamt**: contains the amount a shift instruction will shift by. Shifting a 32-bit word by more than 31 is useless, so this field is only 5 bits
- This field is set to 0 in all but the shift instructions

R-format Example

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

add \$t0, \$s1, \$s2

Special	\$s1	\$s2	\$t0	0	add
0	17	18	8	0	32
000000	10001	10010	01000	00000	100000

$$00000010001100100100000000100000_2 = 02324020_{16}$$

Hexadecimal

◆ Base 16

- Compact representation of bit strings
- 4 bits per hex digit

0	0000	4	0100	8	1000	c	1100
1	0001	5	0101	9	1001	d	1101
2	0010	6	0110	a	1010	e	1110
3	0011	7	0111	b	1011	f	1111

◆ Example: eca8 6420

- 1110 1100 1010 1000 0110 0100 0010 0000

I-Format Instructions

- ◆ Define the following “fields”:

6	5	5	16
opcode	rs	rt	immediate

- opcode: uniquely specifies an I-format instruction
- rs: specifies the *only* register operand
- rt: specifies register which will receive result of computation (*target register*)
- addi, slti, immediate is **sign-extended** to 32 bits, and treated as a signed integer
- 16 bits → can be used to represent immediate up to 2^{16} different values

MIPS I-format Instructions

- ◆ ***Design Principle 4: Good design demands good compromises***
 - Different formats complicate decoding, but allow 32-bit instructions uniformly
 - Keep formats as similar as possible

I-Format Example 1

◆ MIPS Instruction:

`addi $21, $22, -50`

- opcode = 8 (look up in table)
- rs = 22 (register containing operand)
- rt = 21 (target register)
- immediate = -50 (by default, this is decimal)

decimal representation:

8	22	21	-50
---	----	----	-----

binary representation:

001000	10110	10101	1111111111001110
--------	-------	-------	------------------

I-Format Example 2

◆ MIPS Instruction:

`lw $t0, 1200($t1)`

- opcode = 35 (look up in table)
- rs = 9 (base register)
- rt = 8 (destination register)
- immediate = 1200 (offset)

decimal representation:

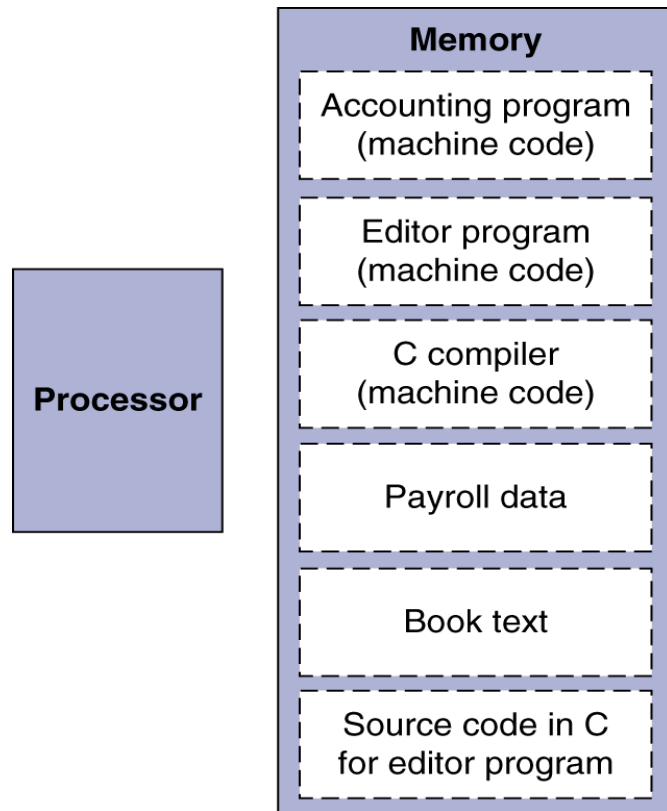
35	9	8	1200
----	---	---	------

binary representation:

100011	01001	01000	0000010010110000
--------	-------	-------	------------------

Stored Program Computers

The BIG Picture



- ◆ Instructions represented in binary, just like data
- ◆ Instructions and data stored in memory
- ◆ Programs can operate on programs
 - e.g., compilers, linkers, ...
- ◆ Binary compatibility allows compiled programs to work on different computers
 - Standardized ISAs

Outline

- ◆ Instruction set architecture
- ◆ Operands
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions
- ◆ Operations
 - Logical (Sec 2.6)
 - Decision making and branches
- ◆ Supporting procedures in hardware
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program
- ◆ A sort example
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets

Bitwise Operations

- ◆ Up until now, we've done arithmetic (add, sub, addi) and memory access (lw and sw)
- ◆ All of these instructions view contents of register as a single quantity (such as a signed or unsigned integer)
- ◆ **New perspective**: View contents of register as 32 bits rather than as a single 32-bit number
- ◆ Since registers are composed of 32 bits, we may want to access individual bits rather than the whole.
- ◆ Introduce two new classes of instructions:
 - **Shift instructions**
 - **Logical operators**

Logical Operations

◆ Instructions for bitwise manipulation

Operation	C	Java	MIPS
Shift left	<<	<<	sll
Shift right	>>	>>>	srl
Bitwise AND	&	&	and, andi
Bitwise OR			or, ori
Bitwise NOT	~	~	nor

◆ Useful for extracting and inserting groups of bits in a word

Shift Operations

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

- ◆ **shamt: how many positions to shift**
- ◆ **Shift left logical**
 - Shift left and fill with 0 bits
 - $sl\ i$ by i bits multiplies by 2^i
- ◆ **Shift right logical**
 - Shift right and fill with 0 bits
 - $sr\ i$ by i bits divides by 2^i (unsigned only)

Shift Instructions (1/3)

◆ Shift Instruction Syntax:

1 2 3 4
sll **\$t2, \$s0, 4**

1) operation name

2) register that will receive value

3) first operand (register)

4) shift amount (constant)

◆ MIPS has three shift instructions:

- **sll** (shift left logical): shifts left, fills empties with 0s
- **sr1** (shift right logical): shifts right, fills empties with 0s
- **sra** (shift right arithmetic): shifts right, fills empties by sign extending

Shift Instructions (2/3)

- ◆ Move (shift) all the bits in a word to the left or right by a number of bits, filling the emptied bits with 0s.
- ◆ Example: shift right by 8 bits

0001 0010 0011 0100 0101 0110 0111 1000

0000 0000 0001 0010 0011 0100 0101 0110


- ◆ Example: shift **left** by 8 bits

0001 0010 0011 0100 0101 0110 0111 1000


0011 0100 0101 0110 0111 1000 0000 0000

Shift Instructions (3/3)

- ◆ Example: shift right arithmetic by 8 bits



- ◆ Example: shift right arithmetic by 8 bits



Uses for Shift Instructions

- ◆ Shift for multiplication: in binary
 - Multiplying by 4 is same as shifting left by 2:
 - $11_2 \times 100_2 = 1100_2$
 - $1010_2 \times 100_2 = 101000_2$
 - Multiplying by 2^n is same as shifting left by n
- ◆ Since shifting is so much faster than multiplication (you can imagine how complicated multiplication is), a good compiler usually notices when C code multiplies by a power of 2 and compiles it to a shift instruction:

`a *= 8;` (in C)

would compile to:

`sll $s0, $s0, 3` (in MIPS)

AND Operations

- ◆ Useful to mask bits in a word
 - Select some bits, clear others to 0
- and \$t0, \$t1, \$t2

\$t2	0000 0000 0000 0000 0000 1101 1100 0000
\$t1	0000 0000 0000 0000 0011 1100 0000 0000
\$t0	0000 0000 0000 0000 0000 1100 0000 0000

OR Operations

- ◆ Useful to include bits in a word
 - Set some bits to 1, leave others unchanged
- or \$t0, \$t1, \$t2

\$t2	0000 0000 0000 0000 0000 1101 1100 0000
\$t1	0000 0000 0000 0000 0011 1100 0000 0000
\$t0	0000 0000 0000 0000 0011 1101 1100 0000

NOT Operations

- ◆ Useful to invert bits in a word
 - Change 0 to 1, and 1 to 0
- ◆ MIPS has NOR 3-operand instruction
 - $a \text{ NOR } b == \text{NOT} (a \text{ OR } b)$

`nor $t0, $t1, $zero`

Register 0: always
read as zero

\$t1 0000 0000 0000 0000 0011 1100 0000 0000

\$t0 1111 1111 1111 1111 1100 0011 1111 1111

So Far...

- ◆ All instructions have allowed us to manipulate data.
- ◆ So we've built a calculator.
- ◆ In order to build a computer, we need ability to make decisions...

Outline

- ◆ Instruction set architecture
- ◆ Operands
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions
- ◆ Operations
 - Logical
 - Decision making and branches (Sec 2.7)
- ◆ Supporting procedures in hardware
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program
- ◆ A sort example
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets

MIPS Decision Instructions

`beq register1, register2, L1`

- ◆ Decision instruction in MIPS:

`beq register1, register2, L1`

“Branch if (registers are) equal”

meaning :

`if (register1==register2) goto L1`

- ◆ Complementary MIPS decision instruction

`bne register1, register2, L1`

“Branch if (registers are) not equal”

meaning :

`if (register1!=register2) goto L1`

- ◆ These are called conditional branches

MIPS Goto Instruction

`j label`

- ◆ MIPS has an **unconditional branch**:

`j label`

- Called a Jump Instruction: jump directly to the given label without testing any condition
- meaning :
`goto label`

- ◆ Technically, it's the same as:

`beq $0, $0, label`

since it always satisfies the condition

- ◆ It has the j-type instruction format

Compiling C if into MIPS

- ◆ Compile by hand

```
if (i == j) f=g+h;  
else f=g-h;
```


- ◆ Use this mapping:

```
f, g, ., j : $s0, $s1, $s2,  
           $s3, $s4
```

- ◆ Final compiled MIPS code:

	<code>bne</code>	<code>\$s3, \$s4, Else</code>	<code># branch i != j</code>
	<code>add</code>	<code>\$s0, \$s1, \$s2</code>	<code># f=g+h (true)</code>
	<code>j</code>	<code>Exit</code>	<code># go to Exit</code>
<code>Else:</code>	<code>sub</code>	<code>\$s0, \$s1, \$s2</code>	<code># f=g-h (false)</code>
<code>Exit:</code>			

Note: Compiler automatically creates labels to handle decisions (branches) appropriately



Compiling Loop Statements

◆ C code:

```
while (save[i] == k) i += 1;
```

- i in \$s3, k in \$s5, address of save in \$s6


◆ Compiled MIPS code:

```
Loop:      sll    $t1, $s3, 2          # $t1 = i x 4
           add    $t1, $t1, $s6       # $t1 = addr of save[i]
           lw     $t0, 0($t1)         # $t0 = save[i]
           bne    $t0, $s5, Exit      # if save[i] != k goto Exit
           addi   $s3, $s3, 1         # i = i + 1
           j      Loop               # goto Loop

Exit: ...
```

Basic Blocks

- ◆ A basic block is a sequence of instructions with
 - No embedded branches (except at end)
 - No branch targets (except at beginning)



- ◆ A compiler identifies basic blocks for optimization
- ◆ An advanced processor can accelerate execution of basic blocks

Inequalities in MIPS

- ◆ Until now, we've only tested equalities (`==` and `!=` in C), but general programs need to test `<` and `>`
- ◆ Set on Less Than:
 - `slt rd, rs, rt`
 - if (`rs < rt`) `rd = 1`; else `rd = 0`;
 - `slti rt, rs, constant`
 - if (`rs < constant`) `rt = 1`; else `rt = 0`;

Compile by hand: `if (g < h) goto Less;`
Let `g: $s0, h: $s1`

```
slt $t0,$s0,$s1    # $t0 = 1 if g<h
bne $t0,$0,Less     # goto Less if $t0!=0
```

- ◆ MIPS has no “branch on less than” => too complex

Branch Instruction Design

- ◆ Why not b1t, bge, etc?
- ◆ Hardware for $<$, \geq , ... slower than $=$, \neq
 - Combining with branch involves more work per instruction, requiring a slower clock
 - All instructions penalized!
- ◆ beq and bne are the common case
- ◆ This is a good design compromise

Signed vs. Unsigned

- ◆ Signed comparison: `slt`, `slti`
- ◆ Unsigned comparison: `sltu`, `sltui`
- ◆ Example
 - `$s0 = 1111 1111 1111 1111 1111 1111 1111 1111`
 - `$s1 = 0000 0000 0000 0000 0000 0000 0000 0001`
 - `slt $t0, $s0, $s1 # signed`
 - $-1 < +1 \Rightarrow \$t0 = 1$
 - `sltu $t0, $s0, $s1 # unsigned`
 - $+4,294,967,295 > +1 \Rightarrow \$t0 = 0$

Outline

- ◆ Instruction set architecture (using MIPS ISA as an example)
- ◆ Operands
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions
- ◆ Operations
 - Logical
 - Decision making and branches
- ◆ Supporting procedures in hardware (Sec. 2.8)
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program
- ◆ A sort example
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets

Procedure Calling

◆ Steps required

Caller:

1. Place parameters in registers
2. Transfer control to procedure

Callee:

3. Acquire storage for procedure
4. Perform procedure's operations
5. Place result in register for caller
6. Return to place of call

C Function Call Bookkeeping

```
sum = leaf_example(a,b,c,d) . . .
```

```
int leaf_example (int g, h, i, j)
{ int f;
  f = (g + h) - (i + j);
  return f;
}
```

- ◆ Return address \$ra
- ◆ Procedure address Labels
- ◆ Arguments \$a0, \$a1, \$a2, \$a3
- ◆ Return value \$v0, \$v1
- ◆ Local variables \$s0, \$s1, ..., \$s7

Note the use of **register conventions**

Registers Conventions for MIPS

0	zero	constant 0
1	at	reserved for assembler
2	v0	expression evaluation &
3	v1	function results
4	a0	arguments
5	a1	
6	a2	
7	a3	
8	t0	temporary: caller saves
...		(callee can clobber)
15	t7	
16	s0	callee saves
...		(caller can clobber)
23	s7	
24	t8	temporary (cont'd)
25	t9	
26	k0	reserved for OS kernel
27	k1	
28	gp	pointer to global area
29	sp	stack pointer
30	fp	frame pointer
31	ra	return address (HW)

Fig. 2.18

Procedure Call Instructions

- ◆ Procedure call: jump and link
jal ProcedureLabel
 - Address of following instruction put in \$ra
 - Jumps to target address (i.e., ProcedureLabel)
- ◆ Procedure return: jump register
jr \$ra
 - Copies \$ra to program counter
 - Can also be used for computed jumps
 - e.g., for case/switch statements
 - ✧ Jump table is an array of addresses corresponding to labels in codes
 - ✧ Load appropriate entry to register
 - ✧ Jump register

Caller's Code

◆ . . .

sum = leaf_example(a,b,c,d)

. . .

◆ MIPS code: a, ..., d in \$s0, ..., \$s3, and sum in \$s4

:

add	\$a0,	\$0,	\$s0
add	\$a1,	\$0,	\$s1
add	\$a2,	\$0,	\$s2
add	\$a3,	\$0,	\$s3
jal	leaf_example		
add	\$s4,	\$0,	\$v0

Move a,b,c,d to a0..a3

Jump to leaf_example

Move result in v0 to sum

:

Leaf Procedure Example

◆ C code:

```
int leaf_example (int g, h, i, j)
{ int f;
  f = (g + h) - (i + j);
  return f;
}
```

- Arguments g, ..., j in \$a0, ..., \$a3
- f in \$s0 (hence, need to save \$s0 on stack)
- Save \$t1 and \$t2
- Result in \$v0

Leaf Procedure Example

◆ MIPS code:

leaf_example:			
addi	\$sp,	\$sp, -12	
sw	\$s0,	0(\$sp)	Save \$s0 \$t0 \$t1 on stack
sw	\$t0,	4(\$sp)	
sw	\$t1,	8(\$sp)	
add	\$t0,	\$a0, \$a1	Procedure body
add	\$t1,	\$a2, \$a3	
sub	\$s0,	\$t0, \$t1	
add	\$v0,	\$s0, \$zero	Result
lw	\$s0,	0(\$sp)	
lw	\$t0,	4(\$sp)	Restore \$s0 \$t0 \$t1
lw	\$t1,	8(\$sp)	
addi	\$sp,	\$sp, 12	
jr	\$ra		Return

Leaf Procedure Example

◆ C code:

```
int leaf_example (int g, h, i, j)
{ int f;
  f = (g + h) - (i + j);
  return f;
}
```

- Arguments g, ..., j in \$a0, ..., \$a3
- f in \$s0 (hence, need to save \$s0 on stack)
- **\$t1 and \$t2 are not saved on stack**
- Result in \$v0

Leaf Procedure Example

◆ MIPS code:

leaf_example:		
addi	\$sp,	\$sp, -4
sw	\$s0,	0(\$sp)
add	\$t0,	\$a0, \$a1
add	\$t1,	\$a2, \$a3
sub	\$s0,	\$t0, \$t1
add	\$v0,	\$s0, \$zero
lw	\$s0,	0(\$sp)
addi	\$sp,	\$sp, 4
jr	\$ra	

Save \$s0 on stack

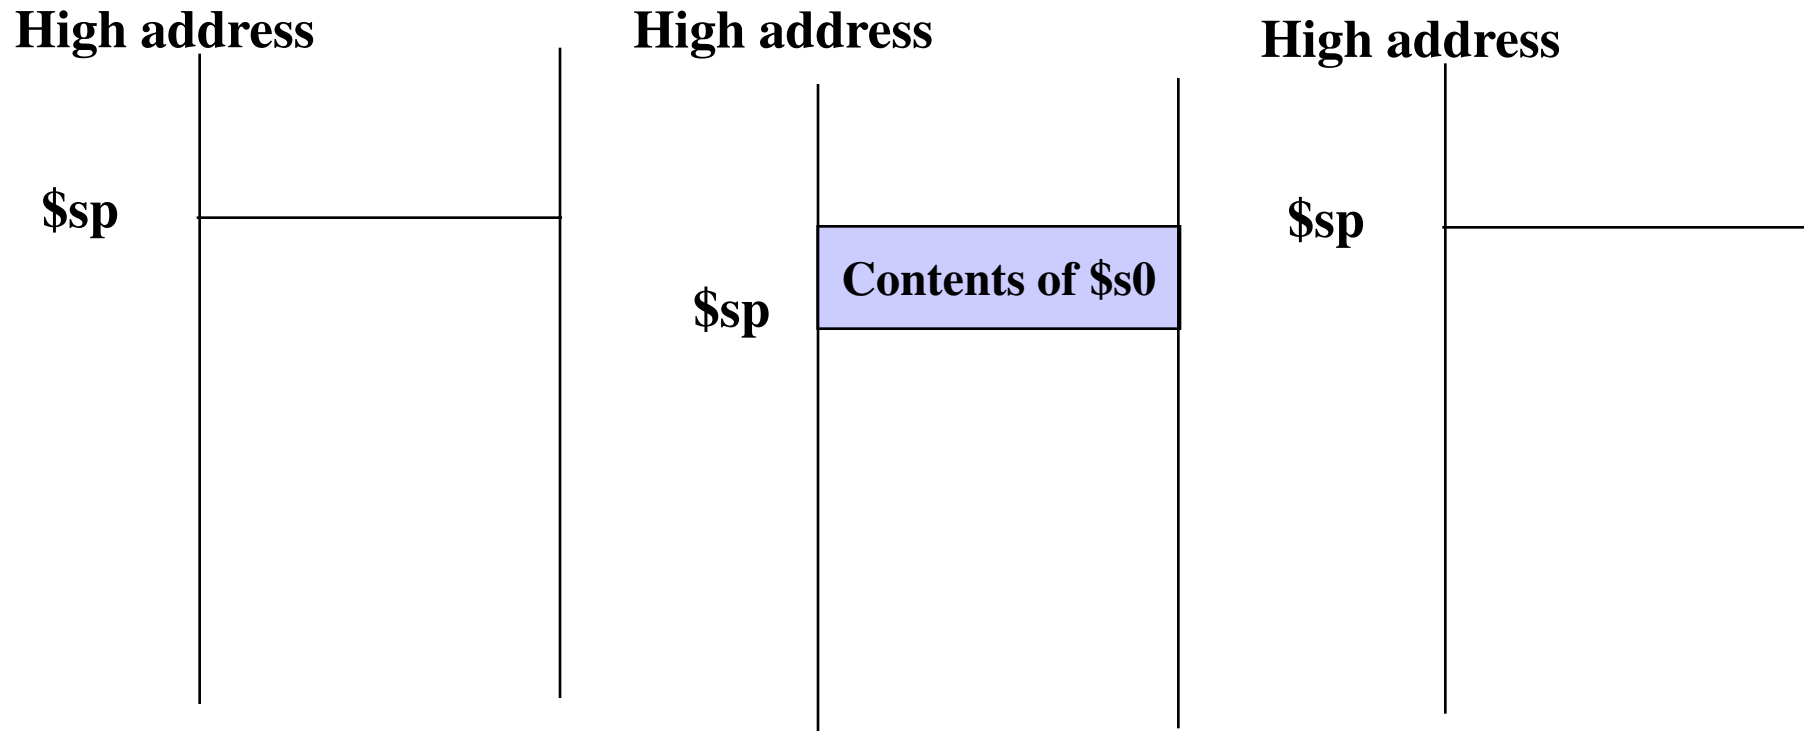
Procedure body

Result

Restore \$s0

Return

Local Data on the Stack



Non-Leaf Procedures

- ◆ Procedures that call other procedures
- ◆ For nested call, caller needs to save on the stack:
 - Its return address
 - Any arguments and temporaries needed after the call (because callee will not save them)
- ◆ Restore from the stack after the call

Non-Leaf Procedure Example

◆ C code:

```
int fact (int n)
{
    if (n < 1) return 1;
    else return n * fact(n - 1);
}
```

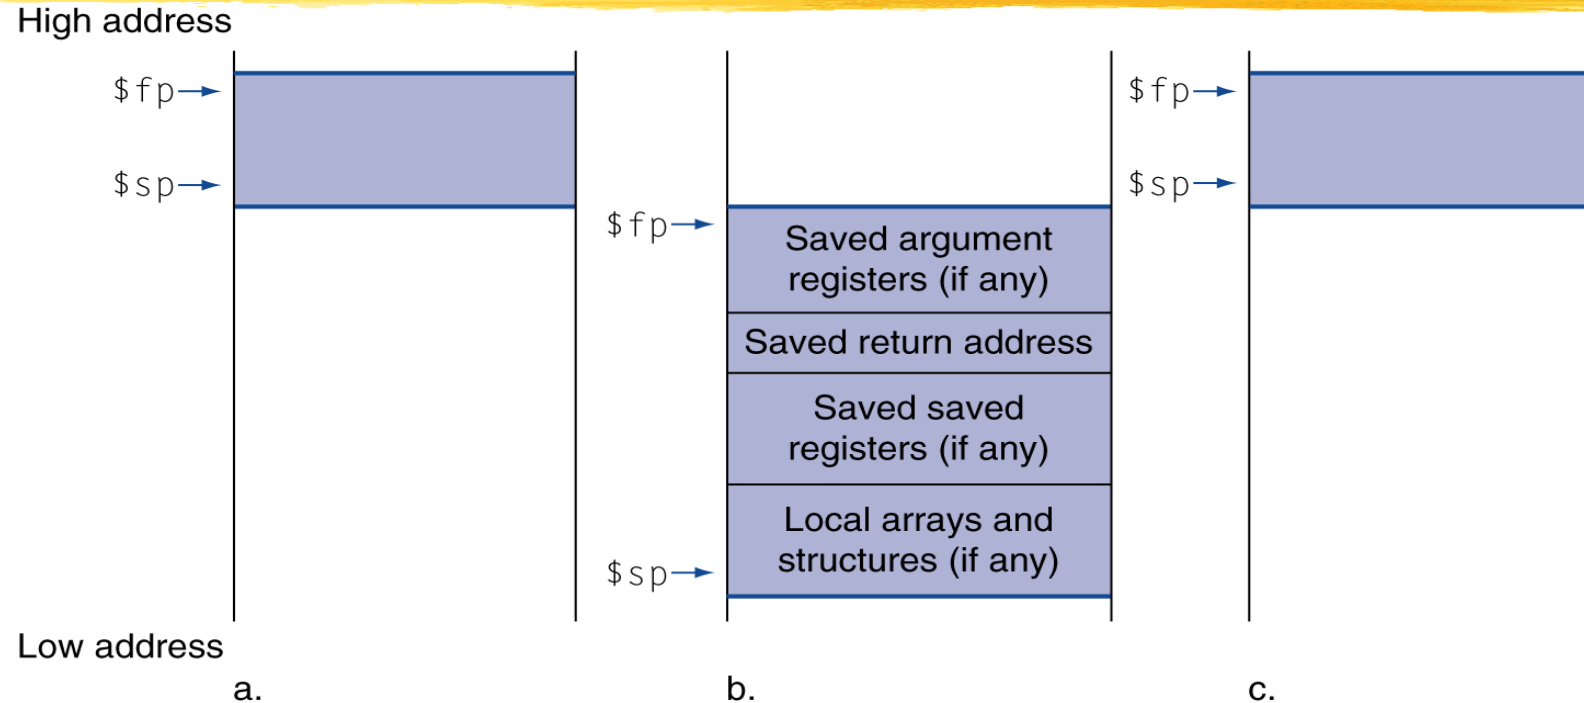
- Argument n in \$a0
- Result in \$v0

Non-Leaf Procedure Example

◆ MIPS code:

fact:		
addi	\$sp, \$sp, -8	# adjust stack for 2 items
sw	\$ra, 4(\$sp)	# save return address
sw	\$a0, 0(\$sp)	# save argument
slti	\$t0, \$a0, 1	# test for n < 1
beq	\$t0, \$zero, L1	
addi	\$v0, \$zero, 1	# if so, result is 1
addi	\$sp, \$sp, 8	# pop 2 items from stack
jr	\$ra	# and return
L1:	addi \$a0, \$a0, -1	# else decrement n
	jal fact	# recursive call
lw	\$a0, 0(\$sp)	# restore original n
lw	\$ra, 4(\$sp)	# and return address
addi	\$sp, \$sp, 8	# pop 2 items from stack
mul	\$v0, \$a0, \$v0	# multiply to get result
jr	\$ra	# and return

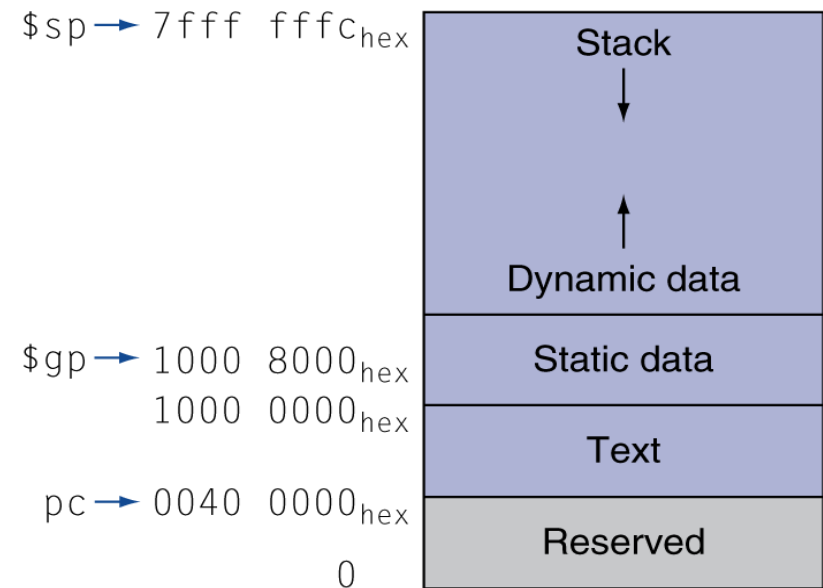
Local Data on the Stack



- ◆ **Local data allocated by callee**
 - e.g., C automatic variables
- ◆ **Procedure frame (activation record)**
 - Used by some compilers to manage stack storage

Memory Layout

- ◆ Text: program code
- ◆ Static data: global variables
 - e.g., static variables in C, constant arrays and strings
 - \$gp initialized to address allowing \pm offsets into this segment
- ◆ Dynamic data: heap
 - E.g., malloc in C, new in Java
- ◆ Stack: automatic storage



Why Procedure Conventions?

- ◆ **Definitions**
 - **Caller:** function making the call, using jal
 - **Callee:** function being called
- ◆ **Procedure conventions as a contract between the Caller and the Callee**
- ◆ **If both the Caller and Callee obey the procedure conventions, there are significant benefits**
 - **People who have never seen or even communicated with each other can write functions that work together**
 - **Recursion functions work correctly**

Caller's Rights, Callee's Rights

◆ Callees' rights:

- Right to use VAT registers freely
- Right to assume arguments are passed correctly

◆ To ensure callees's right, caller saves registers:

- Return address \$ra
- Arguments \$a0, \$a1, \$a2, \$a3
- Return value \$v0, \$v1
- \$t Registers \$t0 - \$t9

◆ Callers' rights:

- Right to use S registers without fear of being overwritten by callee
- Right to assume return value will be returned correctly

◆ To ensure caller's right, callee saves registers:

- \$s Registers \$s0 - \$s7

Contract in Function Calls (1/2)

- ◆ **Caller's responsibilities (how to call a function)**
 - Slide `$sp` down to reserve memory:
e.g., `addi $sp, $sp, -28`
 - Save `$ra` on stack because `jal` clobbers it:
e.g., `sw $ra, 24($sp)`
 - If still need their values after function call, save `$v`, `$a`, `$t` on stack or copy to `$s` registers
 - Put first 4 words of arguments in `$a0-3`, additional arguments go on stack: "`a4`" is 16 (`$sp`)
 - `jal` to the desired function
 - Receive return values in `$v0`, `$v1`
 - Undo first steps: e.g. `lw $t0, 20($sp) lw $ra, 24($sp) addi $sp, $sp, 28`

Contract in Function Calls (2/2)

- ◆ Callee's responsibilities (i.e. how to write a function)
 - If using `$s` or big local structures, slide `$sp` down to reserve memory: e.g., `addi $sp, $sp, -48`
 - If using `$s`, save before using: e.g.,
`sw $s0, 44($sp)`
 - Receive arguments in `$a0-3`, additional arguments on stack
 - Run the procedure body
 - If not void, put return values in `$v0, 1`
 - If applicable, undo first two steps: e.g.,
`lw $s0, 44($sp)` `addi $sp, $sp, 48`
 - `jr $ra`

Outline

- ◆ Instruction set architecture (using MIPS ISA as an example)
- ◆ Operands
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions
- ◆ Operations
 - Logical
 - Decision making and branches
- ◆ Supporting procedures in hardware
- ◆ Communicating with people (Sec. 2.9)
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program
- ◆ A sort example
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets

Character Data

- ◆ **Byte-encoded character sets**
 - **ASCII: 128 characters**
 - 95 graphic, 33 control
 - **Latin-1: 256 characters**
 - ASCII, +96 more graphic characters
- ◆ **Unicode: 32-bit character set**
 - **Used in Java, C++ wide characters, ...**
 - **Most of the world's alphabets, plus symbols**
 - **UTF-8, UTF-16: variable-length encodings**

Byte/Halfword Operations

- ◆ Could use bitwise operations
- ◆ MIPS byte/halfword load/store
- ◆ String processing is a common case

`lb rt, offset(rs)` `lh rt, offset(rs)`

- Sign extend to 32 bits in `rt`


`lbu rt, offset(rs)` `lhu rt, offset(rs)`

- Zero extend to 32 bits in `rt`


`sb rt, offset(rs)` `sh rt, offset(rs)`

- Store just rightmost byte/halfword

Load Byte Signed/Unsigned




lb \$t1, 0(\$t0)



Sign-extended

lbu \$t2, 0(\$t0)



Zero-extended

String Copy Example

◆ C code (naïve):

- Null-terminated string

```
void strcpy (char x[], char y[])  
{ int i;  
  i = 0;  
  while ((x[i]=y[i])!='\0')  
    i += 1;  
}
```

- Addresses of x, y in \$a0, \$a1
- i in \$s0

String Copy Example

◆ MIPS code:

strcpy:		
	addi \$sp, \$sp, -4	# adjust stack for 1 item
	sw \$s0, 0(\$sp)	# save \$s0
	add \$s0, \$zero, \$zero	# i = 0
L1:	add \$t1, \$s0, \$a1	# addr of y[i] in \$t1
	lbu \$t2, 0(\$t1)	# \$t2 = y[i]
	add \$t3, \$s0, \$a0	# addr of x[i] in \$t3
	sb \$t2, 0(\$t3)	# x[i] = y[i]
	beq \$t2, \$zero, L2	# exit loop if y[i] == 0
	addi \$s0, \$s0, 1	# i = i + 1
	j L1	# next iteration of loop
L2:	lw \$s0, 0(\$sp)	# restore saved \$s0
	addi \$sp, \$sp, 4	# pop 1 item from stack
	jr \$ra	# and return

Outline

- ◆ Instruction set architecture (using MIPS ISA as an example)
- ◆ Operands
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions
- ◆ Operations
 - Logical
 - Decision making and branches
- ◆ Supporting procedures in hardware
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses (Sec. 2.10)
- ◆ Translating and starting a program
- ◆ A sort example
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets


32-bit Constants

- ◆ Most constants are small
 - 16-bit immediate is sufficient
- ◆ For the occasional 32-bit constant
- ◆ Load Upper Immediate:
lui rt, constant
 - Copies 16-bit constant to left 16 bits of rt
 - Clears right 16 bits of rt to 0

machine version of lui

lui \$s0, 61

ori \$s0, \$s0, 2304



Branch Addressing (1)

◆ Use I-format:

opcode	rs	rt	immediate
--------	----	----	-----------

- opcode specifies beq or bne
- Rs and Rt specify registers to compare

◆ What can *immediate* specify? PC-relative addressing

- *Immediate* is only 16 bits, but PC is 32-bit
=> *immediate* cannot specify entire address
- Loops are generally small: < 50 instructions
 - Though we want to branch to anywhere in memory, a single branch only need to change **PC** by a small amount
- How to use PC-relative addressing
 - 16-bit *immediate* as a signed two's complement integer to be added to the PC if branch taken

Branch Addressing (2)

- ◆ *Immediate* specifies **word address**
 - Instructions are word aligned (byte address is always a multiple of 4, i.e., it ends with 00 in binary)
 - The number of bytes to add to the PC will always be a multiple of 4
 - Specify the *immediate* in words (confusing?)
 - Now, we can branch $\pm 2^{15}$ **words** from the PC (or $\pm 2^{17}$ bytes), handle loops 4 times as large
- ◆ *Immediate* specifies **PC + 4**
 - Due to hardware, add *immediate* to (PC+4), not to PC
 - If branch not taken: $PC = PC + 4$
 - If branch taken: $PC = (PC+4) + (\text{immediate} * 4)$

Branch Example

◆ MIPS Code:

```
Loop: beq    $9, $0, End
      add    $8, $8, $10
      addi   $9, $9, -1
      j      Loop
```

End:

◆ Branch is I-Format:

opcode	rs	rt	immediate
--------	----	----	-----------

opcode = 4 (look up in table)

rs = 9 (first operand)

rt = 0 (second operand)

immediate = ???

- Number of instructions to add to (or subtract from) the PC, starting at the instruction **following** the branch
=> immediate = ?

Branch Example

◆ MIPS Code:

```
Loop: beq    $9, $0, End
      add    $8, $8, $10
      addi   $9, $9, -1
      j      Loop
```

End:

decimal representation:

4	9	0	3
---	---	---	---

binary representation:

000100	01001	00000	000000000000000011
--------	-------	-------	--------------------

Jump Addressing (1/3)

- ◆ For branches, we assumed that we won't want to branch too far, so we can specify change in PC.
- ◆ For general jumps (j and jal), we may jump to anywhere in memory.
- ◆ Ideally, we could specify a 32-bit memory address to jump to.
- ◆ Unfortunately, we can't fit both a 6-bit opcode and a 32-bit address into a single 32-bit word, so we compromise.

Jump Addressing (2/3)

- ◆ Define “fields” of the following number of bits each:

6 bits	26 bits
--------	---------

- ◆ As usual, each field has a name:

opcode	target address
--------	----------------

- ◆ **Key concepts:**

- Keep opcode field identical to R-format and I-format for consistency
- Combine other fields to make room for target address

- ◆ **Optimization:**

- Jumps only jump to word aligned addresses
 - last two bits are always 00 (in binary)
 - specify 28 bits of the 32-bit bit address

Jump Addressing (3/3)

- ◆ Where do we get the other 4 bits?
 - Take the 4 highest order bits from the PC
 - Technically, this means that we cannot jump to anywhere in memory, but it's adequate 99.9999...% of the time, since programs aren't that long
 - Linker and loader avoid placing a program across an address boundary of 256 MB
- ◆ Summary:
 - New PC = PC[31..28] || target address (26 bits) || 00
 - Note: || means concatenation
4 bits || 26 bits || 2 bits = 32-bit address
- ◆ If we absolutely need to specify a 32-bit address:
 - Use *jr \$ra* # jump to the address specified by \$ra

Target Addressing Example

- ◆ Loop code from earlier example
 - Assume Loop at location 80000

Loop: sll	\$t1, \$s3, 2	80000	0	0	19	9	4	0
add	\$t1, \$t1, \$s6	80004	0	9	22	9	0	32
lw	\$t0, 0(\$t1)	80008	35	9	8	0		
bne	\$t0, \$s5, Exit	80012	5	8	21	2		
addi	\$s3, \$s3, 1	80016	8	19	19	1		
j	Loop	80020	2	20000				
Exit: ...		80024						

- ◆ $80016 + 2 \times 4 = 80024$
- ◆ $20000 \times 4 = 80000$


Branching Far Away

- ◆ If branch target is too far to encode with 16-bit offset, assembler rewrites the code
- ◆ Example


```
        beq $s0,$s1, L1
          ↓
        bne $s0,$s1, L2
        j  L1
L2:     ...
```

MIPS Addressing Mode




1. Immediate addressing



2. Register addressing




3. Base addressing




MPIS Addressing Modes

4. PC-relative addressing



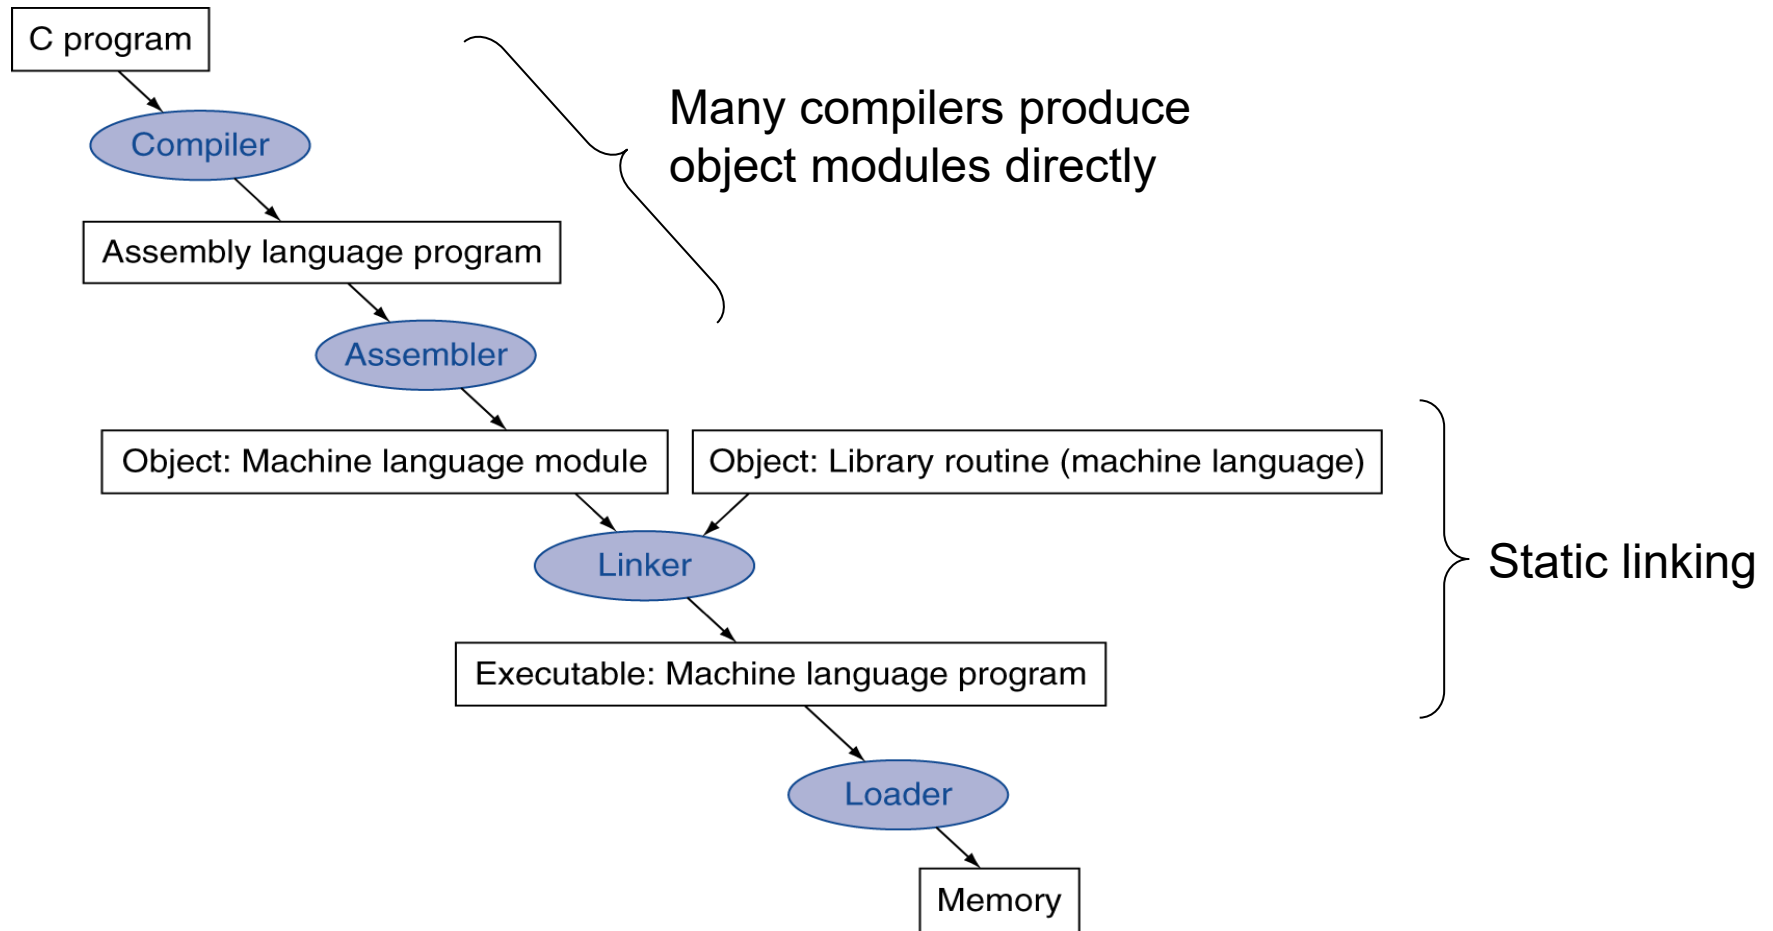
5. Pseudodirect addressing



Outline

- ◆ Instruction set architecture (using MIPS ISA as an example)
- ◆ Operands
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions
- ◆ Operations
 - Logical
 - Decision making and branches
- ◆ Supporting procedures in hardware
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program (Sec. 2.12)
- ◆ A sort example
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets

Translation and Startup



Assembler Pseudoinstructions

- ◆ Most assembler instructions represent machine instructions one-to-one
- ◆ Pseudo instructions: figments of the assembler's imagination

`move $t0, $t1` \rightarrow `add $t0, $zero, $t1`
`blt $t0, $t1, L` \rightarrow `slt $at, $t0, $t1`
 `bne $at, $zero, L`

- `$at` (register 1): assembler temporary

Producing an Object Module

- ◆ Assembler (or compiler) translates program into machine instructions
- ◆ Provides information for building a complete program from the pieces
 - Header: described contents of object module
 - Text segment: translated instructions
 - Static data segment: data allocated for the life of the program
 - Relocation info: for contents that depend on absolute location of loaded program
 - Symbol table: global definitions and external refs
 - Debug info: for associating with source code

Linking Object Modules

- ◆ Produces an executable image
 1. Merges segments
 2. Resolve labels (determine their addresses)
 3. Patch location-dependent and external refs
- ◆ Could leave location dependencies for fixing by a relocating loader
 - But with virtual memory, no need to do this
 - Program can be loaded into absolute location in virtual memory space

Loading a Program

- ◆ Load from image file on disk into memory
 1. Read header to determine segment sizes
 2. Create virtual address space
 3. Copy text and initialized data into memory
 - Or set page table entries so they can be faulted in
 4. Set up arguments on stack
 5. Initialize registers (including \$sp, \$fp, \$gp)
 6. Jump to startup routine
 - Copies arguments to \$a0, ... and calls main
 - When main returns, do exit syscall

Outline

- ◆ Instruction set architecture (using MIPS ISA as an example)
- ◆ Operands
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions
- ◆ Operations
 - Logical
 - Decision making and branches
- ◆ Supporting procedures in hardware
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program
- ◆ A sort example (Sec. 2.13)
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets

C Sort Example

- ◆ Illustrates use of assembly instructions for a C bubble sort function

- ◆ Swap procedure (leaf)

```
void swap(int v[], int k)
{
    int temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}
```

- v in \$a0, k in \$a1, temp in \$t0

The Procedure Swap

swap:	sll \$t1, \$a1, 2	# \$t1 = k * 4
	add \$t1, \$a0, \$t1	# \$t1 = v+(k*4)
		# (address of v[k])
	lw \$t0, 0(\$t1)	# \$t0 (temp) = v[k]
	lw \$t2, 4(\$t1)	# \$t2 = v[k+1]
	sw \$t2, 0(\$t1)	# v[k] = \$t2 (v[k+1])
	sw \$t0, 4(\$t1)	# v[k+1] = \$t0 (temp)
	jr \$ra	# return to calling routine

The Sort Procedure in C

- ◆ Non-leaf (calls swap)

```
void sort (int v[], int n)
{
    int i, j;
    for (i = 0; i < n; i += 1) {
        for (j = i - 1;
             j >= 0 && v[j] > v[j + 1];
             j -= 1) {
            swap(v, j);
        }
    }
}
```

- v in \$a0, k in \$a1, i in \$s0, j in \$s1

The Procedure Body

	move \$s2, \$a0	# save \$a0 into \$s2	Move params
	move \$s3, \$a1	# save \$a1 into \$s3	
	move \$s0, \$zero	# i = 0	
for1tst:	slt \$t0, \$s0, \$s3	# \$t0 = 0 if \$s0 ≥ \$s3 (i ≥ n)	Outer loop
	beq \$t0, \$zero, exit1	# go to exit1 if \$s0 ≥ \$s3 (i ≥ n)	
	addi \$s1, \$s0, -1	# j = i - 1	
for2tst:	slti \$t0, \$s1, 0	# \$t0 = 1 if \$s1 < 0 (j < 0)	
	bne \$t0, \$zero, exit2	# go to exit2 if \$s1 < 0 (j < 0)	
	sll \$t1, \$s1, 2	# \$t1 = j * 4	Inner loop
	add \$t2, \$s2, \$t1	# \$t2 = v + (j * 4)	
	lw \$t3, 0(\$t2)	# \$t3 = v[j]	
	lw \$t4, 4(\$t2)	# \$t4 = v[j + 1]	
	slt \$t0, \$t4, \$t3	# \$t0 = 0 if \$t4 ≥ \$t3	
	beq \$t0, \$zero, exit2	# go to exit2 if \$t4 ≥ \$t3	
	move \$a0, \$s2	# 1st param of swap is v (old \$a0)	Pass params & call
	move \$a1, \$s1	# 2nd param of swap is j	
	jal swap	# call swap procedure	
	addi \$s1, \$s1, -1	# j -= 1	
	j for2tst	# jump to test of inner loop	Inner loop
exit2:	addi \$s0, \$s0, 1	# i += 1	
	j for1tst	# jump to test of outer loop	Outer loop





Architecture

The Full Procedure




sort:	addi \$sp,\$sp, -20	# make room on stack for 5 registers
	sw \$ra, 16(\$sp)	# save \$ra on stack
	sw \$s3,12(\$sp)	# save \$s3 on stack
	sw \$s2, 8(\$sp)	# save \$s2 on stack
	sw \$s1, 4(\$sp)	# save \$s1 on stack
	sw \$s0, 0(\$sp)	# save \$s0 on stack
	...	# procedure body
	...	
exit1:	lw \$s0, 0(\$sp)	# restore \$s0 from stack
	lw \$s1, 4(\$sp)	# restore \$s1 from stack
	lw \$s2, 8(\$sp)	# restore \$s2 from stack
	lw \$s3,12(\$sp)	# restore \$s3 from stack
	lw \$ra,16(\$sp)	# restore \$ra from stack
	addi \$sp,\$sp, 20	# restore stack pointer
	jr \$ra	# return to calling routine

Effect of Compiler Optimization

Compiled with gcc for Pentium 4 under Linux



Effect of Language and Algorithm



Lessons Learnt

- ◆ **Instruction count and CPI are not good performance indicators in isolation**
- ◆ **Compiler optimizations are sensitive to the algorithm**
- ◆ **Nothing can fix a dumb algorithm!**

Outline

- ◆ Instruction set architecture (using MIPS ISA as an example)
- ◆ Operands
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions
- ◆ Operations
 - Logical
 - Decision making and branches
- ◆ Supporting procedures in hardware
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program (Sec. 2.12)
- ◆ A sort example
- ◆ Arrays versus pointers (Sec. 2.14)
- ◆ ARM and x86 instruction sets

Arrays vs. Pointers

- ◆ Array indexing involves
 - Multiplying index by element size
 - Adding to array base address
- ◆ Pointers correspond directly to memory addresses
 - Can avoid indexing complexity

Example: Clearing an Array

```
clear1(int array[], int size) {  
    int i;  
    for (i = 0; i < size; i += 1)  
        array[i] = 0;  
}
```

```
        move $t0,$zero    # i = 0  
loop1: sll $t1,$t0,2      # $t1 = i * 4  
        add $t2,$a0,$t1  # $t2 =  
                        # &array[i]  
        sw $zero, 0($t2) # array[i] = 0  
        addi $t0,$t0,1   # i = i + 1  
        slt $t3,$t0,$a1  # $t3 =  
                        # (i < size)  
        bne $t3,$zero,loop1 # if (...)  
                        # goto loop1
```

```
clear2(int *array, int size) {  
    int *p;  
    for (p = &array[0]; p < &array[size];  
        p = p + 1)  
        *p = 0;  
}
```

```
        move $t0, $a0    # p = & array[0]  
        sll $t1, $a1, 2  # $t1 = size * 4  
        add $t2,$a0,$t1  # $t2 =  
                        # &array[size]  
loop2: sw $zero,0($t0) # Memory[p] = 0  
        addi $t0,$t0,4   # p = p + 4  
        slt $t3,$t0,$t2  # $t3 =  
                        # (p < &array[size])  
        bne $t3,$zero,loop2 # if (...)  
                        # goto loop2
```

Comparison of Array vs. Ptr

- ◆ Multiply “strength reduced” to shift (strength reduction)
- ◆ Array version requires shift to be inside loop
 - Part of index calculation for incremented i
 - c.f. incrementing pointer
- ◆ Compiler can achieve same effect as manual use of pointers
 - Eliminating array address calculations within loop (induction variable elimination): 6 instructions reduced to 4 in loop
 - Better to make program clearer and safer

Outline

- ◆ Instruction set architecture (using MIPS ISA as an example)
- ◆ Operands
 - Register operands and their organization
 - Memory operands, data transfer
 - Immediate operands
- ◆ Signed and unsigned numbers
- ◆ Representing instructions
- ◆ Operations
 - Logical
 - Decision making and branches
- ◆ Supporting procedures in hardware
- ◆ Communicating with people
- ◆ Addressing for 32-bit immediate and addresses
- ◆ Translating and starting a program
- ◆ A sort example
- ◆ Arrays versus pointers
- ◆ ARM and x86 instruction sets (Sec. 2.15, 2.16)

ARM & MIPS Similarities

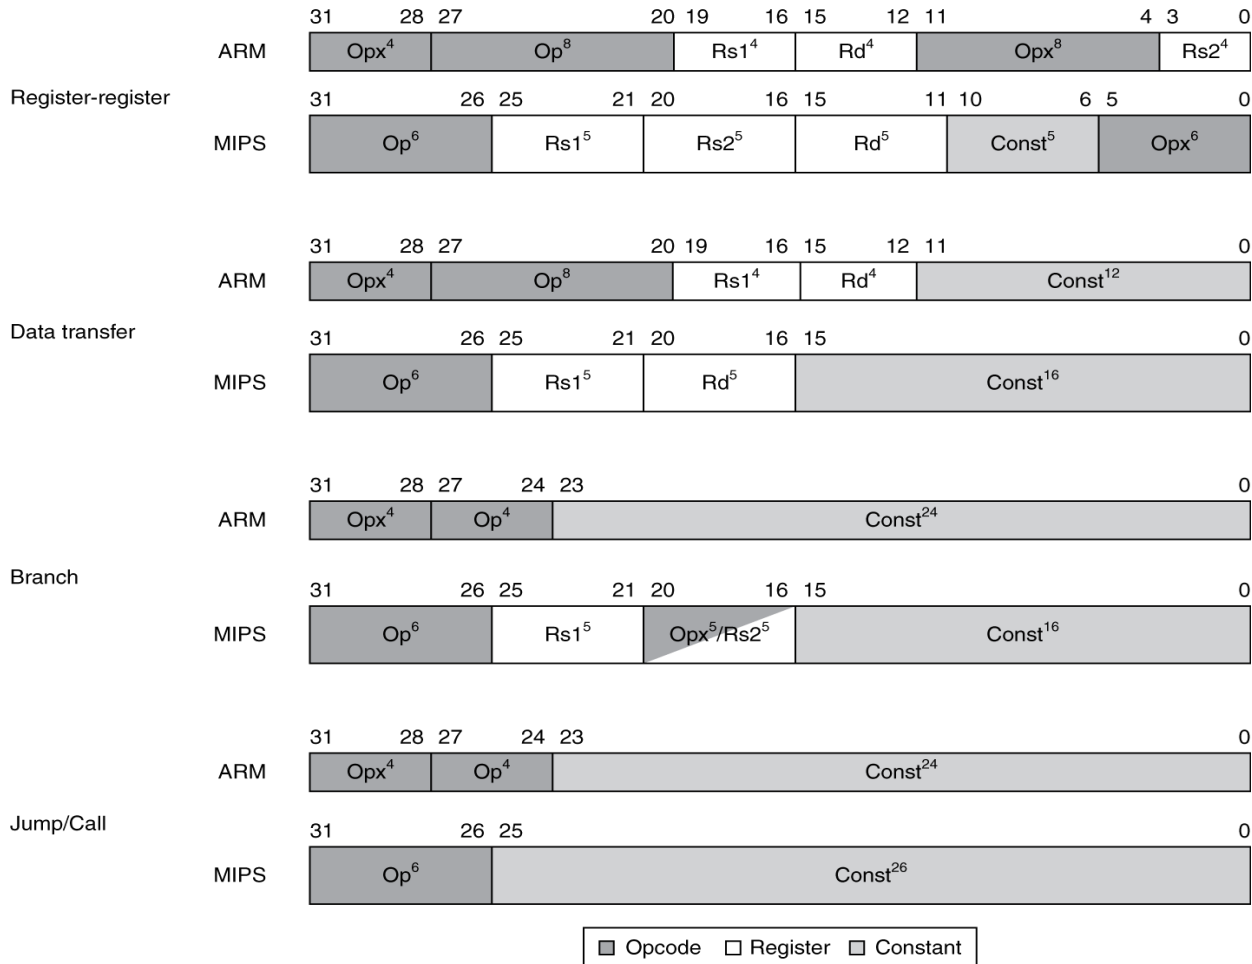
- ◆ ARM: the most popular embedded core
- ◆ Similar basic set of instructions to MIPS

	ARM	MIPS
Date announced	1985	1985
Instruction size	32 bits	32 bits
Address space	32-bit flat	32-bit flat
Data alignment	Aligned	Aligned
Data addressing modes	9	3
Registers	15 × 32-bit	31 × 32-bit
Input/output	Memory mapped	Memory mapped

Compare and Branch in ARM

- ◆ Uses condition codes for result of an arithmetic/logical instruction
 - Negative, zero, carry, overflow
 - Compare instructions to set condition codes without keeping the result
- ◆ Each instruction can be conditional
 - Top 4 bits of instruction word: condition value
 - Can avoid branches over single instructions

Instruction Encoding



The Intel x86 ISA

- ◆ Evolution with backward compatibility
 - 8080 (1974): 8-bit microprocessor
 - Accumulator, plus 3 index-register pairs
 - 8086 (1978): 16-bit extension to 8080
 - Complex instruction set (CISC)
 - 8087 (1980): floating-point coprocessor
 - Adds FP instructions and register stack
 - 80286 (1982): 24-bit addresses, MMU
 - Segmented memory mapping and protection
 - 80386 (1985): 32-bit extension (now IA-32)
 - Additional addressing modes and operations
 - Paged memory mapping as well as segments

The Intel x86 ISA

◆ Further evolution...

- i486 (1989): pipelined, on-chip caches and FPU
 - Compatible competitors: AMD, Cyrix, ...
- Pentium (1993): superscalar, 64-bit datapath
 - Later versions added MMX (Multi-Media eXtension) instructions
 - The infamous FDIV bug
- Pentium Pro (1995), Pentium II (1997)
 - New microarchitecture (see Colwell, *The Pentium Chronicles*)
- Pentium III (1999)
 - Added SSE (Streaming SIMD Extensions) and associated registers
- Pentium 4 (2001)
 - New microarchitecture
 - Added SSE2 instructions

The Intel x86 ISA

◆ And further...

- **AMD64 (2003): extended architecture to 64 bits**
 - **EM64T – Extended Memory 64 Technology (2004)**
 - AMD64 adopted by Intel (with refinements)
 - Added SSE3 instructions
 - **Intel Core (2006)**
 - Added SSE4 instructions, virtual machine support
 - **AMD64 (announced 2007): SSE5 instructions**
 - Intel declined to follow, instead...
 - **Advanced Vector Extension (announced 2008)**
 - Longer SSE registers, more instructions
- ◆ If Intel didn't extend with compatibility, its competitors would!
- **Technical elegance ≠ market success**

Basic x86 Registers

Name	31	0	Use
EAX			GPR 0
ECX			GPR 1
EDX			GPR 2
EBX			GPR 3
ESP			GPR 4
EBP			GPR 5
ESI			GPR 6
EDI			GPR 7
	CS		Code segment pointer
	SS		Stack segment pointer (top of stack)
	DS		Data segment pointer 0
	ES		Data segment pointer 1
	FS		Data segment pointer 2
	GS		Data segment pointer 3
EIP			Instruction pointer (PC)
EFLAGS			Condition codes

Basic x86 Addressing Modes

◆ Two operands per instruction


Source/dest operand	Second source operand
Register	Register
Register	Immediate
Register	Memory
Memory	Register
Memory	Immediate

◆ Memory addressing modes

- Address in register
- Address = $R_{\text{base}} + \text{displacement}$
- Address = $R_{\text{base}} + 2^{\text{scale}} \times R_{\text{index}}$ (scale = 0, 1, 2, or 3)
- Address = $R_{\text{base}} + 2^{\text{scale}} \times R_{\text{index}} + \text{displacement}$

x86 Instruction Encoding


a. JE EIP + displacement




b. CALL




c. MOV EBX, [EDI + 45]



d. PUSH ESI



e. ADD EAX, #6765



f. TEST EDX, #42



◆ Variable length encoding

- Postfix bytes specify addressing mode
- Prefix bytes modify operation
 - Operand length, repetition, locking, ...

Implementing IA-32


- ◆ **Complex instruction set makes implementation difficult**
 - **Hardware translates instructions to simpler microoperations**
 - **Simple instructions: 1–1**
 - **Complex instructions: 1–many**
 - **Microengine similar to RISC**
 - **Market share makes this economically viable**
- ◆ **Comparable performance to RISC**
 - **Compilers avoid complex instructions**

Fallacies

- ◆ **Powerful instruction \Rightarrow higher performance**
 - Fewer instructions required
 - But complex instructions are hard to implement
 - May slow down all instructions, including simple ones
 - Compilers are good at making fast code from simple instructions
- ◆ **Use assembly code for high performance**
 - But modern compilers are better at dealing with modern processors
 - More lines of code \Rightarrow more errors and less productivity

Fallacies

- ◆ Backward compatibility \Rightarrow instruction set doesn't change
 - But they do accrete more instructions



x86 instruction set

Pitfalls

- ◆ Sequential words are not at sequential addresses
 - Increment by 4, not by 1!
- ◆ Keeping a pointer to an automatic variable after procedure returns
 - e.g., passing pointer back via an argument
 - Pointer becomes invalid when stack popped

Concluding Remarks

- ◆ Design principles
 1. Simplicity favors regularity
 2. Smaller is faster
 3. Make the common case fast
 4. Good design demands good compromises
- ◆ Layers of software/hardware
 - Compiler, assembler, hardware
- ◆ MIPS: typical of RISC ISAs
 - c.f. x86

Concluding Remarks

- ◆ Measure MIPS instruction executions in benchmark programs
 - Consider making the common case fast
 - Consider compromises

Instruction class	MIPS examples	SPEC2006 Int	SPEC2006 FP
Arithmetic	add, sub, addi	16%	48%
Data transfer	lw, sw, lb, lbu, lh, lhu, sb, lui	35%	36%
Logical	and, or, nor, andi, ori, sll, srl	12%	4%
Cond. Branch	beq, bne, slt, slti, sltiu	34%	8%
Jump	j, jr, jal	2%	0%