

# CS4100: 計算機結構

## Computer Arithmetic

---

國立清華大學資訊工程學系  
一零零學年度第二學期


# Outline

- ◆ Addition and subtraction (Sec. 3.2)
- ◆ Constructing an arithmetic logic unit (Appendix C)
- ◆ Multiplication (Sec. 3.3, Appendix C)
- ◆ Division (Sec. 3.4)
- ◆ Floating point (Sec. 3.5)

# Problem: Designing MIPS ALU

- ◆ Requirements: must support the following arithmetic and logic operations
  - **add, sub**: two's complement adder/subtractor with overflow detection
  - **and, or, nor** : logical AND, logical OR, logical NOR
  - **slt** (set on less than): two's complement adder with inverter, check sign bit of result

# Functional Specification



## ALU Control (ALUop)

0000

0001

0010

0110

0111

1100

## Function

and

or

add

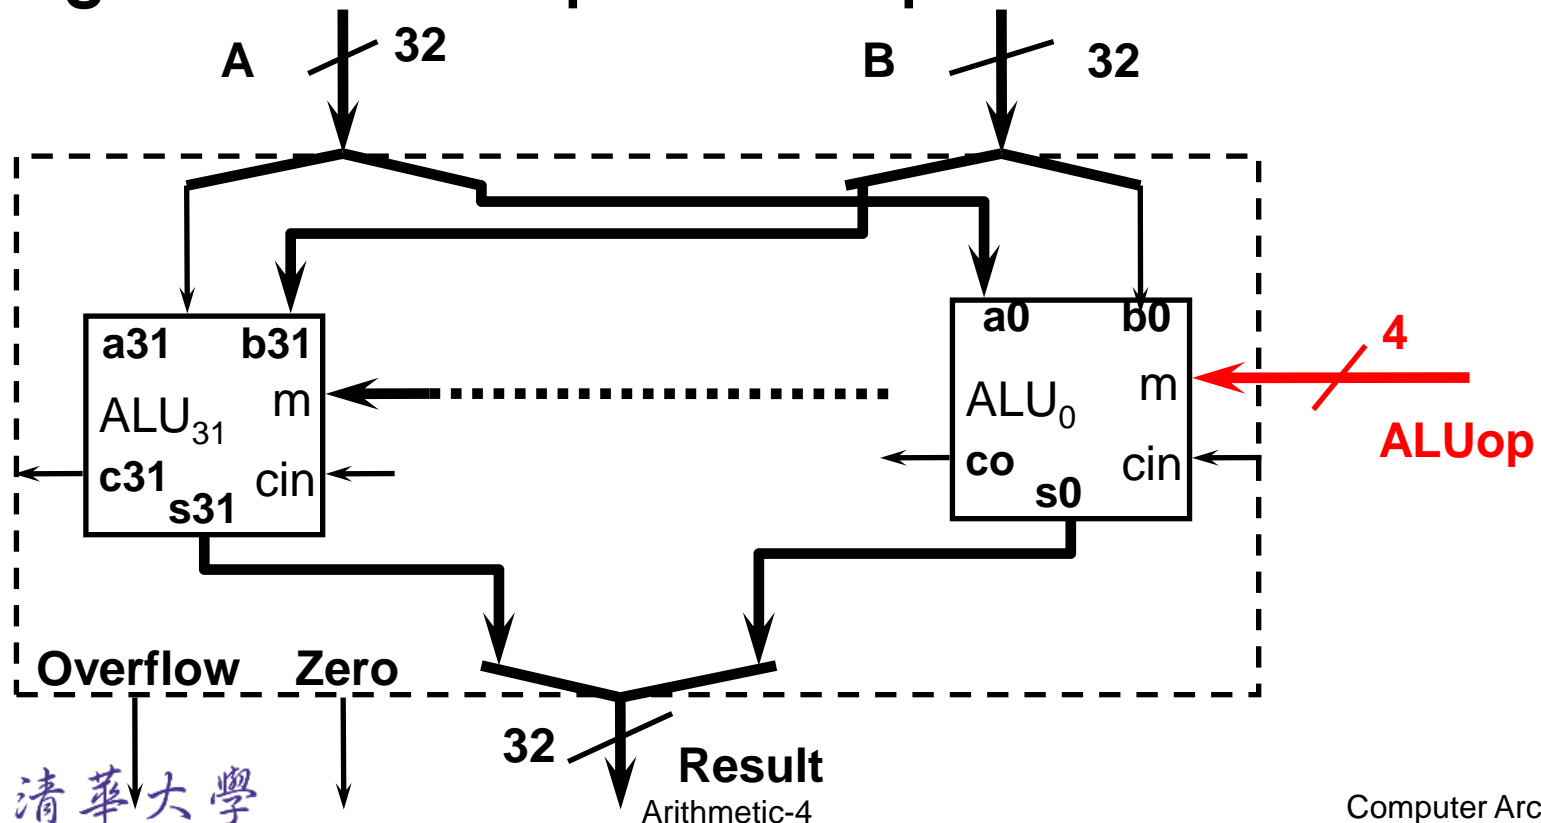
subtract

set-on-less-than

nor

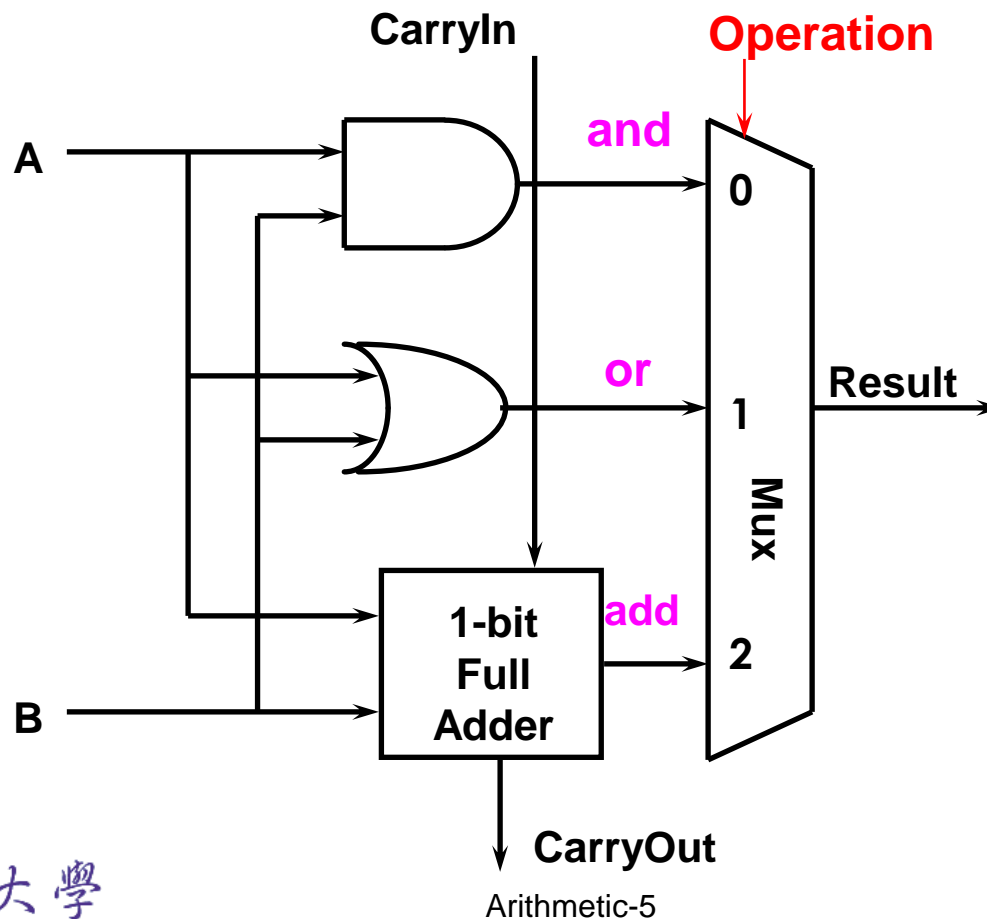
# A Bit-slice ALU

- ◆ Design trick 1: divide and conquer
  - Break the problem into simpler problems, solve them and glue together the solution
- ◆ Design trick 2: solve part of the problem and extend




# A 1-bit ALU

- ◆ Design trick 3: take pieces you know (or can imagine) and try to put them together

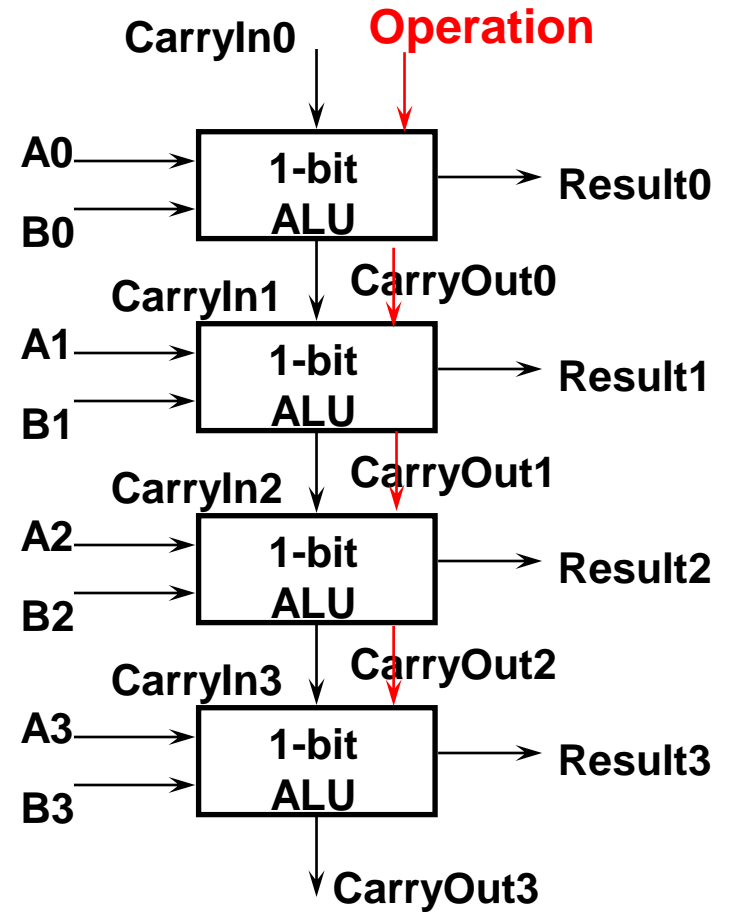


# A 4-bit ALU

## 1-bit ALU




## 4-bit ALU



# How about Subtraction?

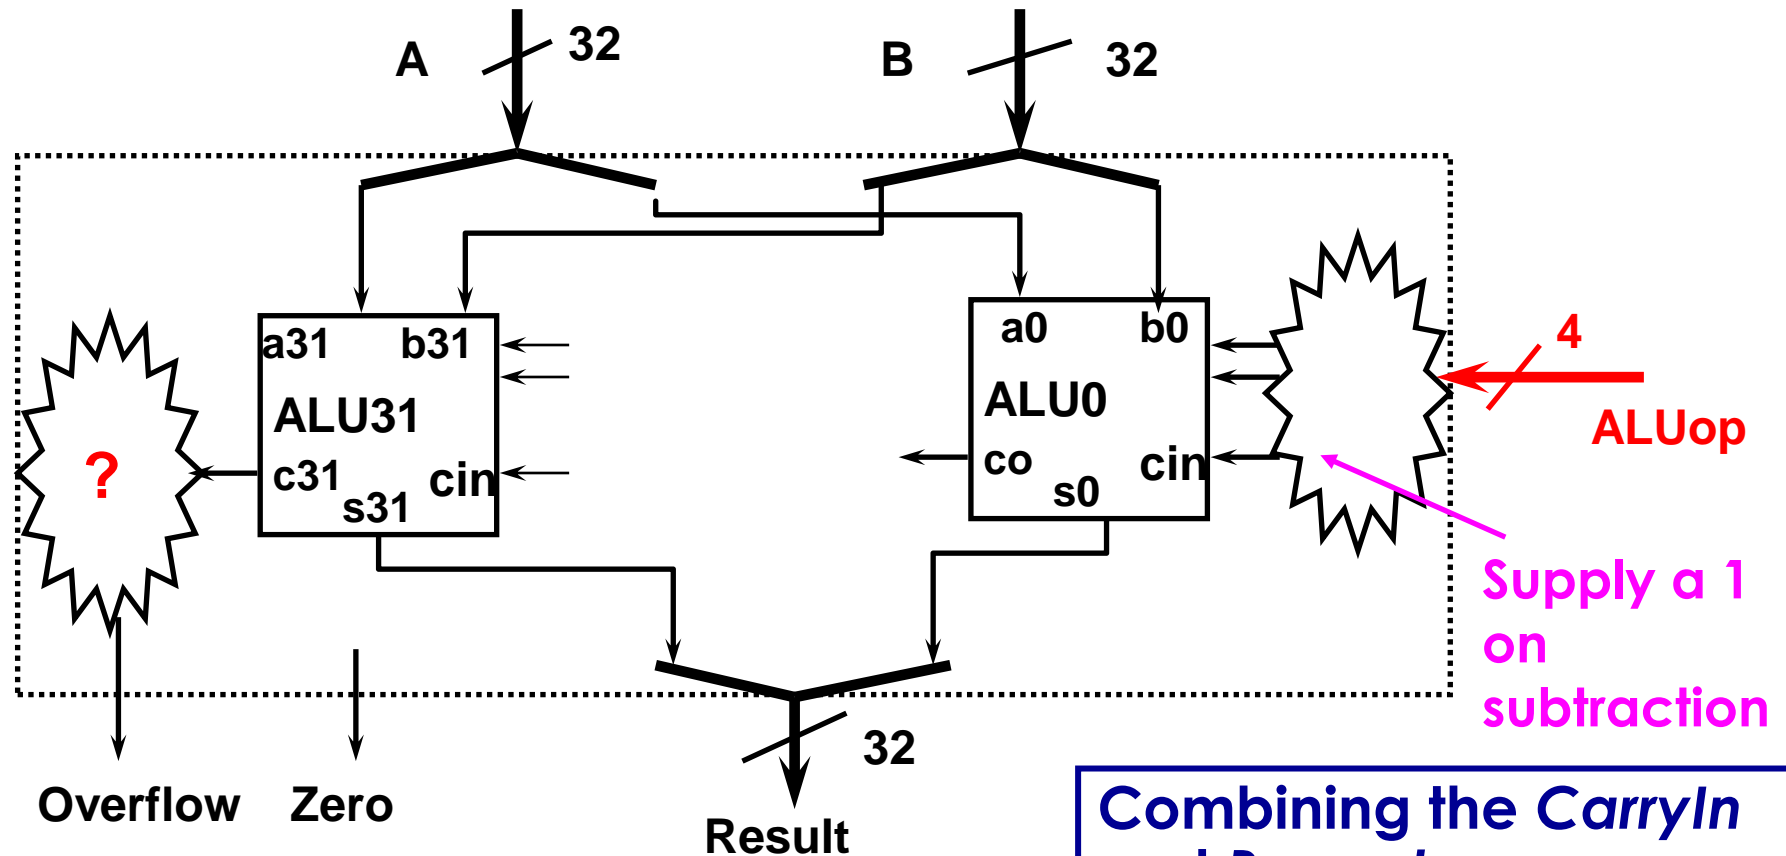
- ◆ 2's complement: take inverse of every bit and add 1 (at  $c_{in}$  of first stage)
  - $A + B' + 1 = A + (B' + 1) = A + (-B) = A - B$
  - Bit-wise inverse of B is  $B'$






# Revised Diagram

- ◆ LSB and MSB need to do a little extra



# Functional Specification



## ALU Control (ALUop)

0000

0001

0010

0110

0111

1100

## Function

and

or

add

subtract

set-on-less-than

nor

# R-Format Instructions (1/2)

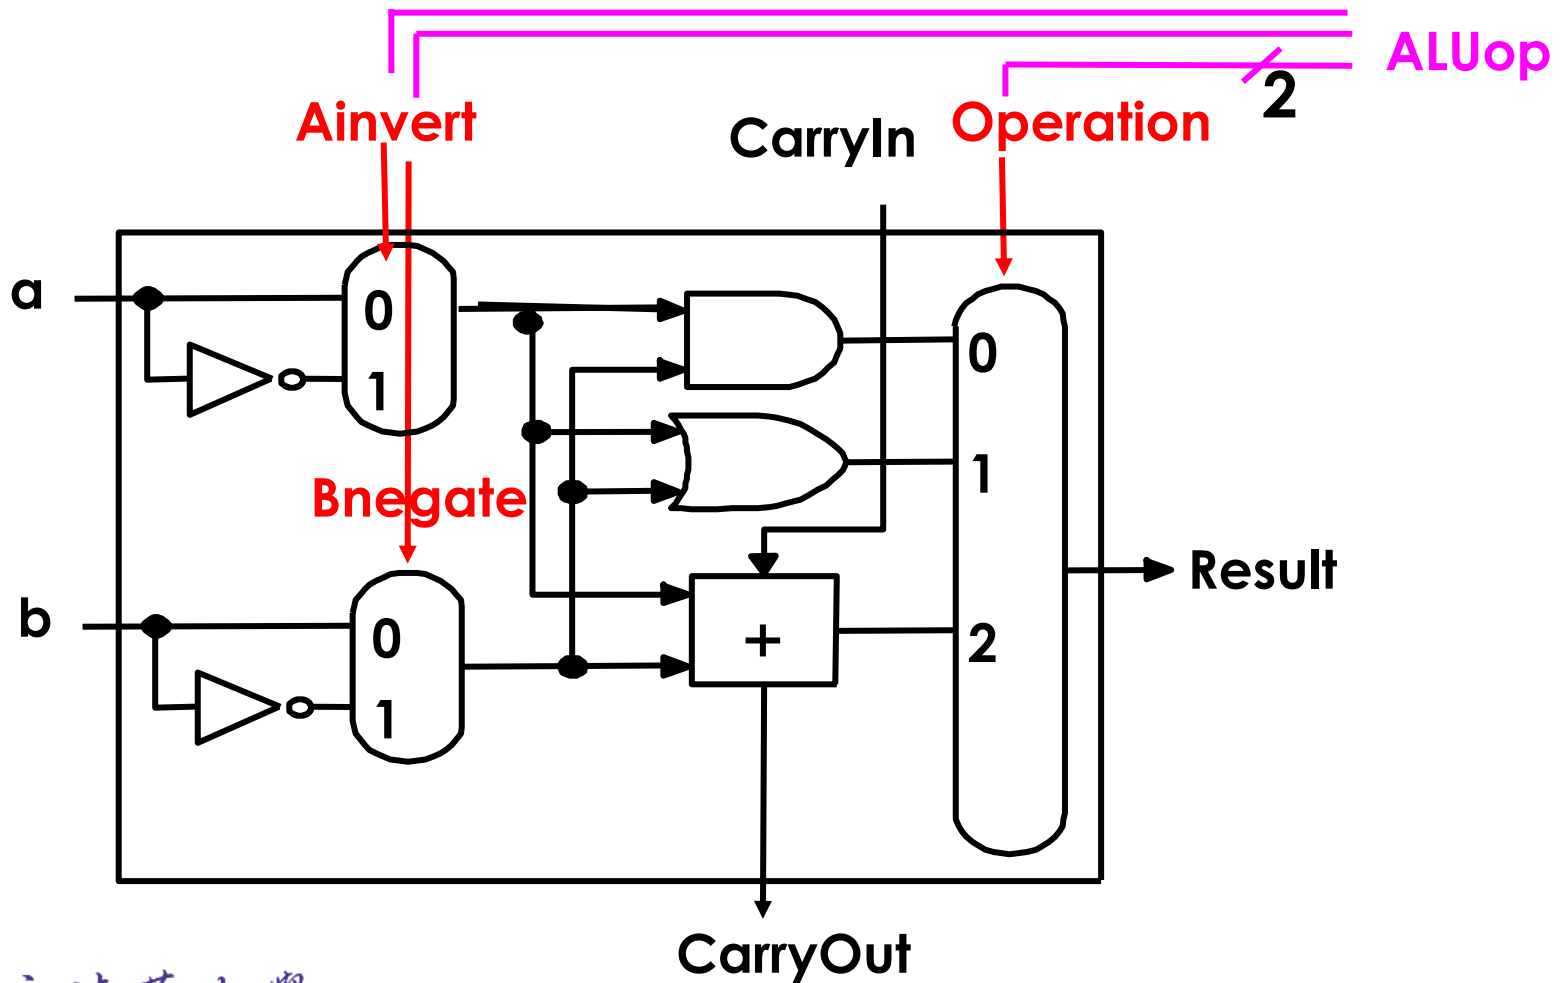
- ◆ Define the following “fields”:

6	5	5	5	5	6
opcode	rs	rt	rd	shamt	funct

- opcode: partially specifies what instruction it is (Note: 0 for all R-Format instructions)
- funct: combined with opcode to specify the instruction  
Question: Why aren't opcode and funct a single 12-bit field?
- rs (Source Register): *generally* used to specify register containing first operand
- rt (Target Register): *generally* used to specify register containing second operand
- rd (Destination Register): *generally* used to specify register which will receive result of computation

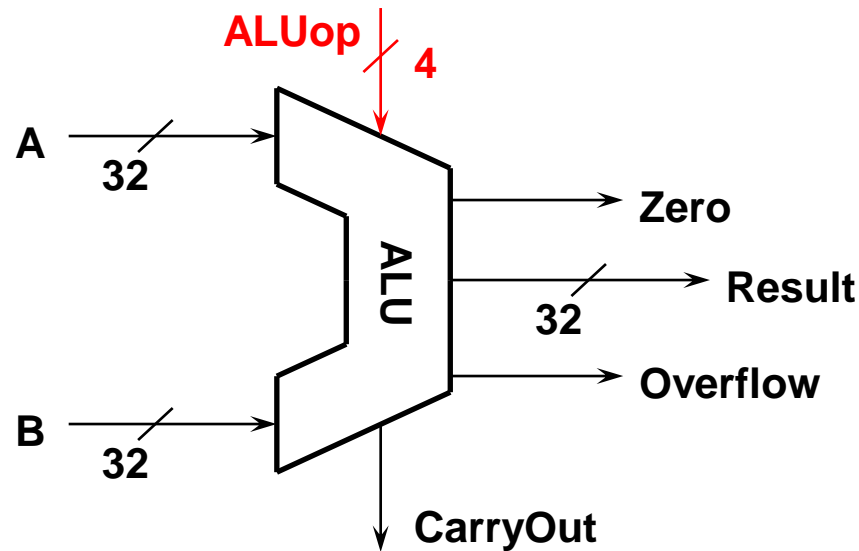
# Nor Operation

- ◆  $A \text{ nor } B = (\text{not } A) \text{ and } (\text{not } B)$



Arithmetic-11

# Functional Specification



## ALU Control (ALUop)

0000

0001

0010

0110

0111

1100

## Function

and

or


add

subtract

set-on-less-than

nor

# Functional Specification



## ALU Control (ALUop)

0000

0001

0010

0110

0111

1100

## Function

and

or

add

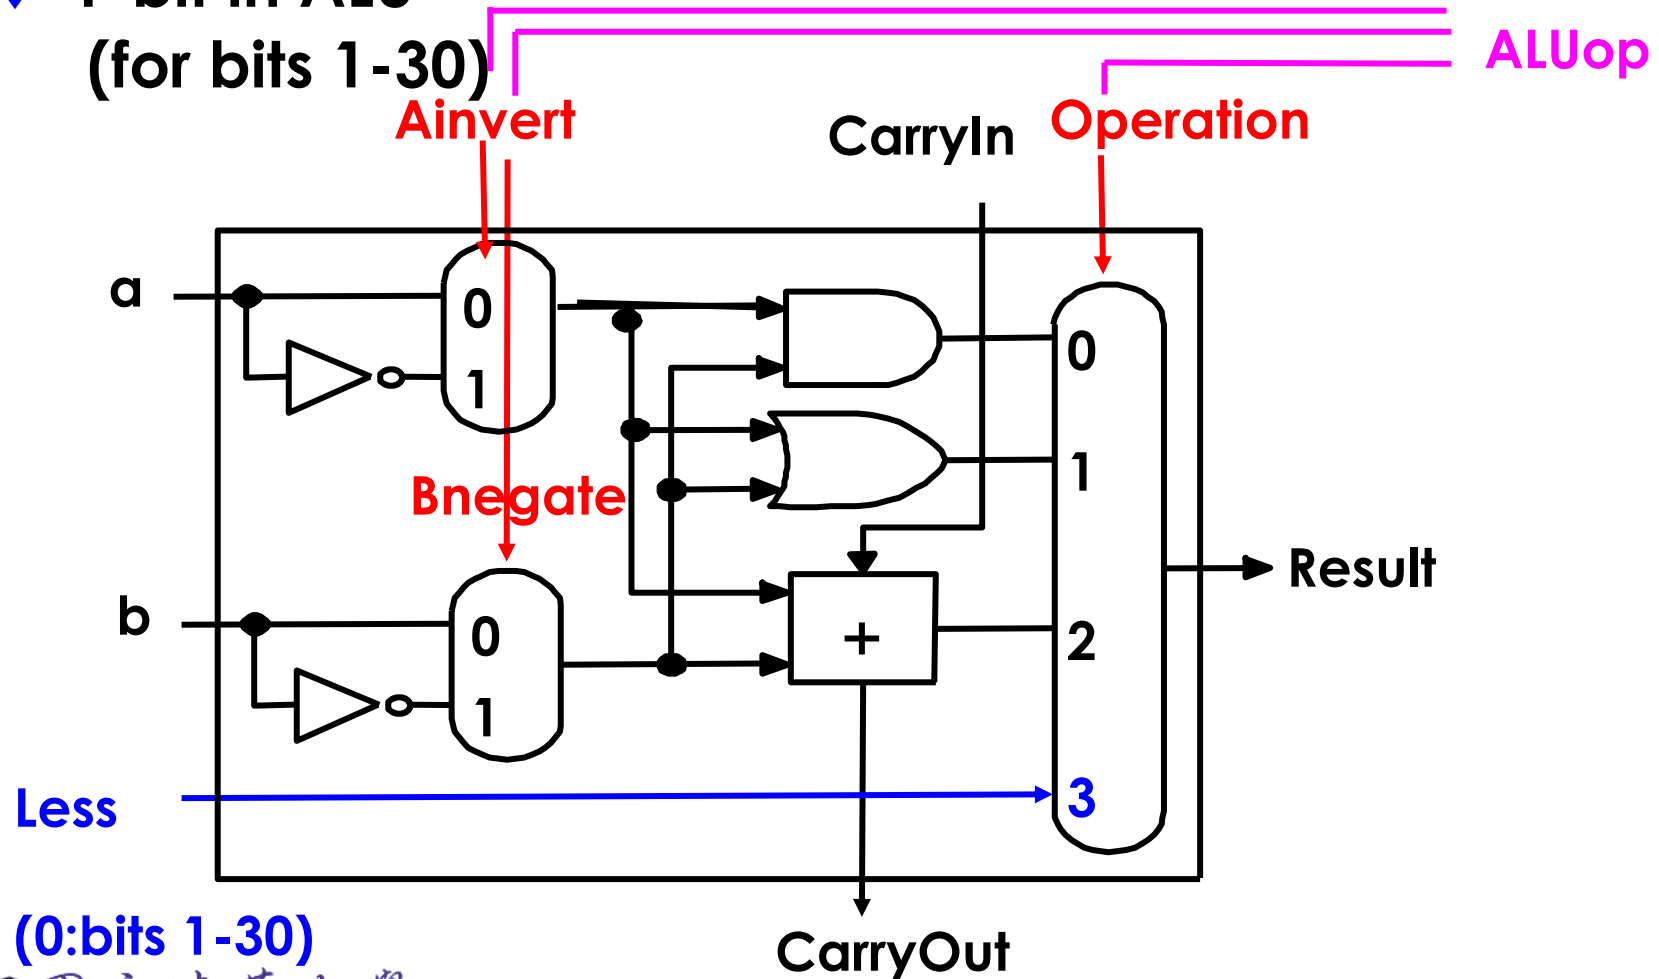
subtract

set-on-less-than

nor

# Set on Less Than (I)

- ◆ **1-bit in ALU**  
**(for bits 1-30)**

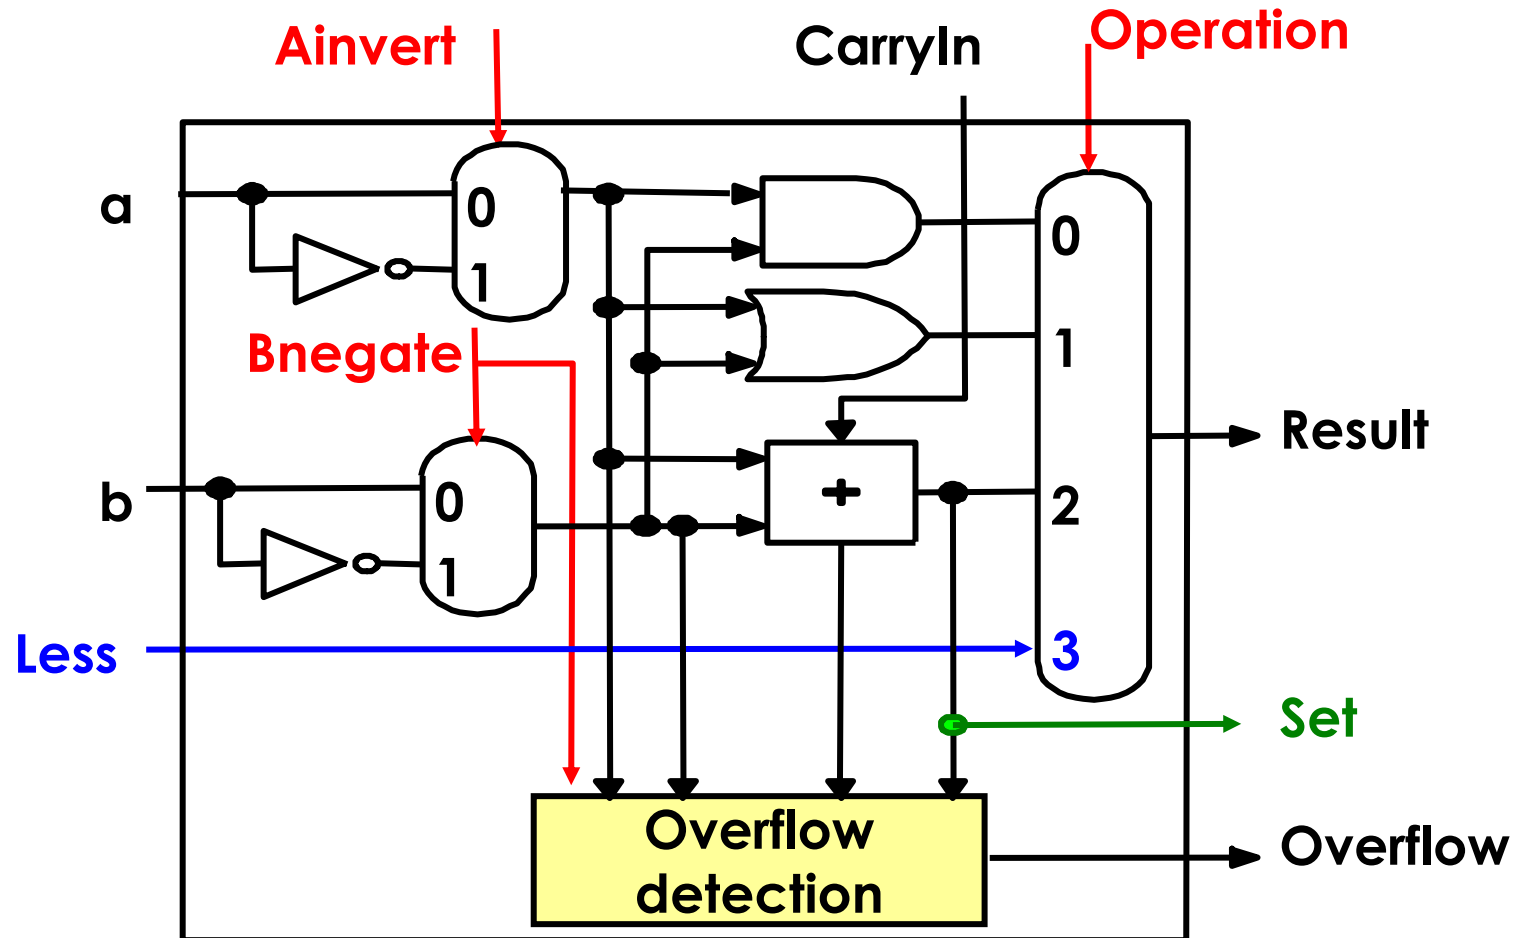


# CarryOut

Arithmetic-14

# Set on Less Than (II)

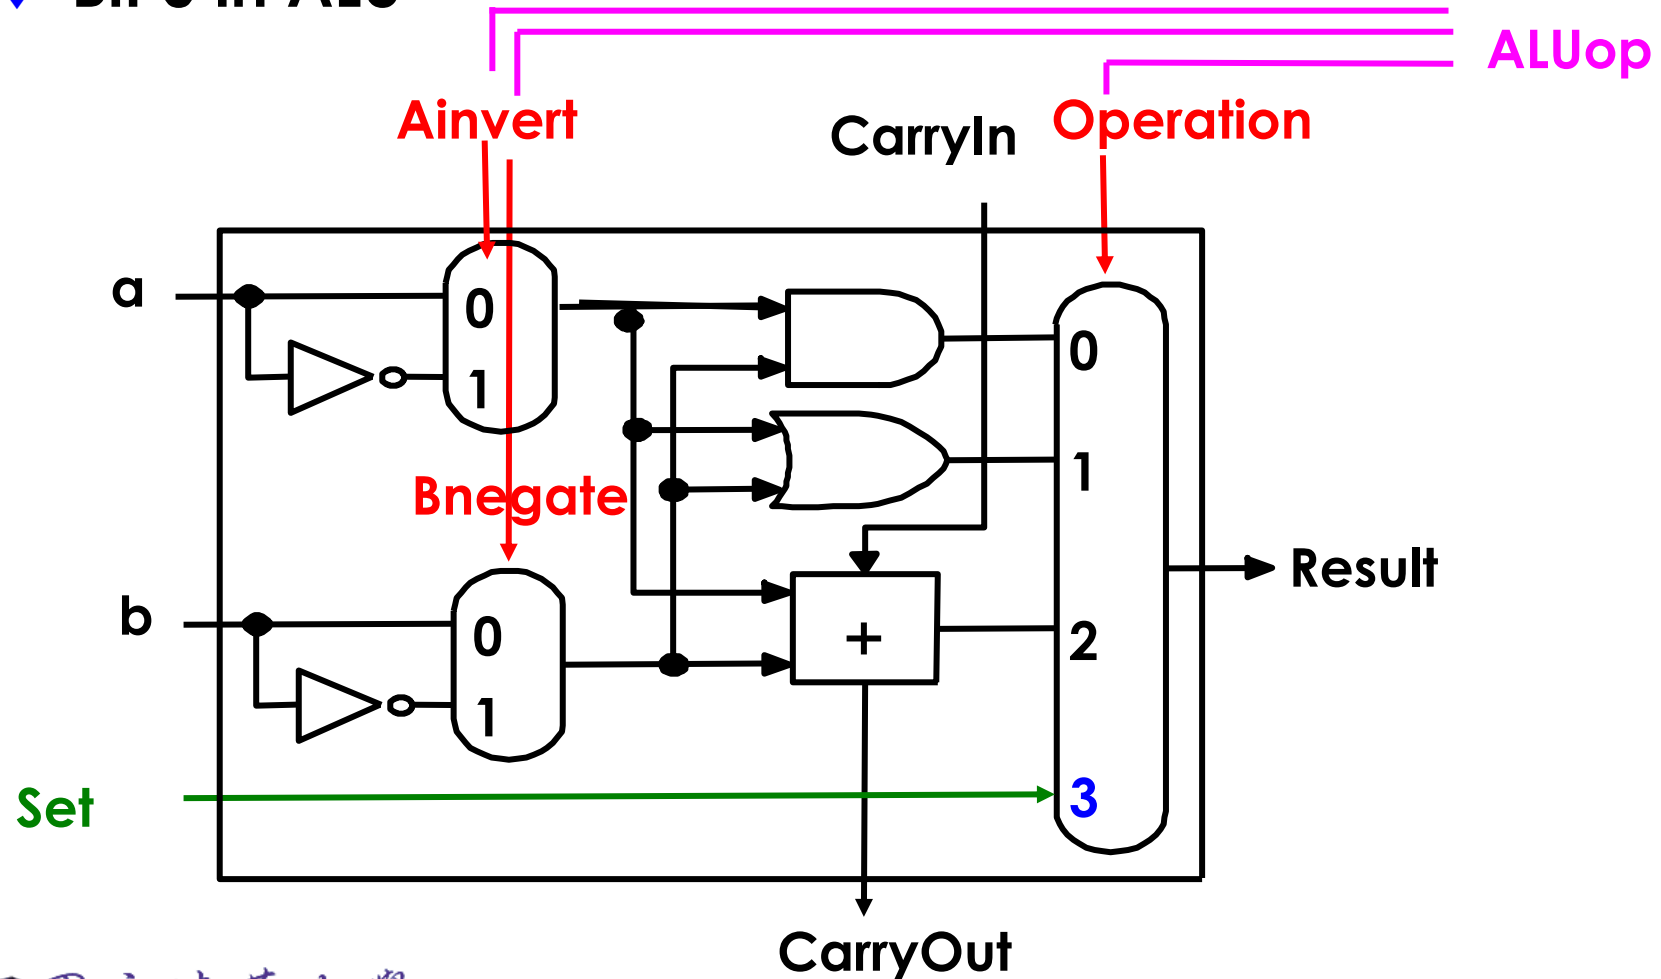
## ◆ Sign bit in ALU



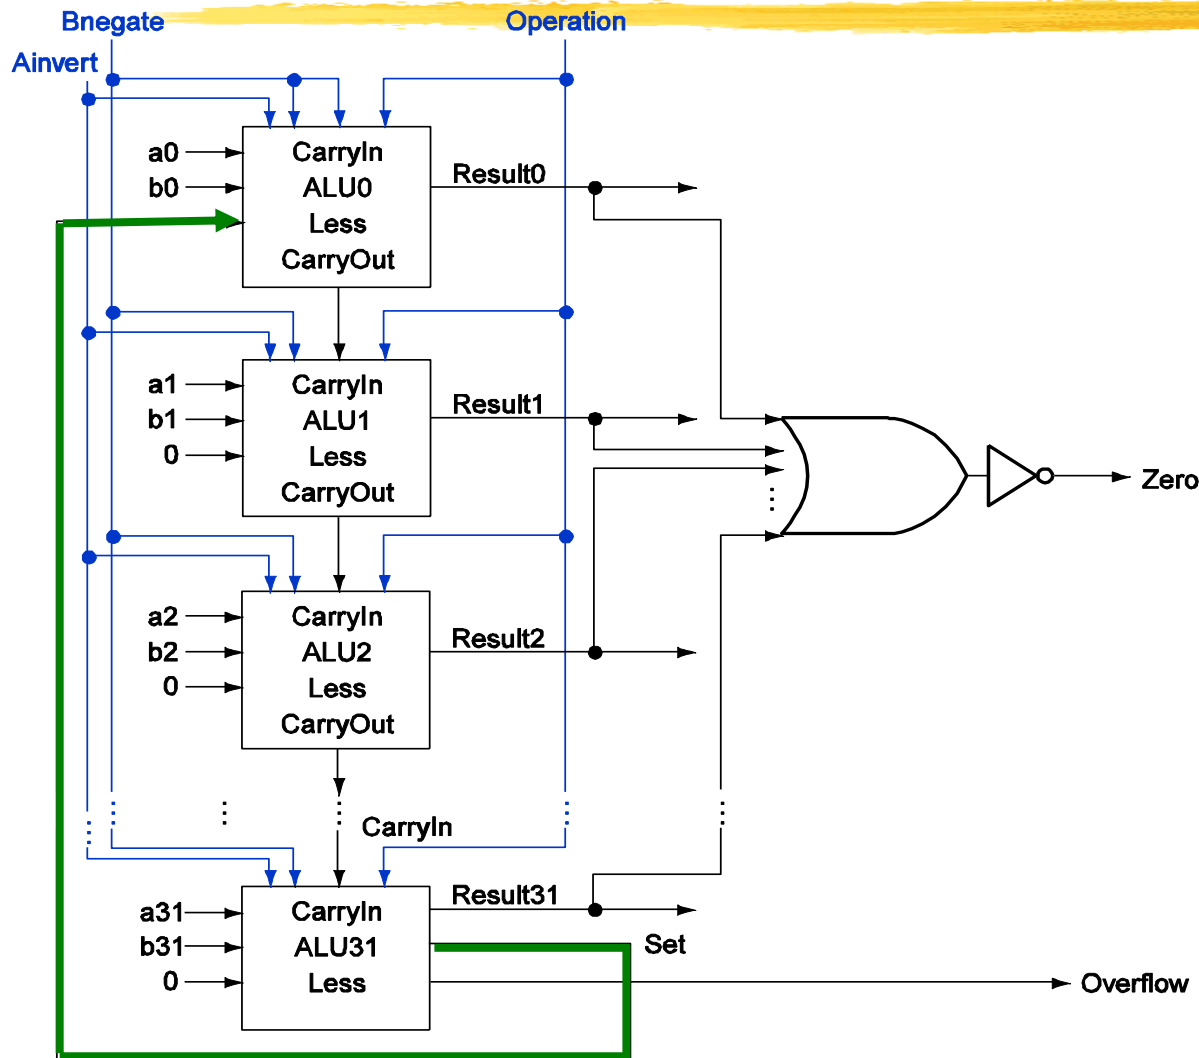


# Set on Less Than (III)

## ◆ Bit 0 in ALU



# A Ripple Carry Adder and Set on Less Than



ALUop	Function
0000	and
0001	or
0010	add
0110	subtract
0111	set-less-than
1100	nor

# Overflow

Decimal	Binary	Decimal	2's complement
0	0000	0	0000
1	0001	-1	1111
2	0010	-2	1110
3	0011	-3	1101
4	0100	-4	1100
5	0101	-5	1011
6	0110	-6	1010
7	0111	-7	1001
		-8	1000

Ex:  $7 + 3 = 10$  but ...

$$\begin{array}{r}
 \begin{array}{ccccccc}
 0 & 1 & 1 & 1 & & & \\
 & \swarrow & \swarrow & \swarrow & \swarrow & & \\
 & 0 & 1 & 1 & 1 & & 7 \\
 + & 0 & 0 & 1 & 1 & & 3 \\
 \hline
 & 1 & 0 & 1 & 0 & & -6
 \end{array}
 \end{array}$$

$-4 - 5 = -9$  but ...

$$\begin{array}{r}
 \begin{array}{ccccccc}
 1 & 0 & 0 & 0 & & & \\
 & \swarrow & \swarrow & \swarrow & \swarrow & & \\
 & 1 & 1 & 0 & 0 & & -4 \\
 + & 1 & 0 & 1 & 1 & & -5 \\
 \hline
 & 0 & 1 & 1 & 1 & & 7
 \end{array}
 \end{array}$$

# Overflow Detection

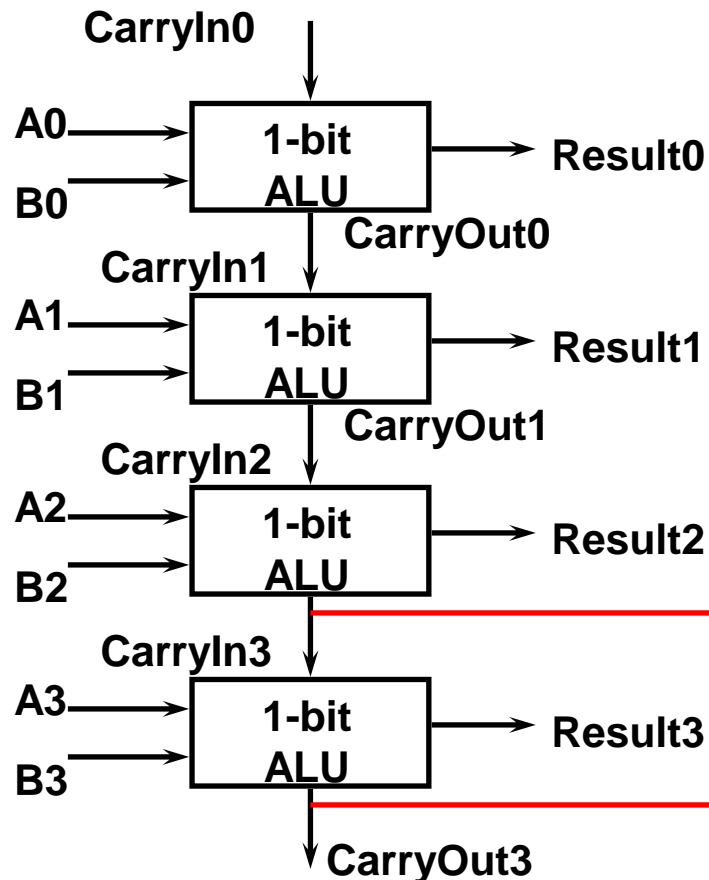
- ◆ **Overflow: result too big/small to represent**
  - $-8 \leq 4\text{-bit binary number} \leq 7$
  - When adding operands with different signs, overflow cannot occur!
  - Overflow occurs when adding:
    - 2 positive numbers and the sum is negative
    - 2 negative numbers and the sum is positive $\Rightarrow$  sign bit is set with the value of the result
  - Overflow if: Carry into MSB  $\neq$  Carry out of MSB

	0	1	1	1	
	0	1	1	1	7
+	0	0	1	1	3
<hr/>					
	1	0	1	0	-6


	1	0	0	0	
	1	1	0	0	-4
+	1	0	1	1	-5
<hr/>					
	0	1	1	1	7

# Overflow Detection Logic

- ◆ **Overflow = CarryIn[N-1] XOR CarryOut[N-1]**



X	Y	X XOR Y
0	0	0
0	1	1
1	0	1
1	1	0

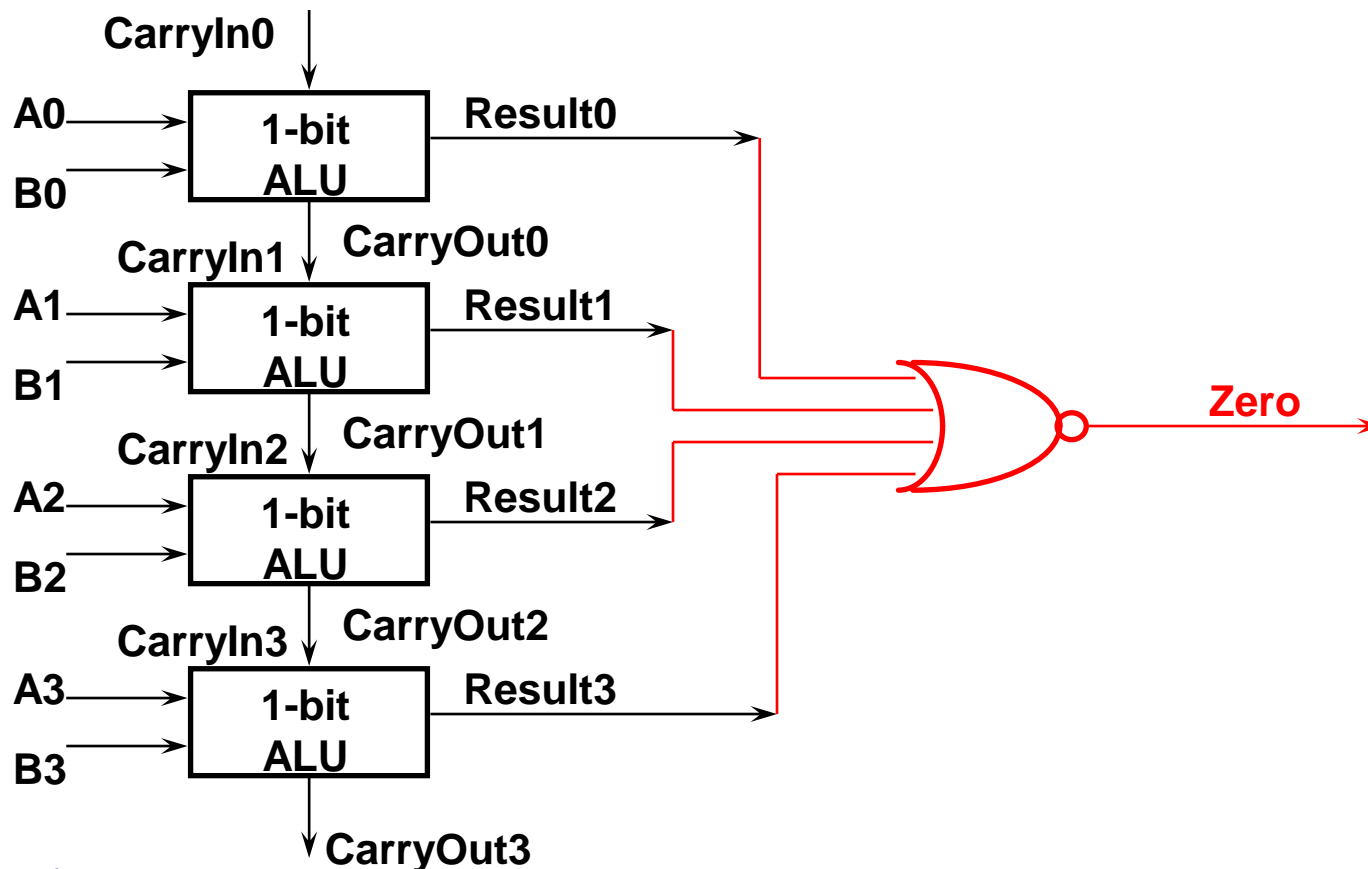


# Dealing with Overflow

- ◆ Some languages (e.g., C) ignore overflow
  - Use MIPS `addu`, `addui`, `subu` instructions
- ◆ Other languages (e.g., Ada, Fortran) require raising an exception
  - Use MIPS `add`, `addi`, `sub` instructions
  - On overflow, invoke exception handler
    - Save PC in exception program counter (EPC) register
    - Jump to predefined handler address
    - `mfc0` (move from coprocessor reg) instruction can retrieve (copy) EPC value (to a general purpose register), to return after corrective action (by jump register instruction)


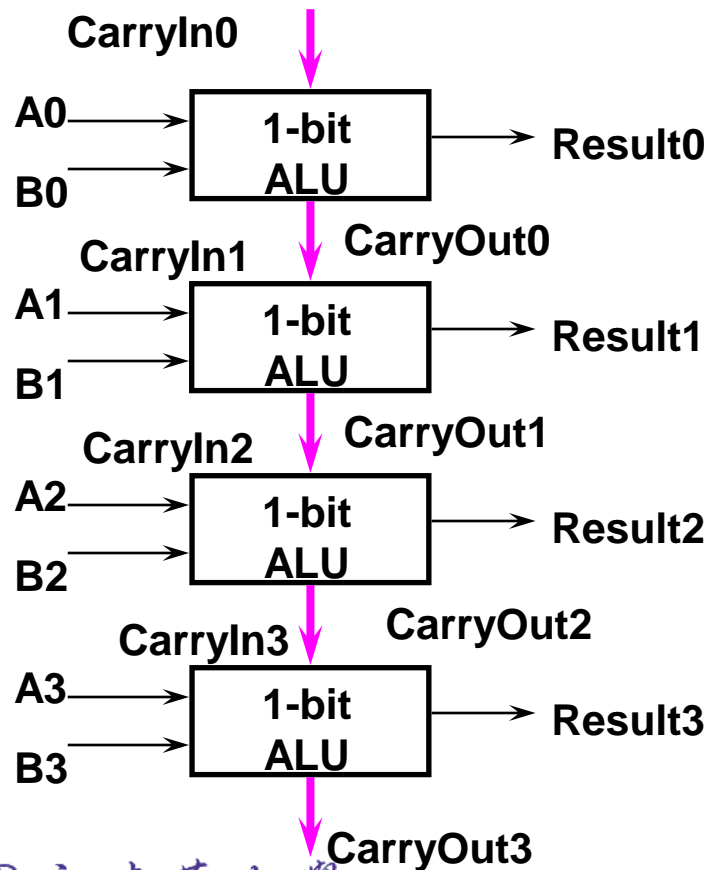
# Zero Detection Logic

- ◆ Zero Detection Logic is a one BIG NOR gate (support conditional jump)



# Problems with Ripple Carry Adder

- ◆ Carry bit may have to propagate from LSB to MSB => worst case delay: N-stage delay




**Design Trick: look for parallelism and throw hardware at it**



# Carry Lookahead: Theory (I)

## (Appendix C)



- ◆  $\text{CarryOut} = (B * \text{CarryIn}) + (A * \text{CarryIn}) + (A * B)$ 
  - $\text{Cin2} = \text{Cout1} = (B1 * \text{Cin1}) + (A1 * \text{Cin1}) + (A1 * B1)$
  - $\text{Cin1} = \text{Cout0} = (B0 * \text{Cin0}) + (A0 * \text{Cin0}) + (A0 * B0)$
- ◆ Substituting Cin1 into Cin2:
  - $\text{Cin2} = (A1 * A0 * B0) + (A1 * A0 * \text{Cin0}) + (A1 * B0 * \text{Cin0}) + (B1 * A0 * B0) + (B1 * A0 * \text{Cin0}) + (B1 * B0 * \text{Cin0}) + (A1 * B1)$

# Carry Lookahead: Theory (II)

## ◆ Now define two new terms:

- Generate Carry at Bit i:  $g_i = A_i * B_i$
- Propagate Carry via Bit i:  $p_i = A_i \text{ xor } B_i$


## ◆ We can rewrite:

- $Cin1 = g_0 + (p_0 * Cin0)$
- $Cin2 = g_1 + (p_1 * g_0) + (p_1 * p_0 * Cin0)$
- $Cin3 = g_2 + (p_2 * g_1) + (p_2 * p_1 * g_0) + (p_2 * p_1 * p_0 * Cin0)$

## ◆ Carry going into bit 3 is 1 if

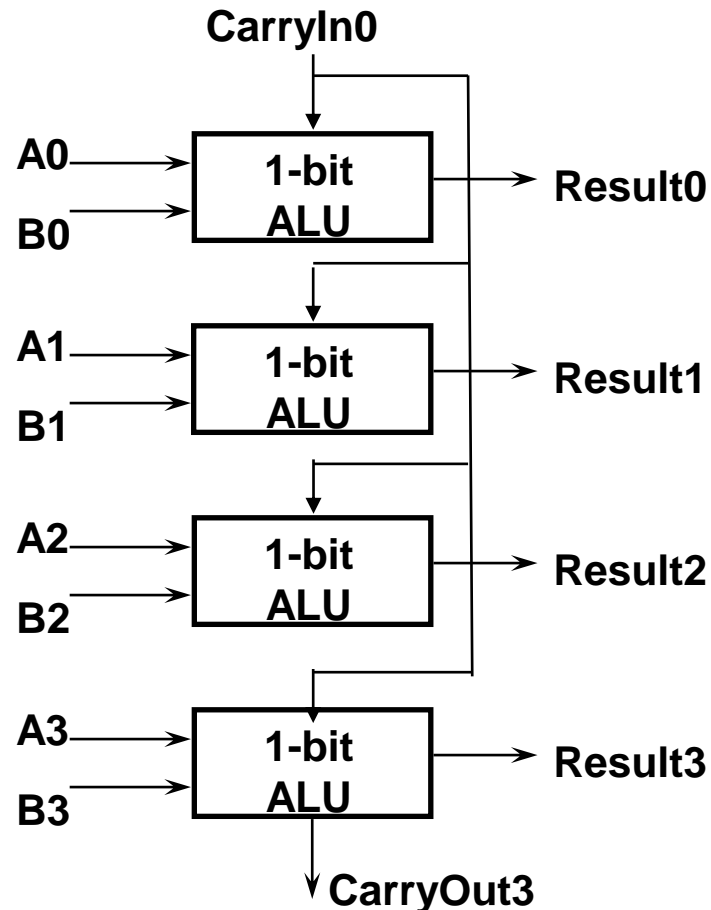
- We generate a carry at bit 2 ( $g_2$ )
- Or we generate a carry at bit 1 ( $g_1$ ) and bit 2 allows it to propagate ( $p_2 * g_1$ )
- Or we generate a carry at bit 0 ( $g_0$ ) and bit 1 as well as bit 2 allows it to propagate .....

# A Plumbing Analogy for Carry Lookahead (1, 2, 4 bits)



# Carry Lookahead Adder

- ◆ No Carry bit propagation from LSB to MSB

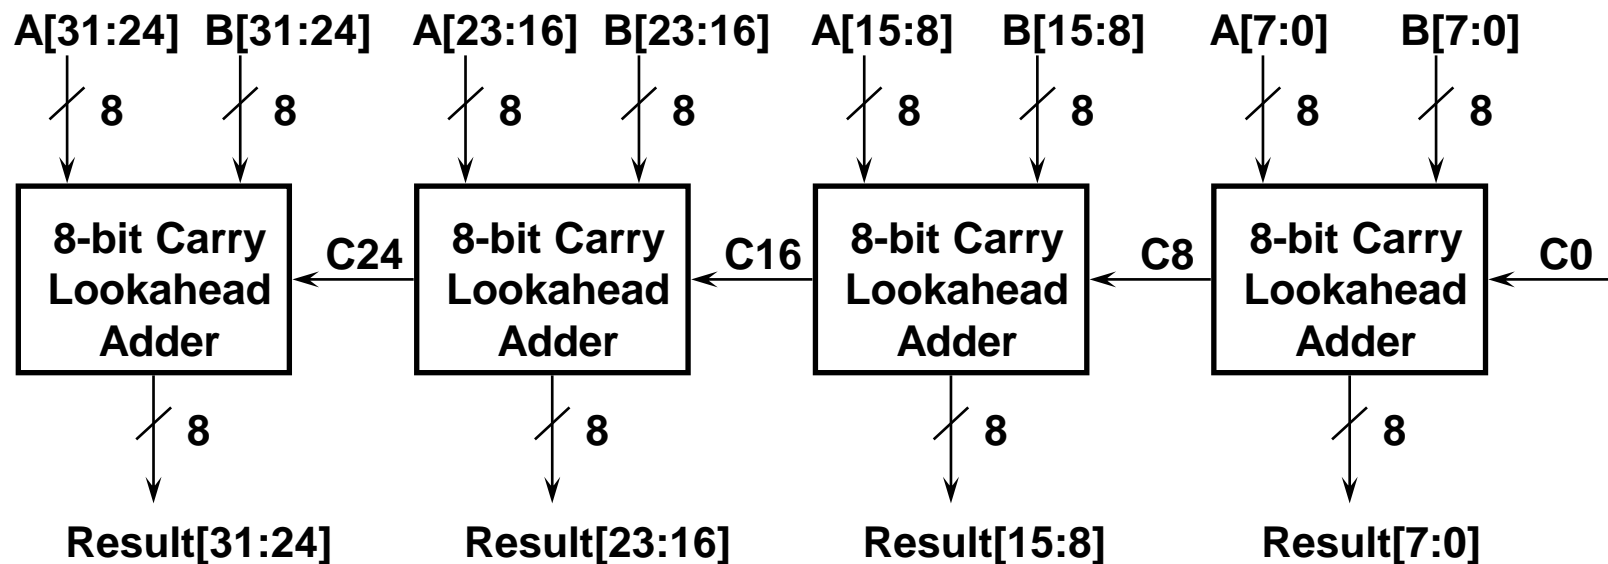


# Common Carry Lookahead Adder


- ◆ Expensive to build a “full” carry lookahead adder
  - Just imagine length of the equation for  $C_{in31}$
- ◆ Common practices:
  - Cascaded carry look-ahead adder
  - Multiple level carry look-ahead adder

# Cascaded Carry Lookahead

- ◆ Connects several N-bit lookahead adders to form a big one

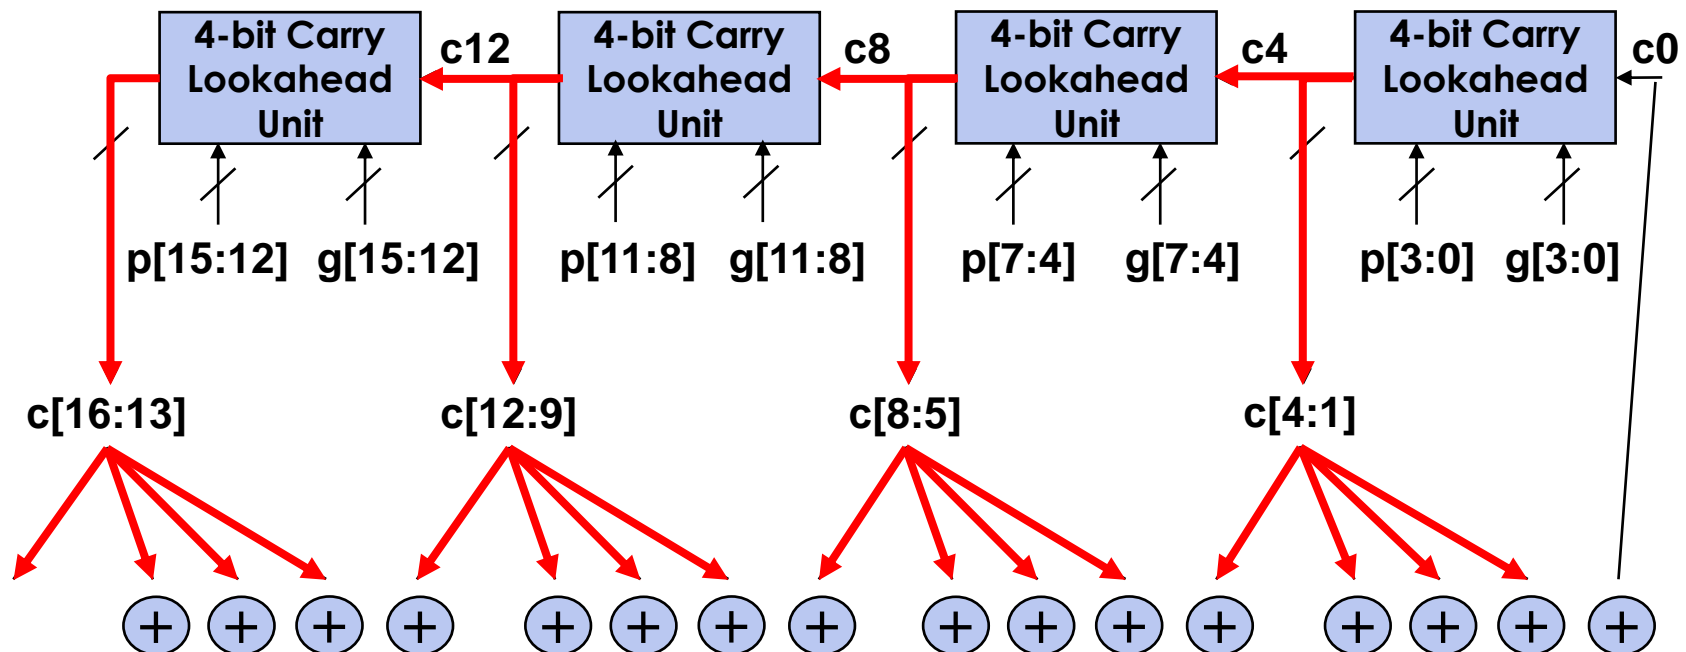


# Example: Carry Lookahead Unit



# Example: Cascaded Carry Lookahead

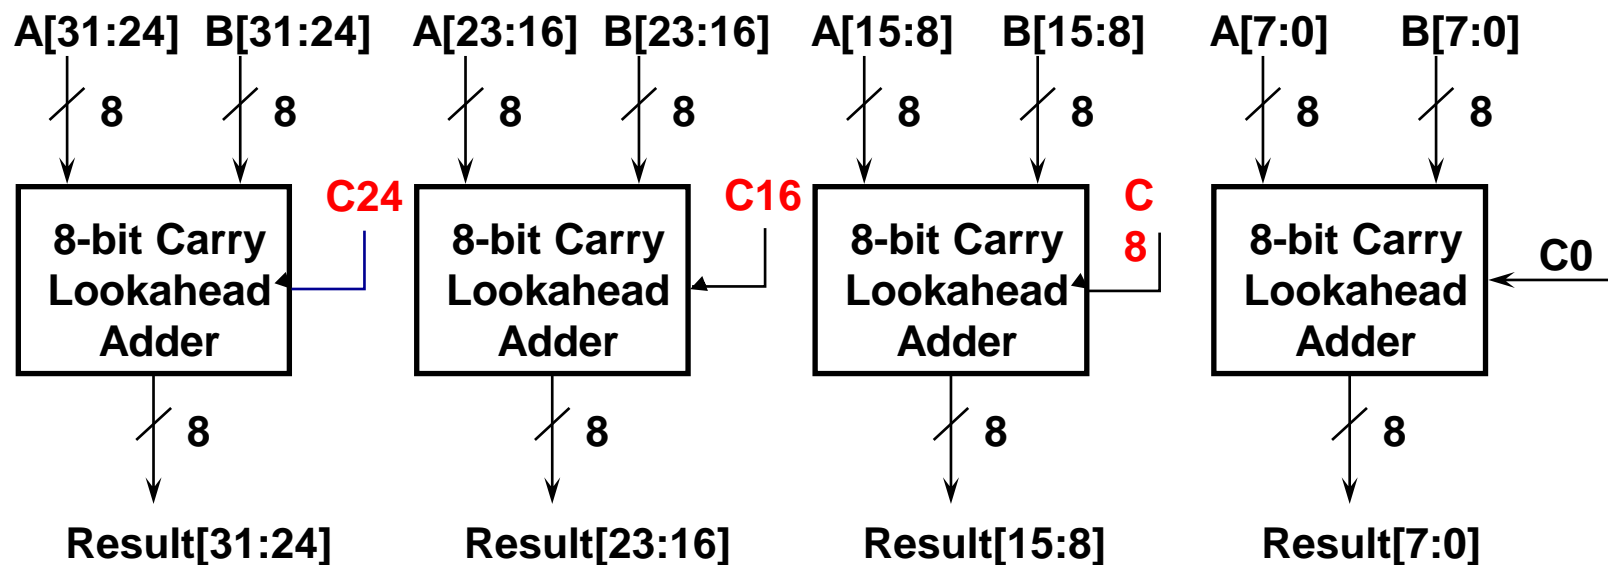
- ◆ Connects several N-bit lookahead adders to form a big one






# Multiple Level Carry Lookahead

- ◆ View an N-bit lookahead adder as a block
- ◆ Where to get Cin of the block ?

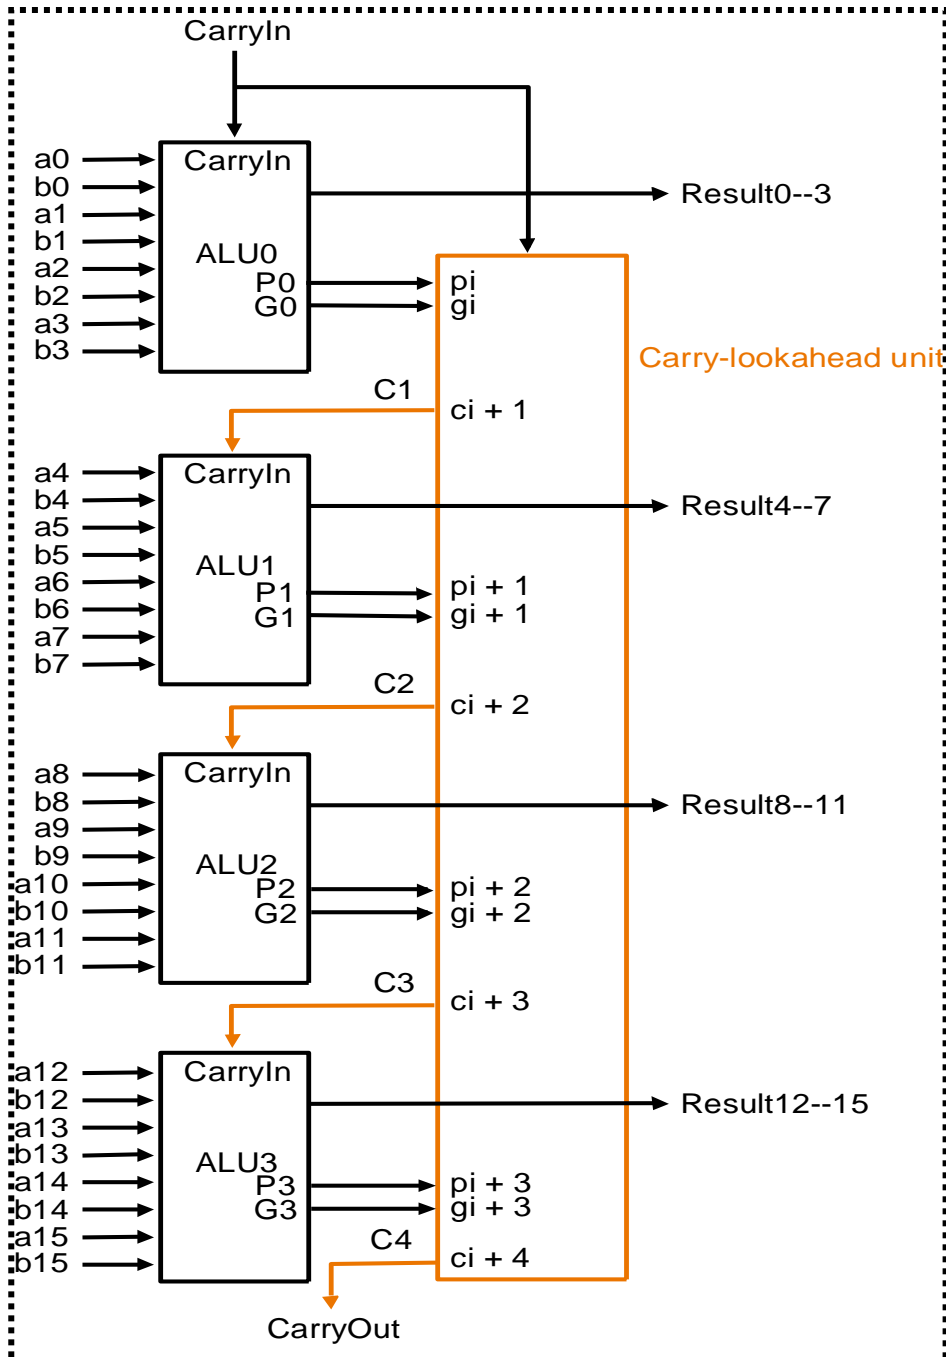


- Generate "super"  $P_i$  and  $G_i$  of the block
- Use next level carry lookahead structure to generate block  $C_{in}$

# A Plumbing Analogy for Carry Lookahead (Next Level P0 and G0)



# A Carry Lookahead Adder




A	B	Cout	
0	0	0	kill
0	1	Cin	propagate
1	0	Cin	propagate
1	1	1	generate

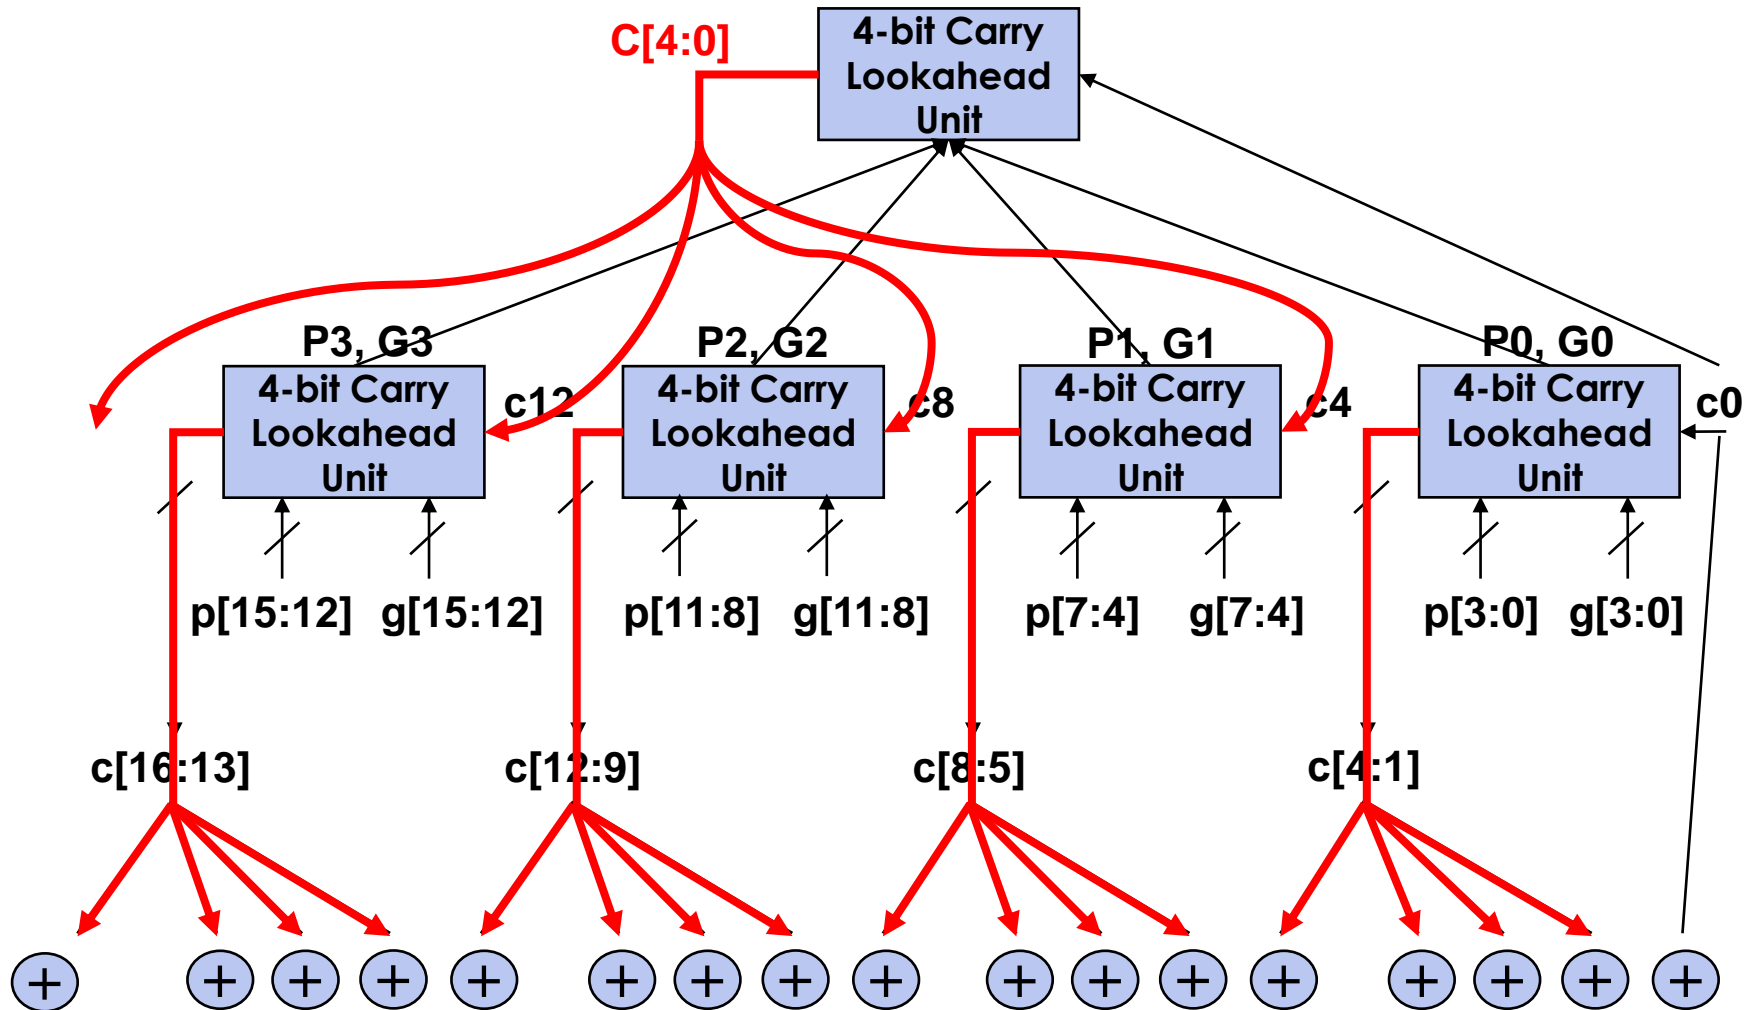
$$G = A * B$$

$$P = A + B$$

# Example: Carry Lookahead Unit




# Example: Multiple Level Carry Lookahead

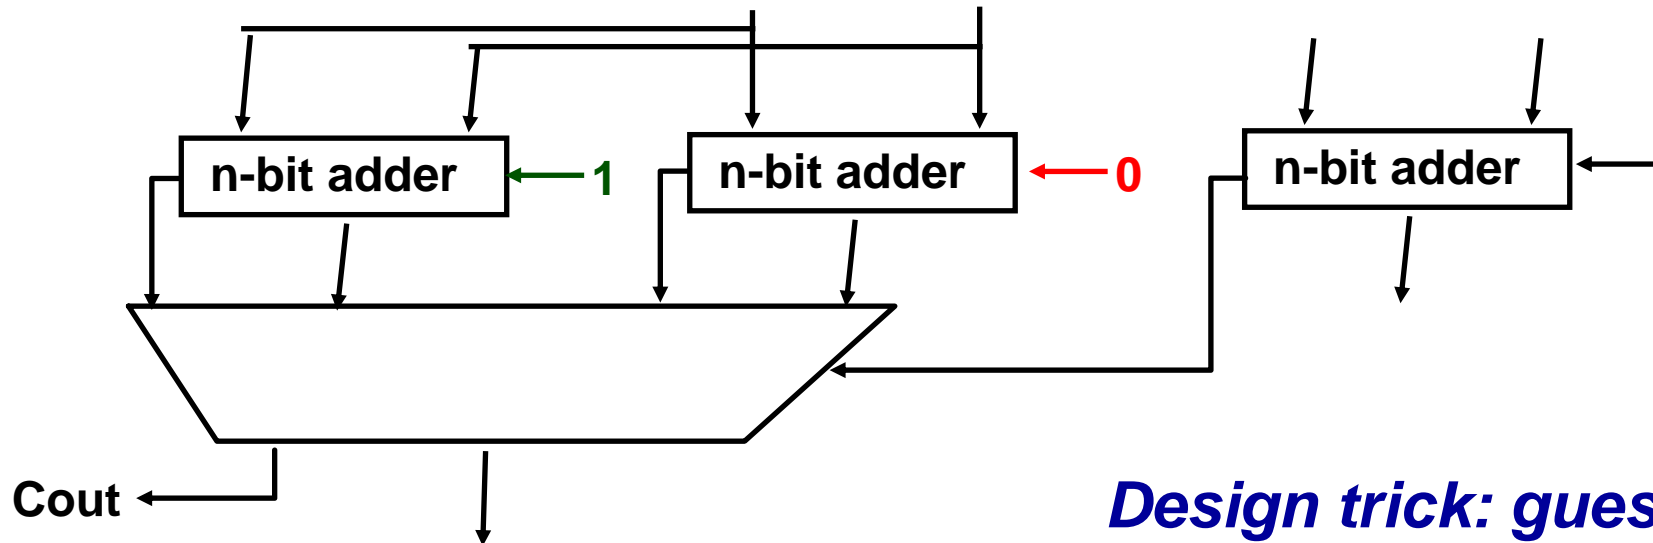


# Carry-select Adder

$$CP(2n) = 2 * CP(n)$$



$$CP(2n) = CP(n) + CP(\text{mux})$$



*Design trick: guess*

# Arithmetic for Multimedia

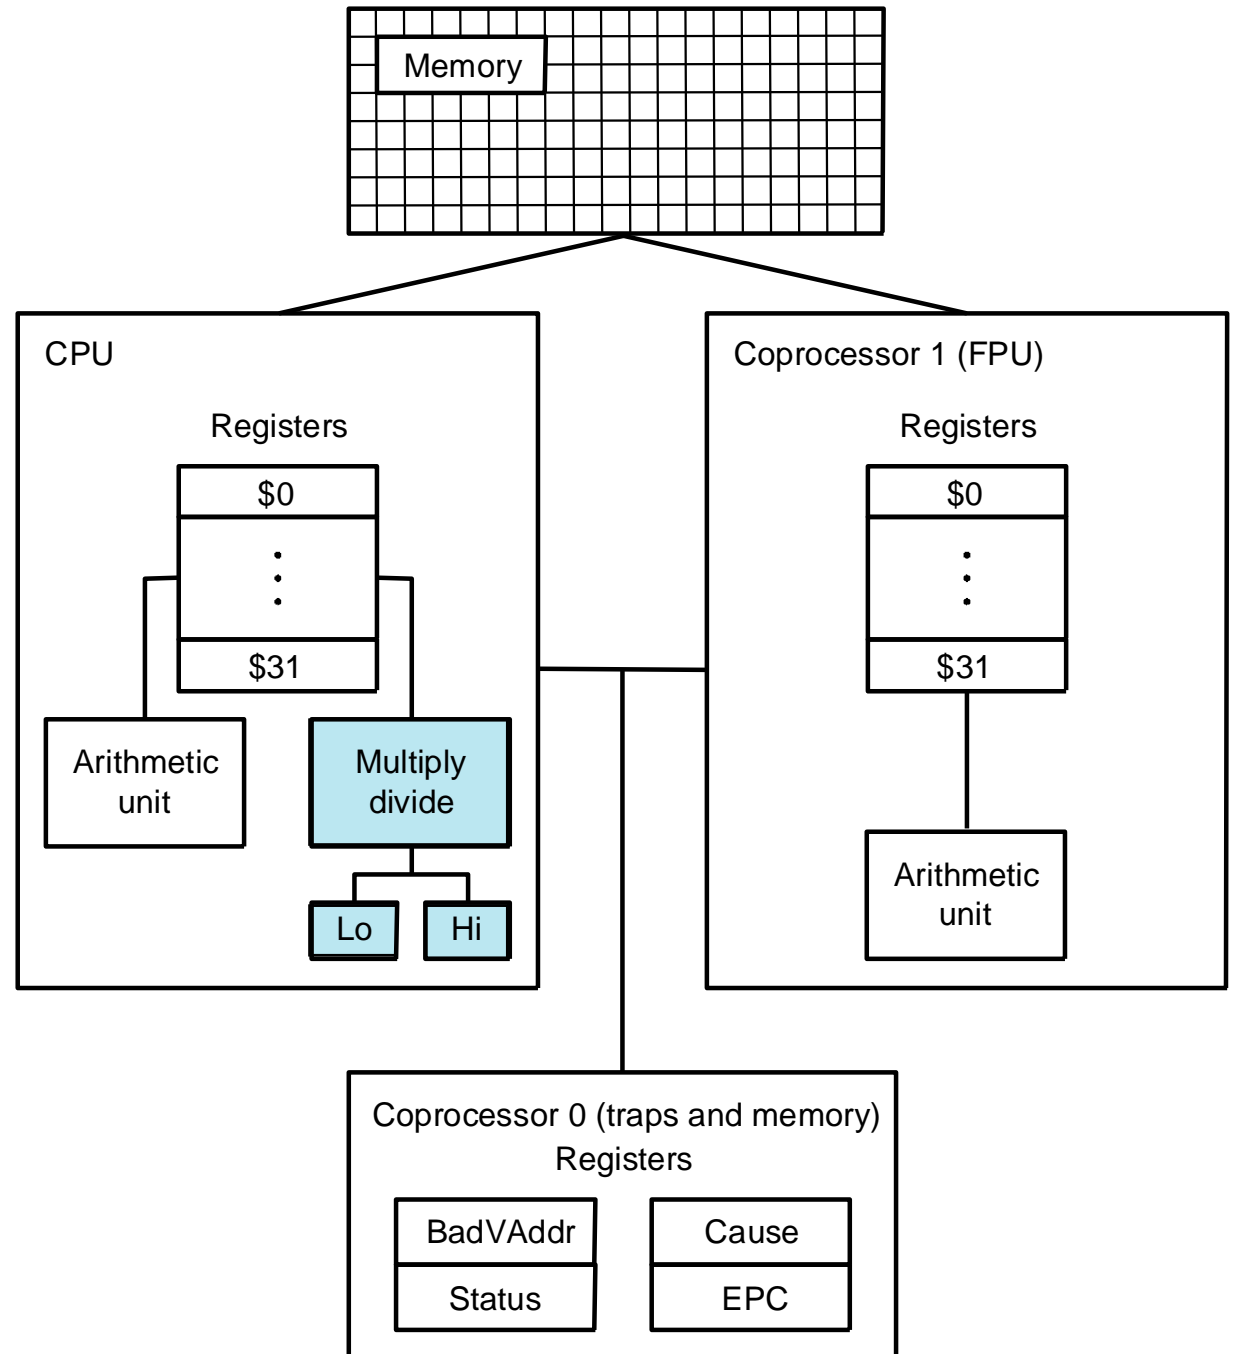
- ◆ Graphics and media processing operates on vectors of 8-bit and 16-bit data
  - Use 64-bit adder, with partitioned carry chain
    - Operate on 8×8-bit, 4×16-bit, or 2×32-bit vectors
  - SIMD (single-instruction, multiple-data)
- ◆ Saturating operations
  - On overflow, result is largest representable value
    - c.f. 2s-complement modulo arithmetic
  - E.g., clipping in audio, saturation in video

# Outline

- ◆ Addition and subtraction (Sec. 3.2)
- ◆ Constructing an arithmetic logic unit (Appendix C)
- ◆ Multiplication (Sec. 3.3, Appendix C)
- ◆ Division (Sec. 3.4)
- ◆ Floating point (Sec. 3.5)



# MIPS R2000 Organization



# Multiplication in MIPS

`mult $t1, $t2`                      `# t1 * t2`

- ◆ No destination register: product could be  $\sim 2^{64}$ ; need two special registers to hold it
- ◆ 3-step process:

\$t1	01111111111111111111111111111111
X \$t2	01000000000000000000000000000000

00011111111111111111111111111111	11000000000000000000000000000000
----------------------------------	----------------------------------

Hi

Lo

`mfhi $t3`

\$t3	00011111111111111111111111111111
------	----------------------------------

`mflo $t4`

\$t4	11000000000000000000000000000000
------	----------------------------------

# MIPS Multiplication

- ◆ Two 32-bit registers for product
  - HI: most-significant 32 bits
  - LO: least-significant 32-bits
- ◆ Instructions
  - `mult rs, rt` / `multu rs, rt`
    - 64-bit product in HI/LO
  - `mfhi rd` / `mflo rd`
    - Move from HI/LO to rd
    - Can test HI value to see if product overflows 32 bits
  - `mul rd, rs, rt`
    - Least-significant 32 bits of product → rd

# Unsigned Multiply

- ◆ Paper and pencil example (unsigned):

**Multiplicand**  
**Multiplier**

$$\begin{array}{r} 1000_{\text{ten}} \\ \times 1001_{\text{ten}} \\ \hline 1000 \end{array}$$

0000

0000

1000

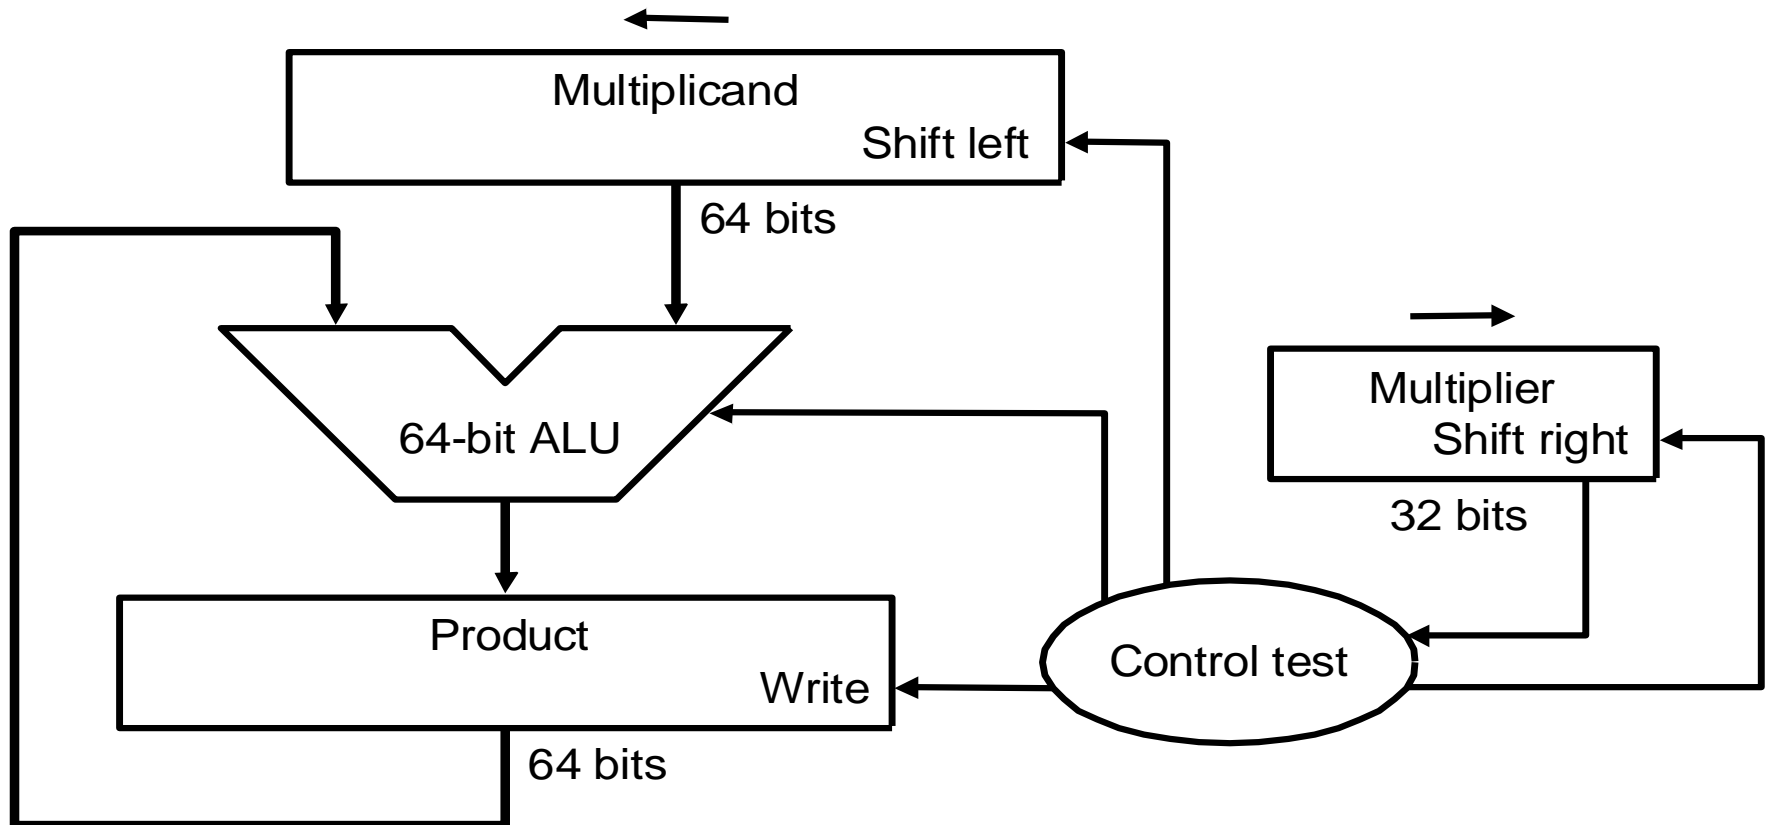
**Product**

$$\begin{array}{r} 1000 \\ \hline 01001000_{\text{ten}} \end{array}$$

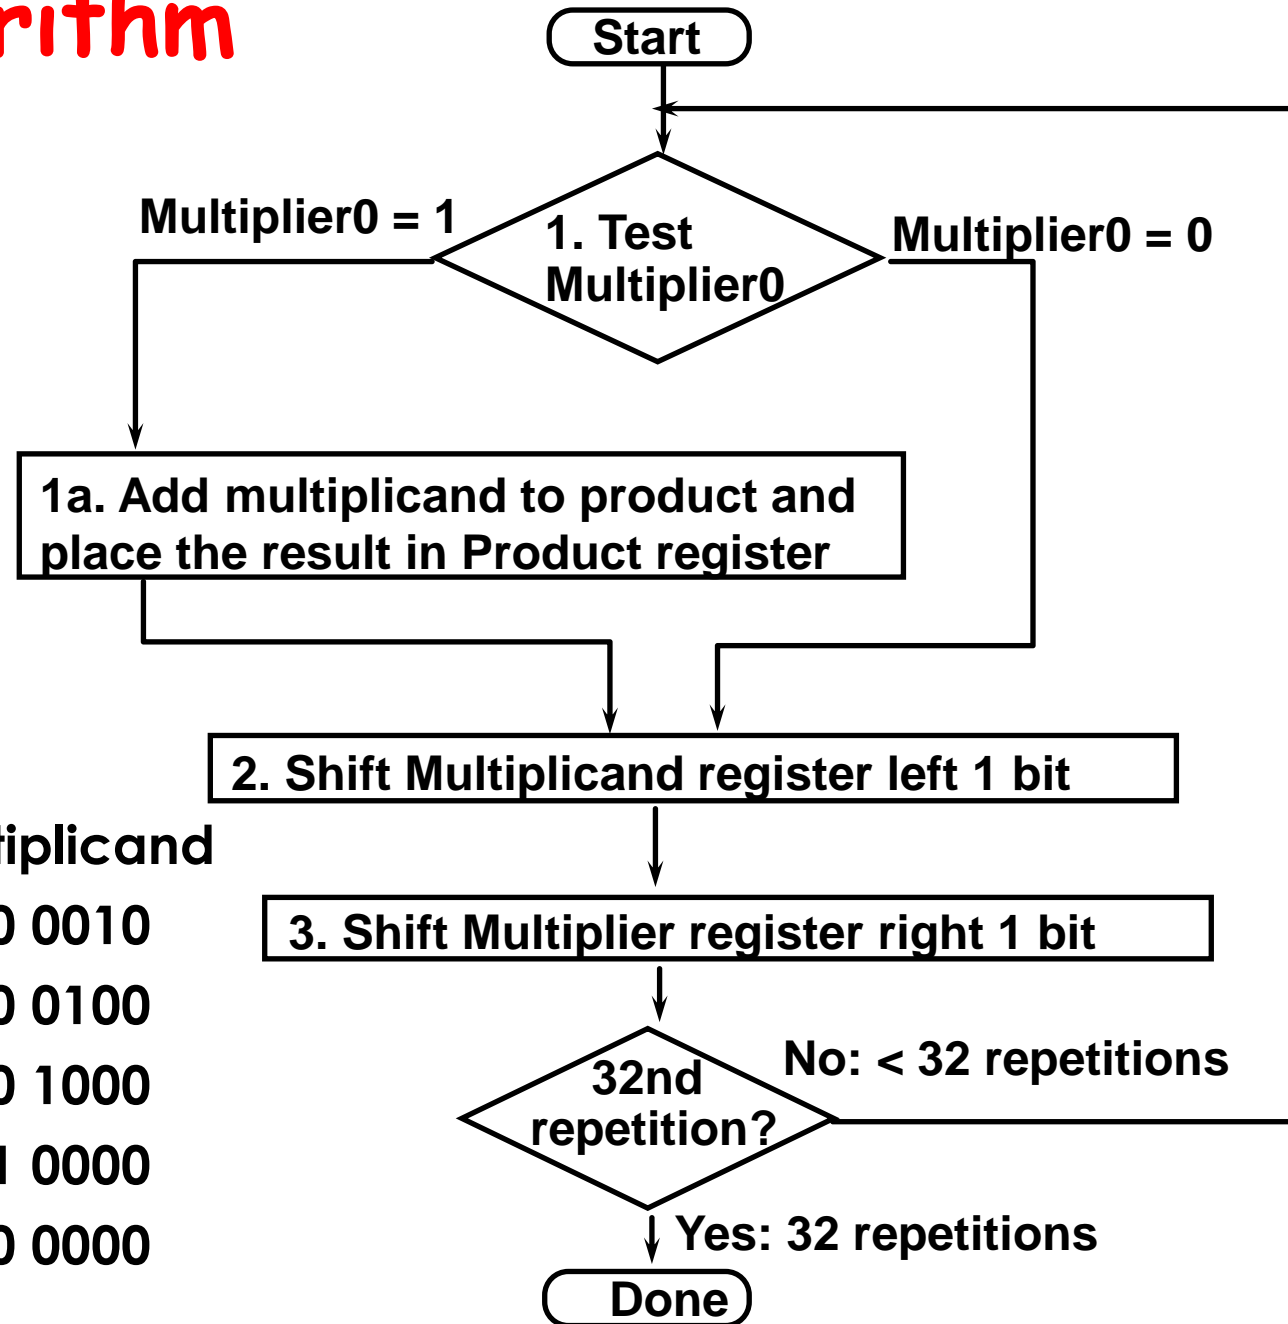
- ◆  $m$  bits  $\times$   $n$  bits =  $m+n$  bit product
- ◆ Binary makes it easy:
  - 0  $\Rightarrow$  place 0 (0  $\times$  multiplicand)
  - 1  $\Rightarrow$  place a copy (1  $\times$  multiplicand)
- ◆ 2 versions of multiply hardware and algorithm

# Unsigned Multiplier (Ver. 1)

- ◆ 64-bit *multiplicand register* (with 32-bit multiplicand at right half), 64-bit ALU, 64-bit *product register*, 32-bit *multiplier register*



# Multiply Algorithm (Ver. 1)



0010 x 0011

Product	Multiplier	Multiplicand
0000 0000	001 <b>1</b>	0000 0010
0000 0010	000 <b>1</b>	0000 0100
0000 0110	000 <b>0</b>	0000 1000
0000 0110	000 <b>0</b>	0001 0000
<b>0000 0110</b>	0000	0010 0000

# Observations: Multiply Ver. 1

- ◆ 1 clock per cycle => ~100 clocks per multiply
  - Ratio of multiply to add 5:1 to 100:1
- ◆ Half of the bits in multiplicand always 0  
=> 64-bit adder is wasted
- ◆ 0's inserted in right of multiplicand as shifted  
=> least significant bits of product never changed once formed
- ◆ Instead of shifting multiplicand to left, shift product to right?
- ◆ Product register wastes space => combine Multiplier and Product register

# Unsigned Multiply

- ◆ Paper and pencil example (unsigned):

**Multiplicand**  
**Multiplier**

$$\begin{array}{r} 1000_{\text{ten}} \\ \times 1001_{\text{ten}} \\ \hline 1000 \end{array}$$

0000

0000

1000

**Product**

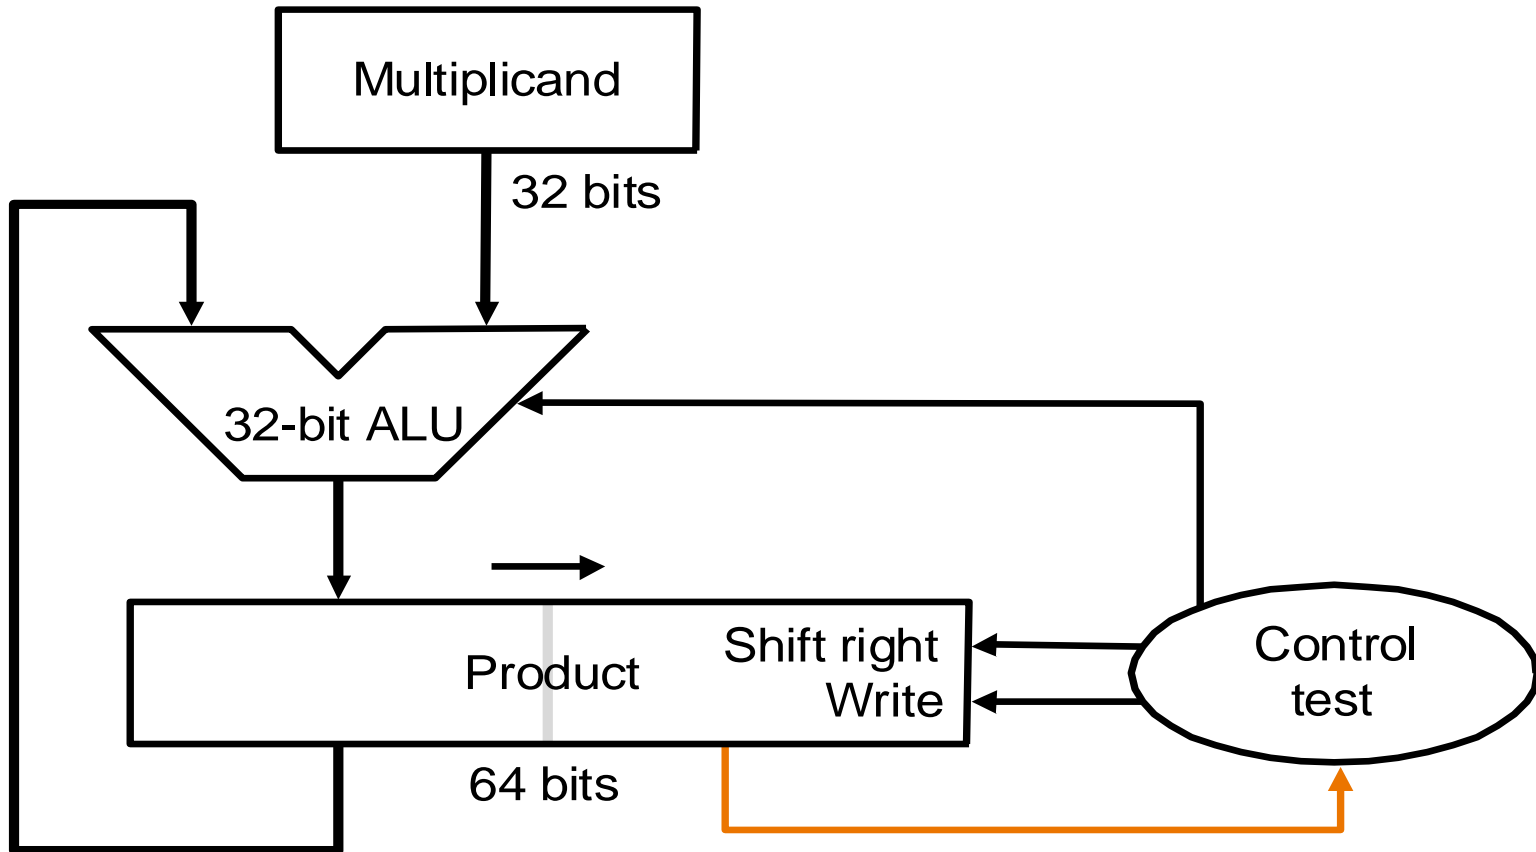
$$\begin{array}{r} 01001000_{\text{ten}} \end{array}$$

- ◆  $m$  bits  $\times$   $n$  bits =  $m+n$  bit product
- ◆ Binary makes it easy:
  - 0  $\Rightarrow$  place 0 (0  $\times$  multiplicand)
  - 1  $\Rightarrow$  place a copy (1  $\times$  multiplicand)
- ◆ 2 versions of multiply hardware and algorithm

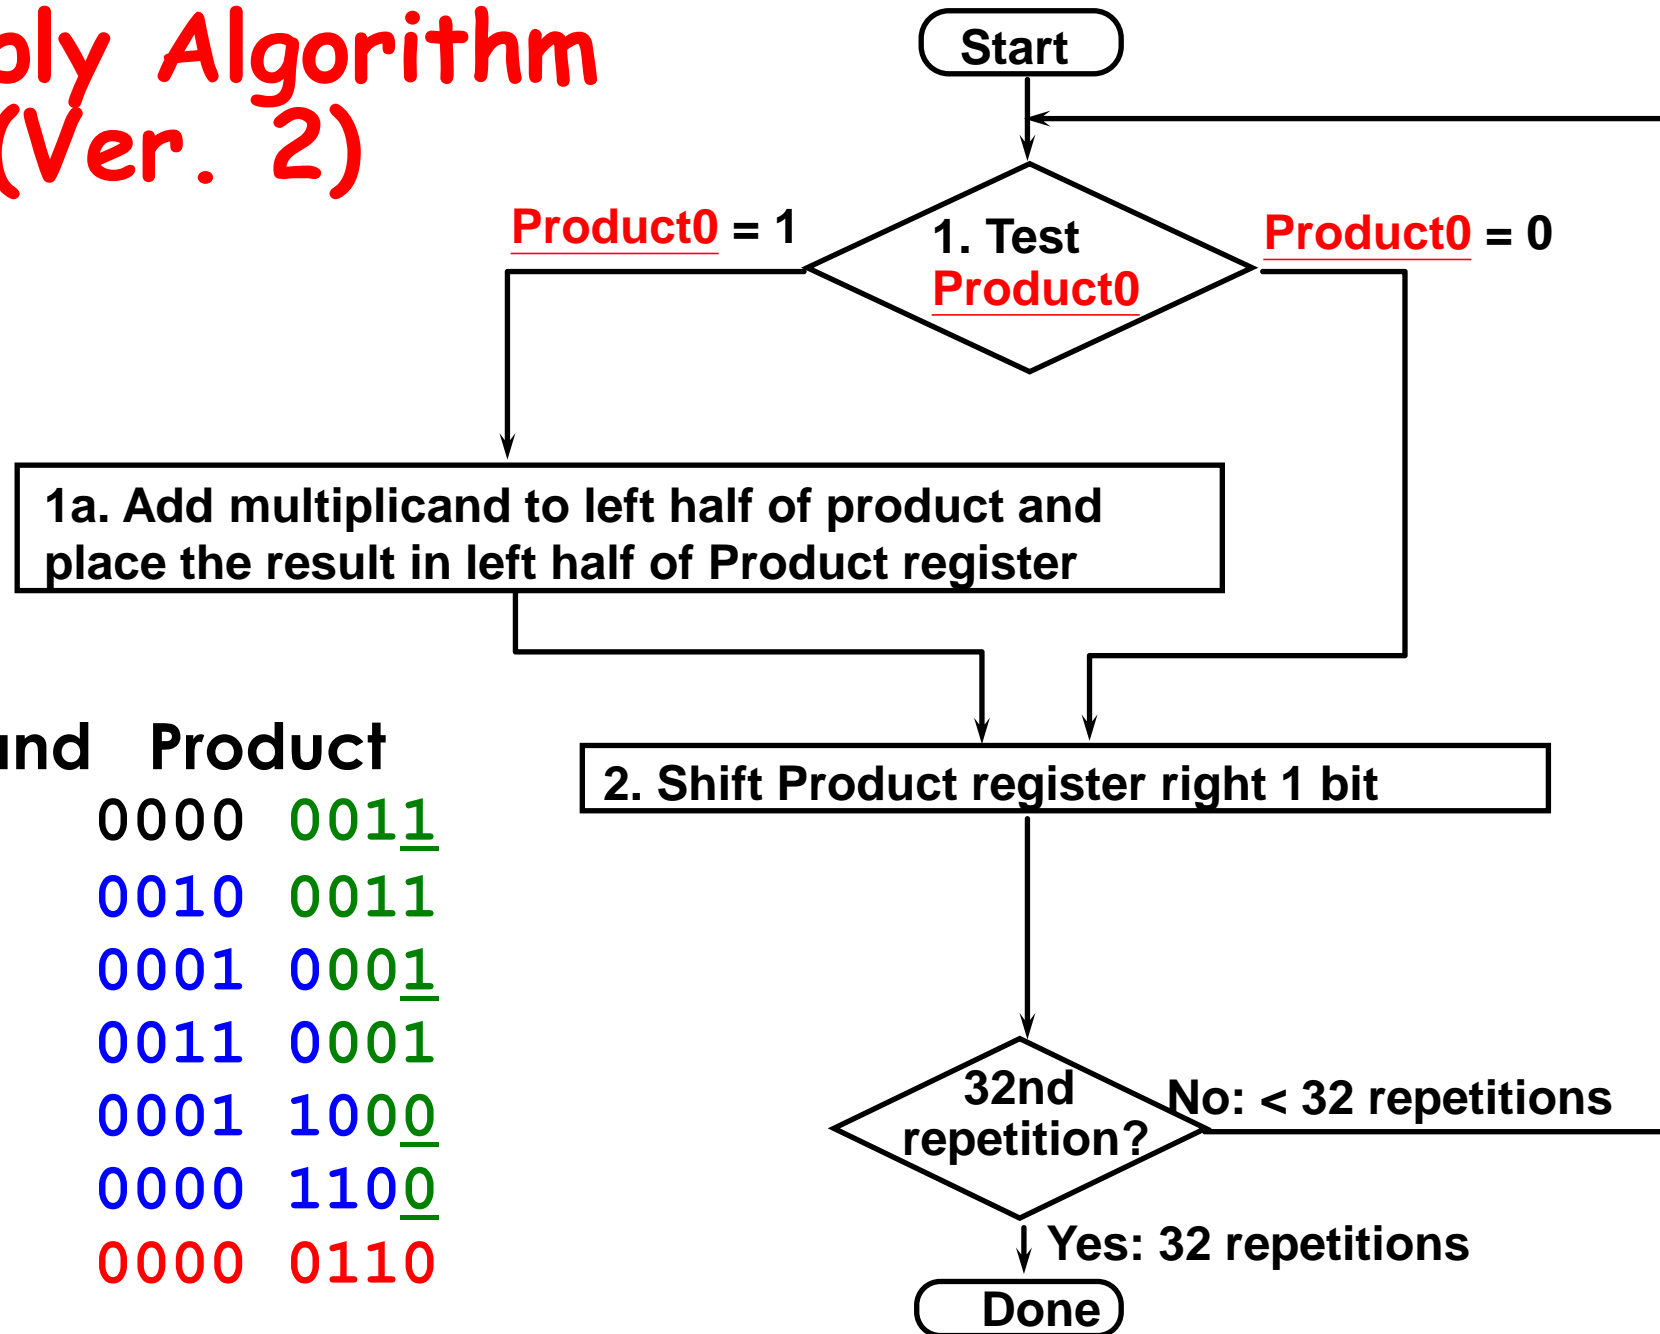


# Unsigned Multiplier (Ver. 2)

- ◆ 32-bit Multiplicand register, 32-bit ALU, 64-bit Product register (HI & LO in MIPS), (0-bit Multiplier register)



# Multiply Algorithm (Ver. 2)



Multiplicand	Product
0010	0000 00 <u>11</u>
	0010 00 <u>11</u>
0010	0001 000 <u>1</u>
	0011 000 <u>1</u>
0010	0001 100 <u>0</u>
0010	0000 110 <u>0</u>
0010	0000 01 <u>10</u>

# Observations: Multiply Ver. 2

- ◆ 2 steps per bit because multiplier and product registers combined
- ◆ MIPS registers Hi and Lo are left and right half of Product register  
=> this gives the MIPS instruction MultU
- ◆ What about signed multiplication?
  - The easiest solution is to make both positive and remember whether to complement product when done (leave out sign bit, run for 31 steps)
  - Apply definition of 2's complement
    - sign-extend partial products and subtract at end
  - Booth's Algorithm is an elegant way to multiply signed numbers using same hardware as before and save cycles

# Signed Multiply

## ♦ Paper and pencil example (signed):

$$\begin{array}{r}
 \text{Multiplicand} \quad 1001 \text{ (-7)} \\
 \text{Multiplier} \quad \times 1001 \text{ (-7)} \\
 \hline
 11111001 \\
 + \quad 0000000 \\
 + \quad 0000000 \\
 - \quad 11001 \\
 \hline
 \text{Product} \quad 00110001 \text{ (49)}
 \end{array}$$

## ♦ Rule 1: Multiplicand sign extended


## ♦ Rule 2: Sign bit (s) of Multiplier

- 0 => 0 x multiplicand
- 1 => -1 x multiplicand

## ♦ Why rule 2 ?

- $X = s x_{n-2} x_{n-3} \dots x_1 x_0$  (2's complement)

- $\text{Value}(X) = -1 \times s \times 2^{n-1} + x_{n-2} \times 2^{n-2} + \dots + x_0 \times 2^0$


$$\begin{array}{r} 00100000 \\ - 00000001 \\ \hline 00011111 \end{array}$$

# Booth's Algorithm: Motivation

- ◆ Example:  $2 \times 6 = 0010 \times 0110$ :

	0010 <sub>two</sub>	
x	0110 <sub>two</sub>	
+	0000	shift (0 in multiplier)
+	0010	add (1 in multiplier)
+	0010	add (1 in multiplier)
+	0000	shift (0 in multiplier)
	0001100 <sub>two</sub>	

- ◆ Can get same result in more than one way:


$$6 = -2 + 8 \qquad 0110 = -00010 + 01000$$

- ◆ Basic idea: replace a string of 1s with an initial subtract on seeing a one and add after last one

	0010 <sub>two</sub>	
x	0110 <sub>two</sub>	
	0000	shift (0 in multiplier)
-	0010	sub (first 1 in multiplier)
	0000	shift (mid string of 1s)
+	0010	add (prior step had last 1)
	00001100 <sub>two</sub>	



# Booth's Algorithm: Rationale



Current Bit to bit	right	Explanation	Example	Op
1	0	Begins run of 1s	0000111 <u>1</u> 000	sub
1	1	Middle run of 1s	00001 <u>11</u> 1000	none
0	1	End of run of 1s	000 <u>01</u> 111000	add
0	0	Middle run of 0s	0 <u>00</u> 01111000	none

Originally for speed (when shift was faster than add)

◆ Why it works?

$$\begin{array}{r}
 -1 \\
 + 10000 \\
 \hline
 01111
 \end{array}$$

# Booth's Algorithm

1. Depending on the current and previous bits, do one of the following:
  - 00: Middle of a string of 0s, no arithmetic op.
  - 01: End of a string of 1s, so add multiplicand to the left half of the product
  - 10: Beginning of a string of 1s, so subtract multiplicand from the left half of the product
  - 11: Middle of a string of 1s, so no arithmetic op.
2. As in the previous algorithm, shift the Product register right (arithmetically) 1 bit



# Booths Example (2 x 7)

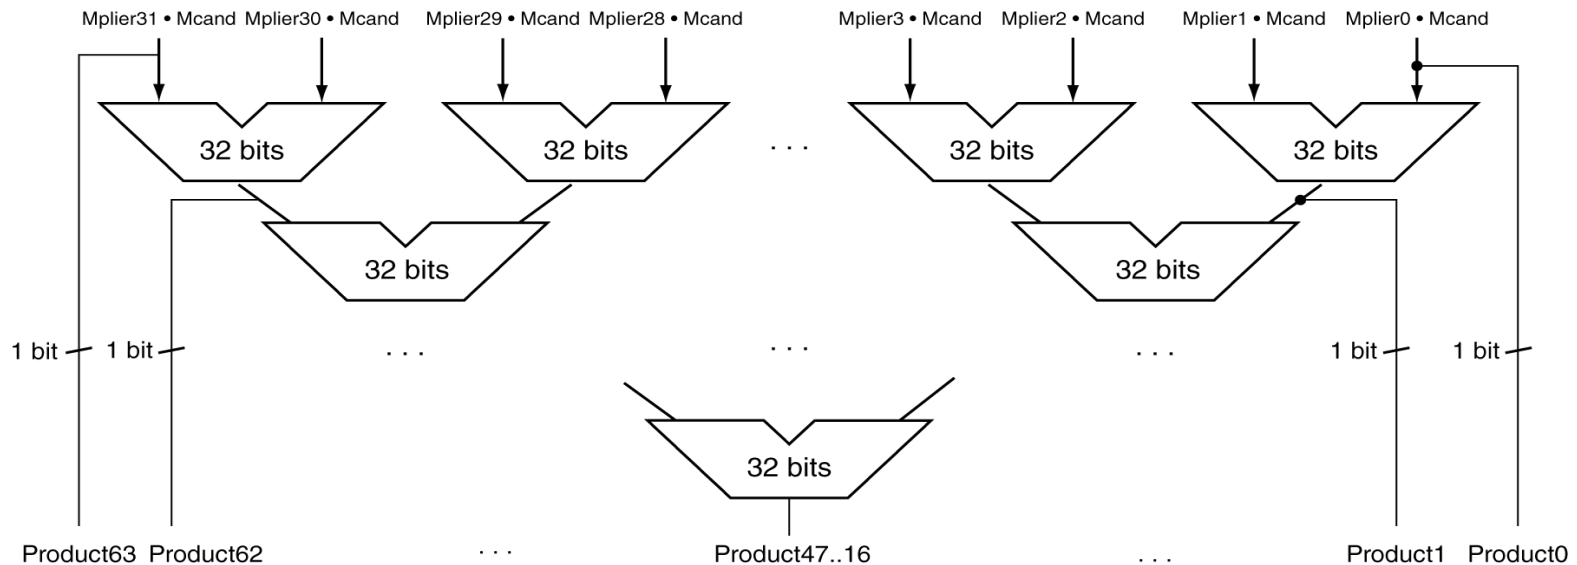
Operation	Multiplicand	Product	next?
0. initial value	0010	0000 <b>0111</b> 0	10 -> sub
1a. $P = P - m$	<b>1110</b>	+1110	
		1110 <b>0111</b> 0	shift P (sign ext)
1b.	0010	<u>1</u> 111 <b>0011</b> <b>1</b>	11 -> nop, shift
2.	0010	1111 1 <b>001</b> <b>1</b>	11 -> nop, shift
3.	0010	1111 11 <b>00</b> <b>1</b>	01 -> add
4a.	<b>0010</b>	+0010	
		0001 11 <b>00</b> <b>1</b>	shift
4b.	0010	0000 111 <b>0</b> <b>0</b>	done

# Booths Example (2 x -3)

Operation	Multiplicand	Product	next?
0. initial value	0010	0000 <b>1101</b> 0	10 -> sub
1a. $P = P - m$	<b>1110</b>	+1110	
		1110 <b>1101</b> 0	shift P (sign ext)
1b.	0010	<u>1</u> 111 0 <b>110</b> <b>1</b>	01 -> add
	<b>0010</b>	+0010	
2a.		0001 0 <b>110</b> <b>1</b>	shift P
2b.	0010	0000 10 <b>11</b> <b>0</b>	10 -> sub
	<b>1110</b>	+1110	
3a.	0010	1110 10 <b>11</b> <b>0</b>	shift
3b.	0010	1111 010 <b>1</b> <b>1</b>	<b>11</b> -> nop
4a.		1111 010 <b>1</b> <b>1</b>	shift
4b.	0010	1111 1010 <b>1</b>	done

# Faster Multiplier

- ◆ A combinational multiplier
- ◆ Use multiple adders
  - Cost/performance tradeoff



- ◆ Can be pipelined
  - Several multiplication performed in parallel

# Wallace Tree Multiplier

- ◆ Use carry save adders: three inputs and two outputs

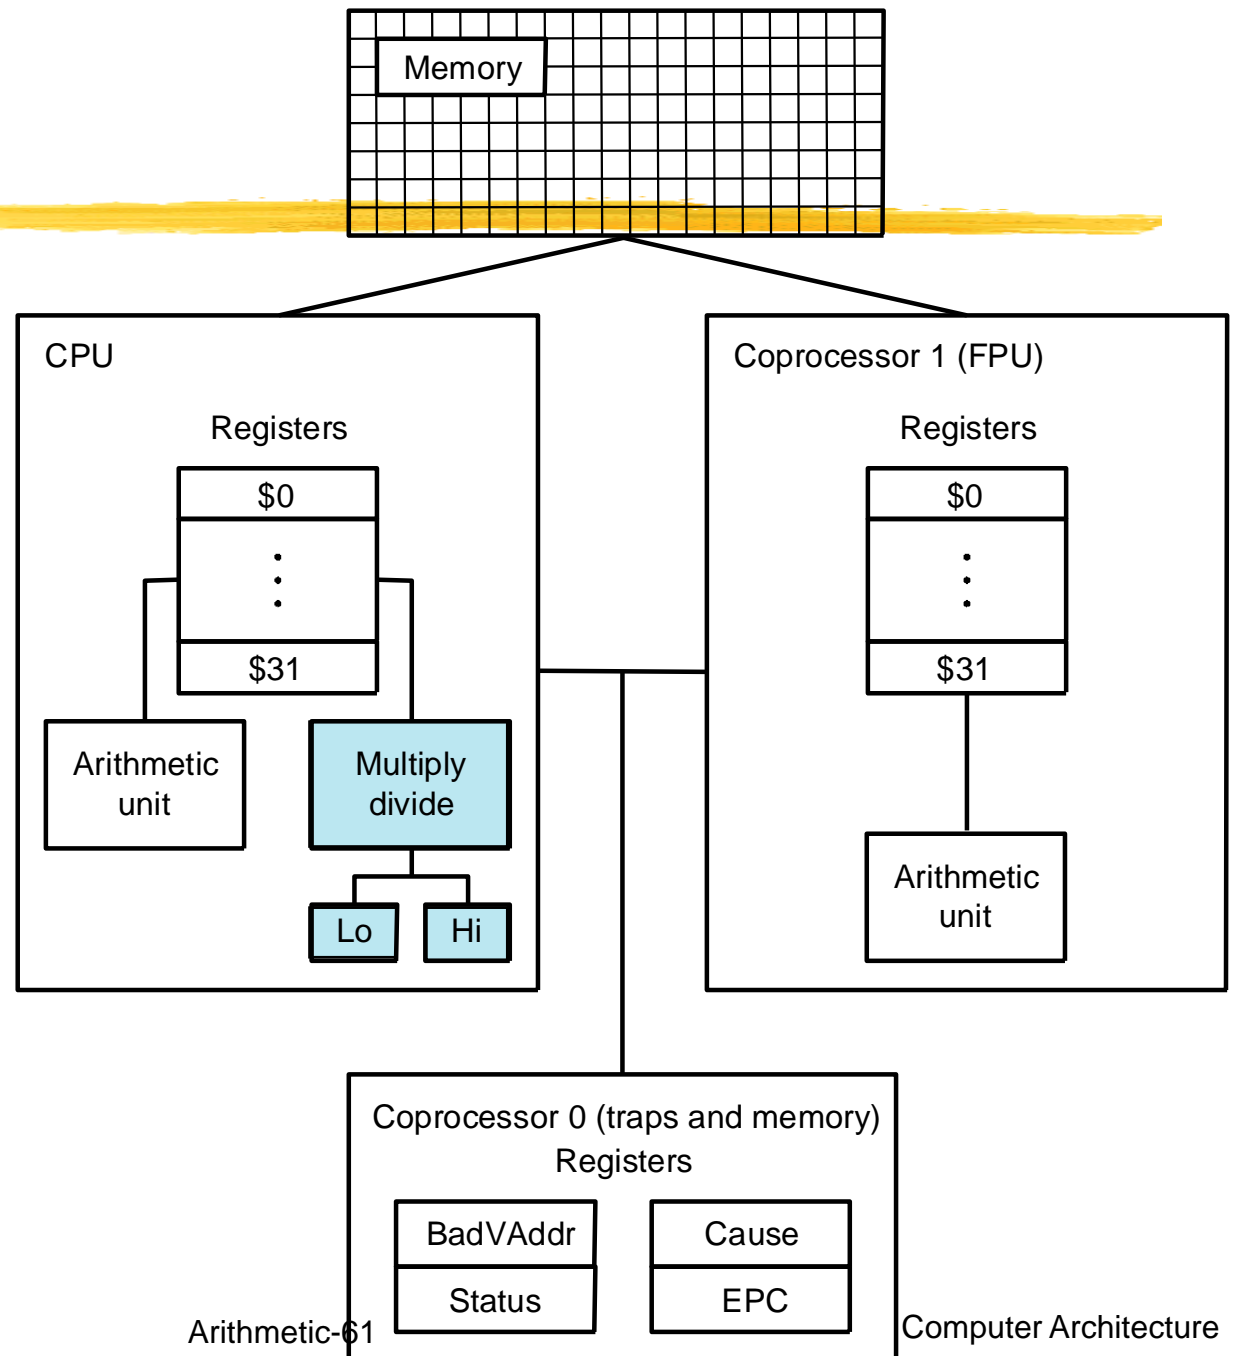
$$\begin{array}{r} 1\ 0\ 1\ 0\ 1\ 1\ 1\ 0 \\ 0\ 0\ 1\ 0\ 0\ 0\ 1\ 1 \\ 1\ 0\ 0\ 0\ 0\ 1\ 1\ 1 \\ \hline 0\ 0\ 0\ 0\ 1\ 0\ 1\ 0\ (\text{sum}) \\ 1\ 0\ 1\ 0\ 0\ 1\ 1\ 1\ (\text{carry}) \end{array}$$

- ◆ 8 full adders
- ◆ One full adder delay (no carry propagation)
- ◆ The last stage is performed by regular adder
- ◆ What is the minimum delay for 16 x 16 multiplier ?

# Outline

- ◆ Addition and subtraction (Sec. 3.2)
- ◆ Constructing an arithmetic logic unit (Appendix C)
- ◆ Multiplication (Sec. 3.3, Appendix C)
- ◆ Division (Sec. 3.4)
- ◆ Floating point (Sec. 3.5)

# MIPS R2000 Organization



# Division in MIPS

`div $t1, $t2           # t1 / t2`

- ◆ Quotient stored in Lo, remainder in Hi

`mflo $t3           #copy quotient to t3`

`mfhi $t4           #copy remainder to t4`

- ◆ 3-step process

- ◆ Unsigned division:

`divu $t1, $t2       # t1 / t2`

- Just like `div`, except now interpret `t1`, `t2` as unsigned integers instead of signed
- Answers are also unsigned, use `mfhi`, `mflo` to access

- ◆ No overflow or divide-by-0 checking

- Software must perform checks if required

# Divide: Paper & Pencil

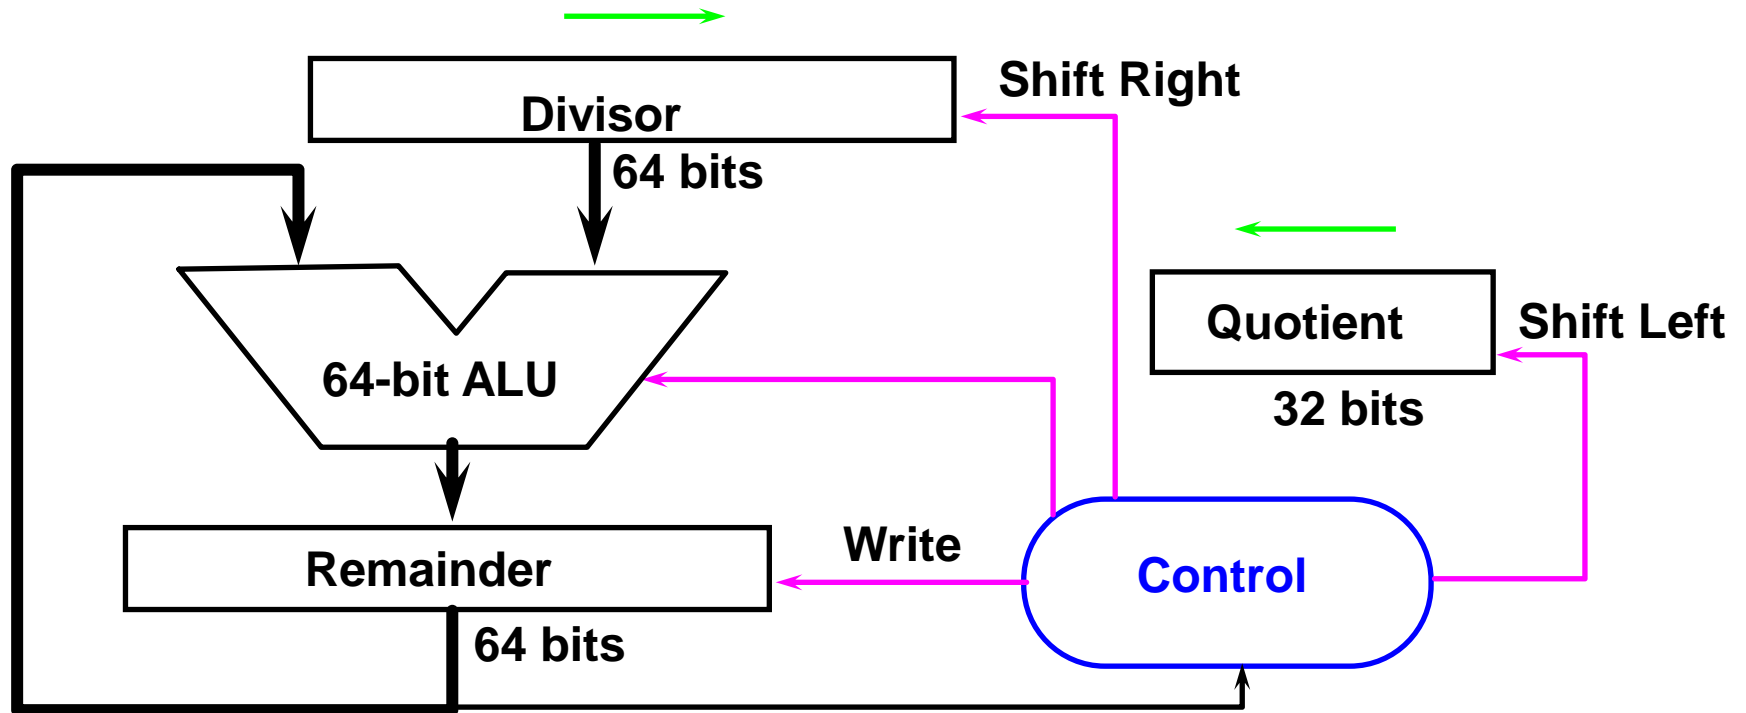
		1001 <sub>ten</sub>	Quotient
Divisor	1000 <sub>ten</sub>	$\overline{)1001010}_{ten}$	Dividend
		-1000	
		0010	
		0101	
		1010	
		-1000	
		10 <sub>ten</sub>	Remainder

- ◆ See how big a number can be subtracted, creating quotient bit on each step  
Binary  $\Rightarrow$  1 \* divisor or 0 \* divisor
- ◆ Two versions of divide, successive refinement
- ◆ Both dividend and divisor are 32-bit positive integers



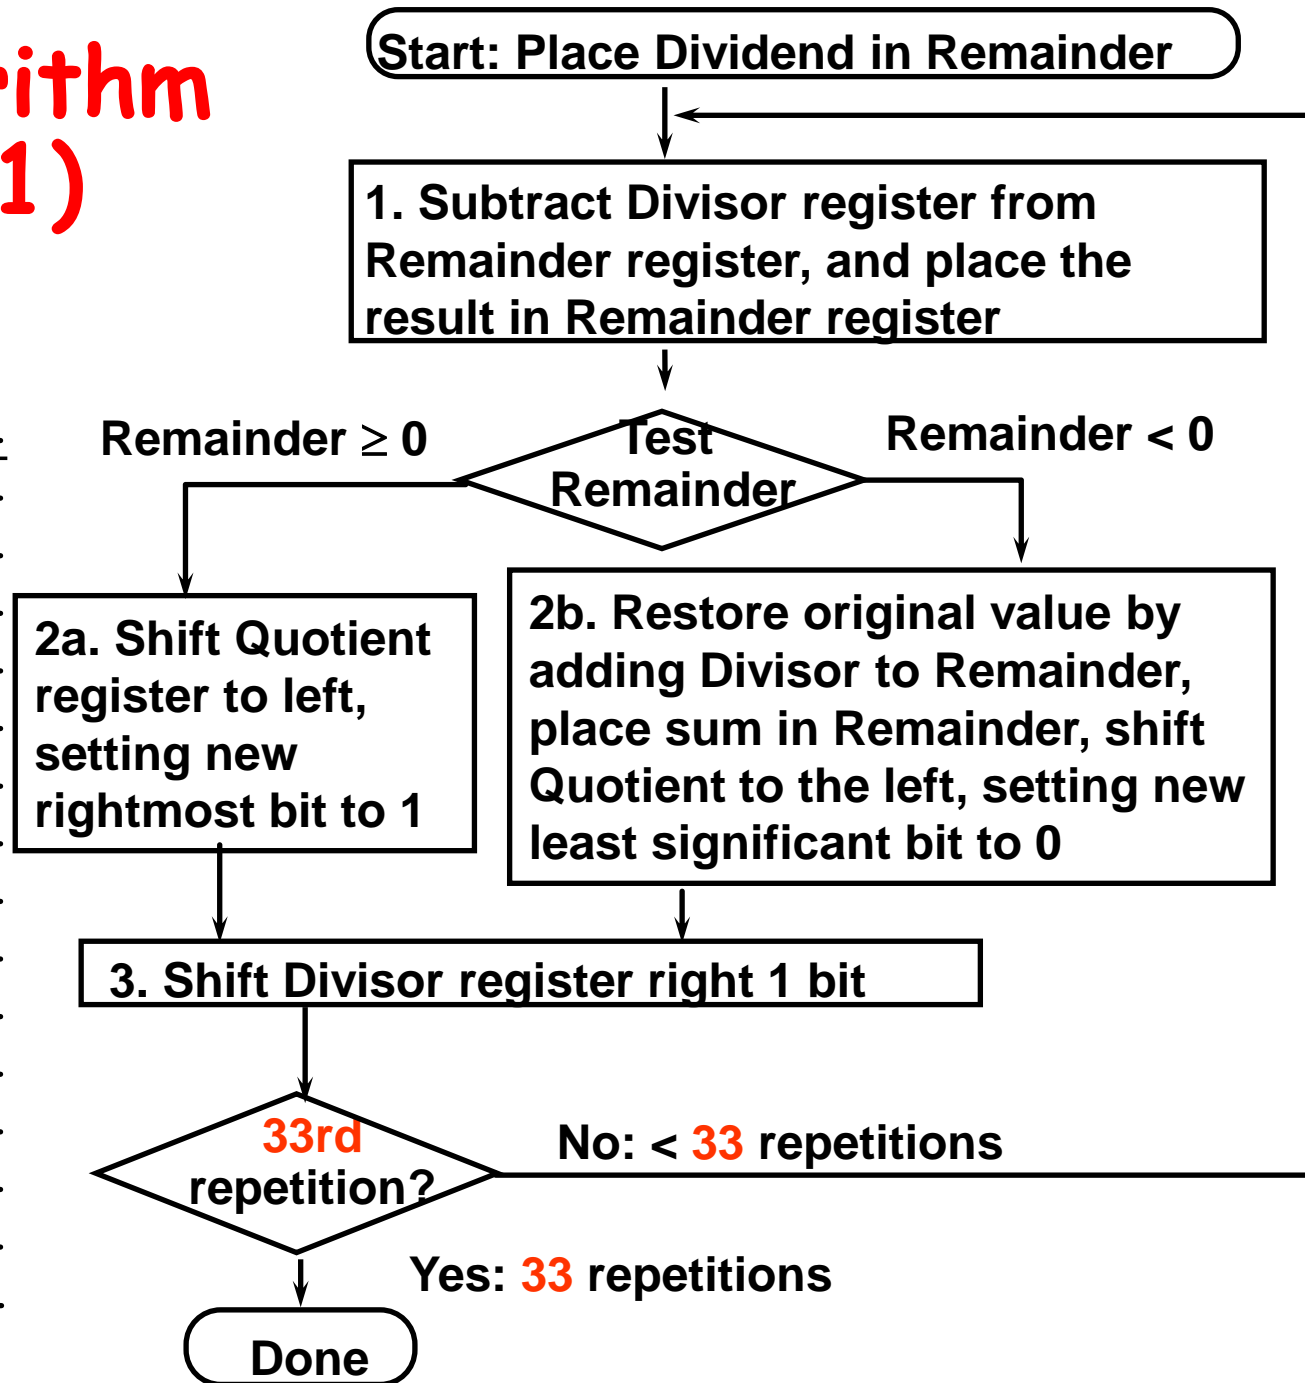
# Divide Hardware (Version 1)

- ◆ 64-bit *Divisor* register (initialized with 32-bit divisor in left half), 64-bit ALU, 64-bit *Remainder* register (initialized with 64-bit dividend), 32-bit *Quotient* register



# Divide Algorithm (Version 1)

Quot.	Divisor	Rem.
0000	<u>00100000</u>	00000111 11100111 00000111
0000	<u>00010000</u>	00000111 11110111 00000111
0000	<u>00001000</u>	00000111 11111111 00000111
0000	<u>00000100</u>	00000111 00000011 00000011
0001		00000011
0001	<u>00000010</u>	00000011 00000001 00000001
0011		00000001
0011	<u>00000001</u>	00000001

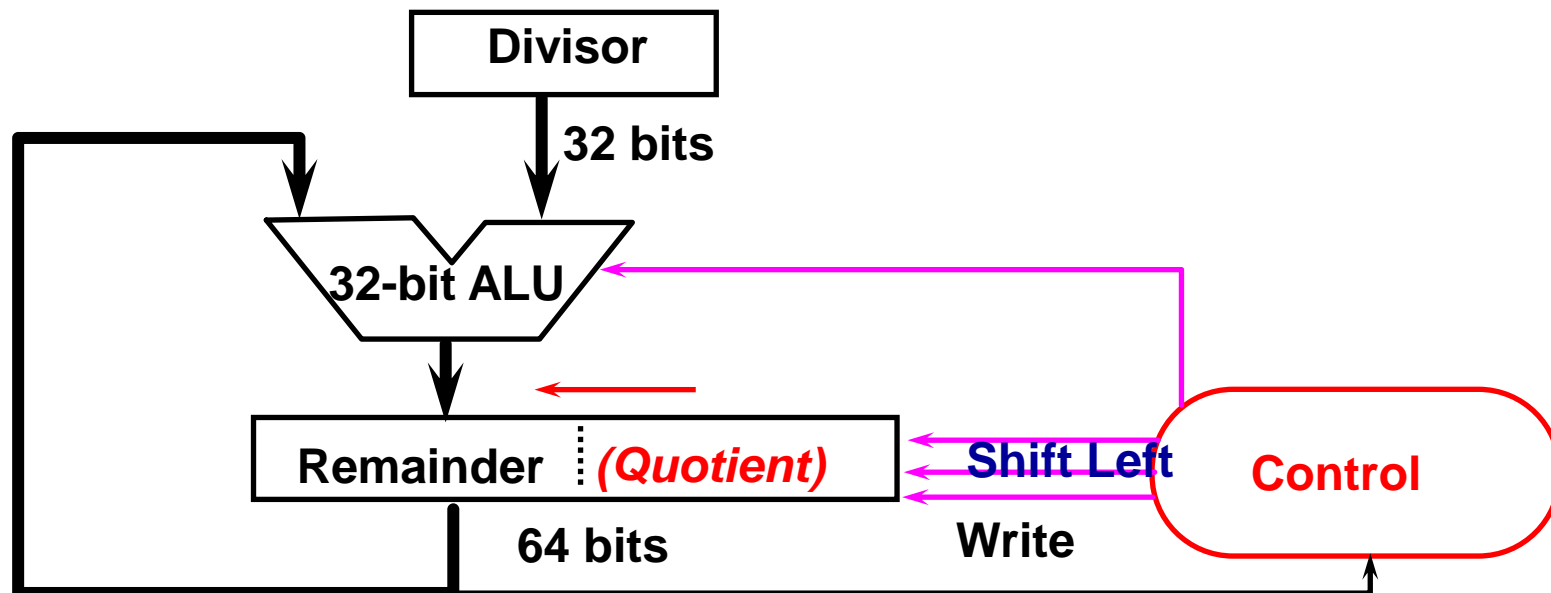


# Observations: Divide Version 1

- ◆ Half of the bits in divisor register always 0
  - => 1/2 of 64-bit adder is wasted
  - => 1/2 of divisor is wasted
- ◆ Instead of shifting divisor to right, shift remainder to left?
- ◆ 1st step cannot produce a 1 in quotient bit (otherwise quotient is too big for the register)
  - => switch order to shift first and then subtract
  - => save 1 iteration
- ◆ Eliminate Quotient register by combining with Remainder register as shifted left

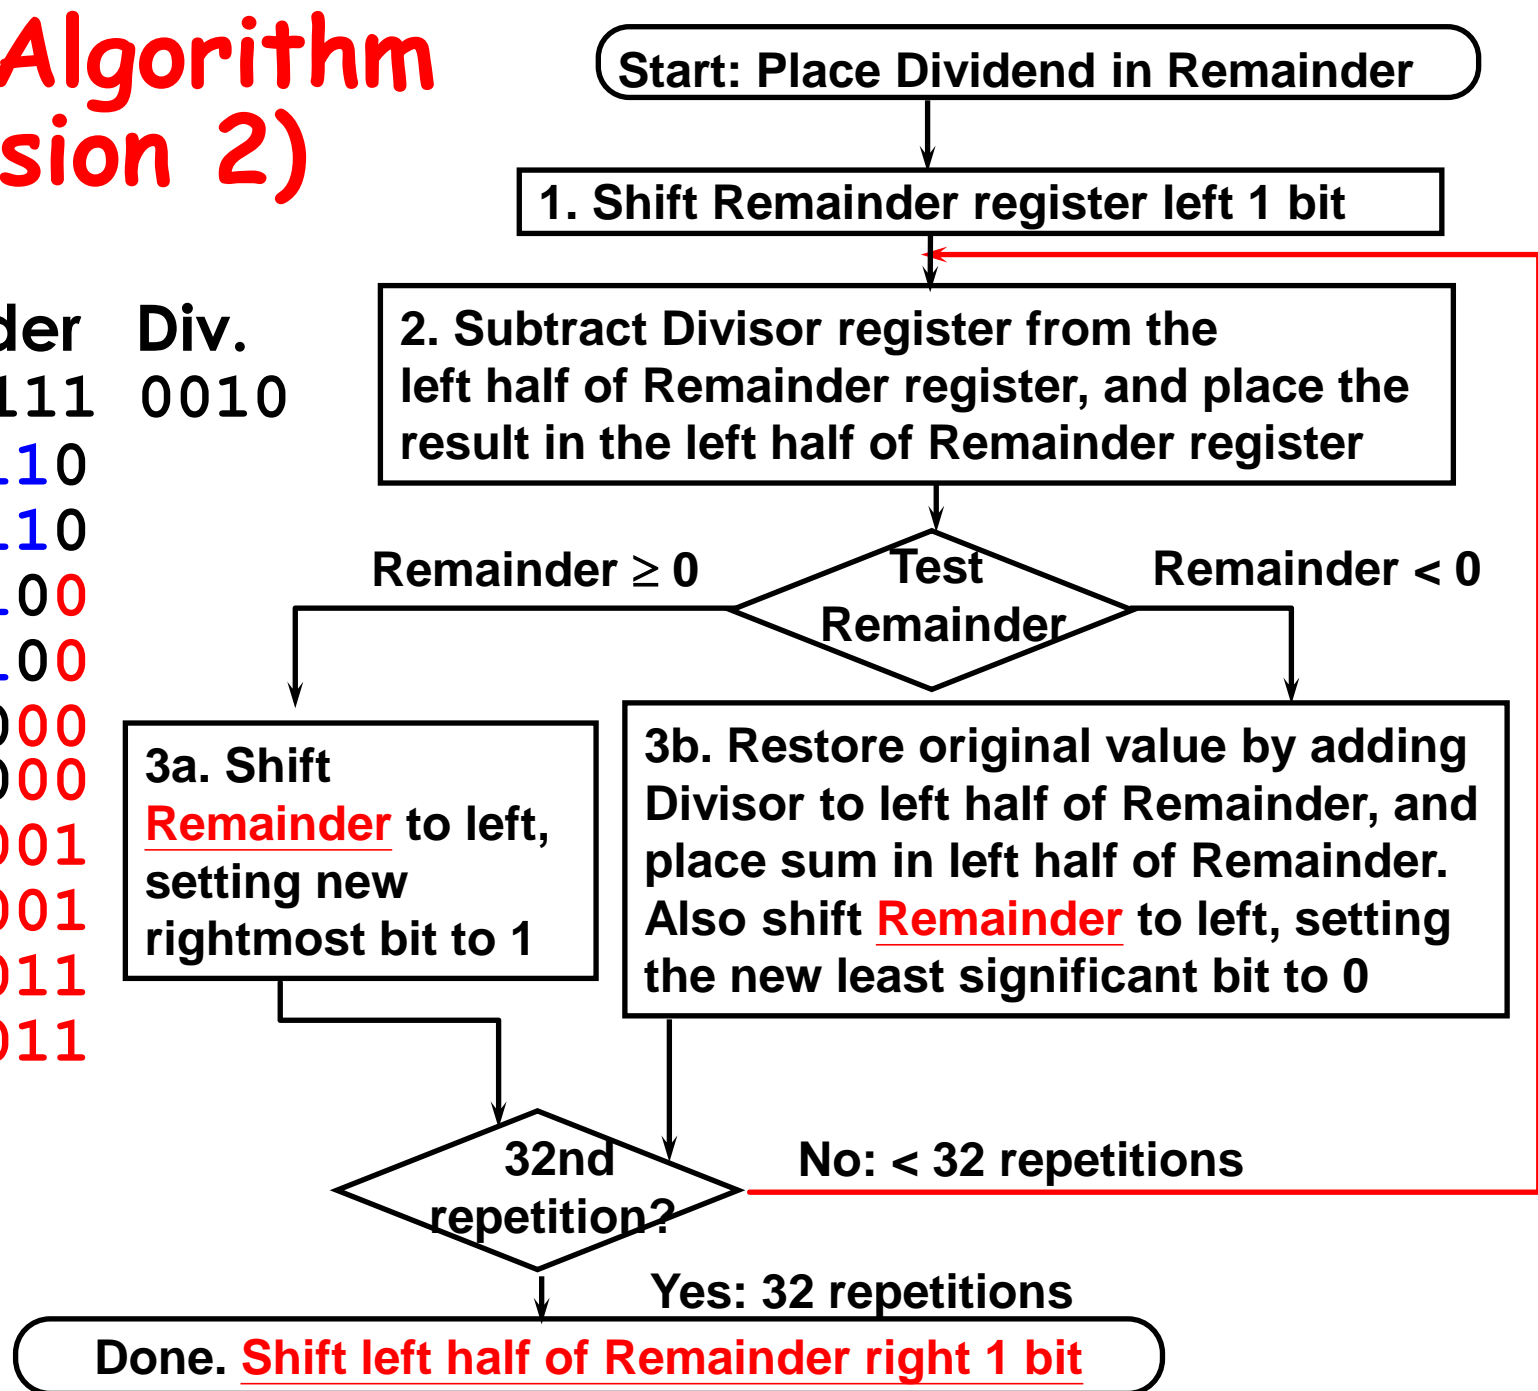
# Divide Hardware (Version 2)

- ◆ 32-bit Divisor register, 32-bit ALU, 64-bit Remainder register, (0-bit Quotient register)



# Divide Algorithm (Version 2)

Step	Remainder	Div.
0	0000 0111	0010
1.1	0000 1110	
1.2	<u>1</u> 110 1110	
1.3b	0001 1100	
2.2	<u>1</u> 111 1100	
2.3b	0011 1000	
3.2	<u>0</u> 001 1000	
3.3a	0011 0001	
4.2	<u>0</u> 001 0001	
4.3a	0010 0011	
	0001 0011	



# Divide

## ◆ Signed Divides:

- Remember signs, make positive, complement quotient and remainder if necessary
- Let Dividend and Remainder have same sign and negate Quotient if Divisor sign & Dividend sign disagree,
- e.g.,  $-7 \div 2 = -3$ , remainder = -1  
 $-7 \div -2 = 3$ , remainder = -1
- Satisfy Dividend = Quotient x Divisor + Remainder

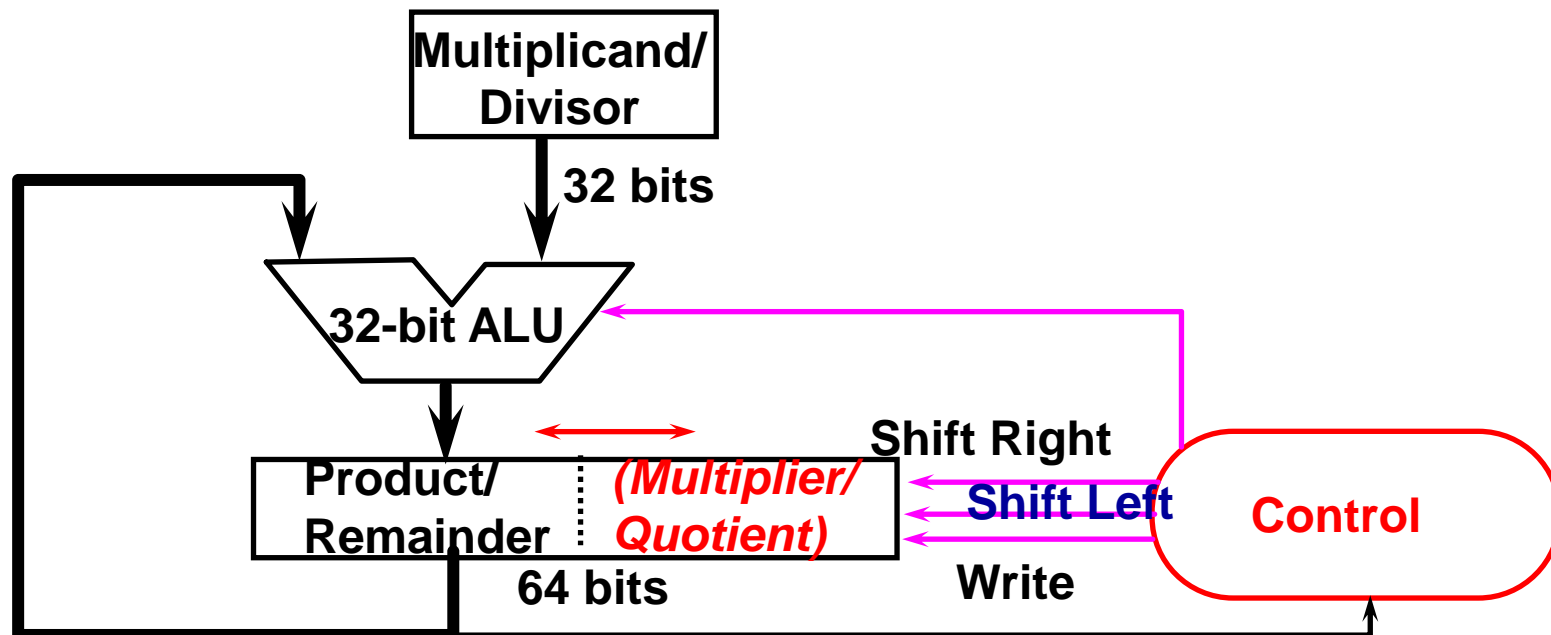
## ◆ Possible for quotient to be too large: if divide 64-bit integer by 1, quotient is 64 bits

# Observations: Multiply and Divide

- ◆ Same hardware as multiply: just need ALU to add or subtract, and 64-bit register to shift left or shift right
- ◆ Hi and Lo registers in MIPS combine to act as 64-bit register for multiply and divide

# Multiply/Divide Hardware

- ◆ 32-bit Multiplicand/Divisor register, 32-bit ALU, 64-bit Product/Remainder register, (0-bit Multiplier/Quotient register)





# Outline

- ◆ Addition and subtraction (Sec. 3.2)
- ◆ Constructing an arithmetic logic unit (Appendix C)
- ◆ Multiplication (Sec. 3.3, Appendix C)
- ◆ Division (Sec. 3.4)
- ◆ Floating point (Sec. 3.5)

# Floating-Point: Motivation

## ◆ What can be represented in N bits?

Unsigned	0	to	$2^n - 1$
2's Complement	$-2^{n-1}$	to	$2^{n-1} - 1$
1's Complement	$-2^{n-1} + 1$	to	$2^{n-1}$
Excess M	-M	to	$2^n - M - 1$

## ◆ But, what about ...

- very large numbers?

9,349,398,989,787,762,244,859,087,678

- very small number?

0.0000000000000000000000000045691

- rationals  $\frac{2}{3}$

- irrationals  $\sqrt{2}$

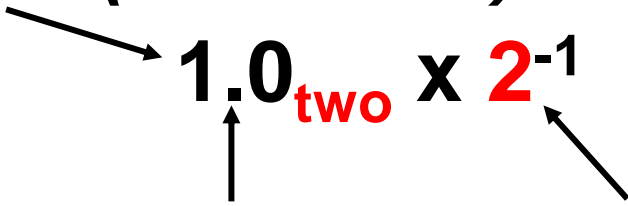
- transcendentals  $e, \pi$

# Scientific Notation: Binary

*Significand (Mantissa)*      *exponent*

$1.0_{\text{two}} \times 2^{-1}$


“*binary point*”      *radix (base)*

A diagram showing the components of the binary scientific notation expression 1.0\_two x 2^-1. An arrow points from the label 'Significand (Mantissa)' to the '1.0' part of the expression. Another arrow points from the label 'exponent' to the '2^-1' part. A third arrow points from the label '“binary point”' to the dot between '1' and '0'. A fourth arrow points from the label 'radix (base)' to the '2' in the exponent.

- ◆ Computer arithmetic that supports it is called floating point, because the binary point is not fixed, as it is for integers
- ◆ Normalized form: no leading 0s (exactly one digit to left of decimal point)
- ◆ Alternatives to represent  $1/1,000,000,000$ 
  - Normalized:  $1.0 \times 10^{-9}$
  - Not normalized:  $0.1 \times 10^{-8}$ ,  $10.0 \times 10^{-10}$

# FP Representation

- ◆ Normal format:  $1.\text{xxxxxxxxxx}_{\text{two}} \times 2^{\text{yyyy}_{\text{two}}}$
- ◆ Want to put it into multiple words: 32 bits for *single-precision* and 64 bits for *double-precision*
- ◆ A simple single-precision representation:



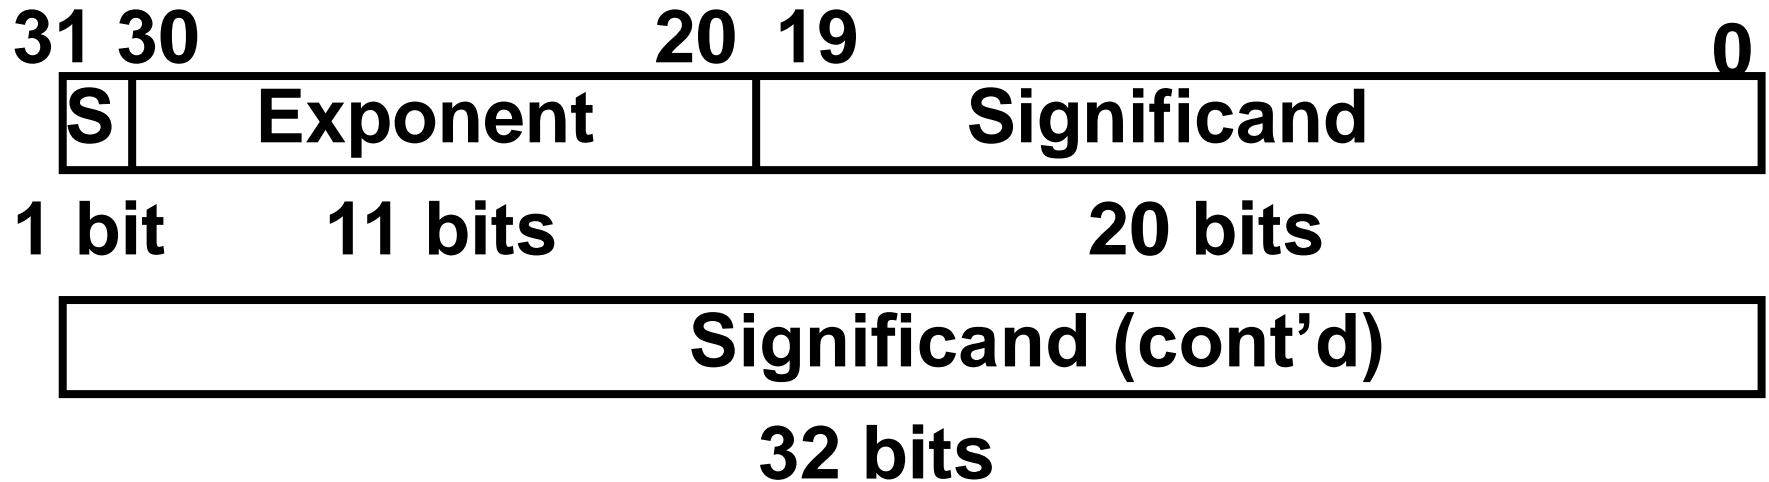
**S** represents sign

**Exponent** represents y's

**Significand** represents x's

# Double Precision Representation

- ◆ Next multiple of word size (64 bits)



- ◆ Double precision (vs. single precision)
  - But primary advantage is greater accuracy due to larger significand

# IEEE 754 Standard (1/4)

- ◆ Regarding single precision, DP similar
- ◆ Sign bit:
  - 1 means negative
  - 0 means positive
- ◆ Significand:
  - To pack more bits, leading 1 implicit for normalized numbers
  - 1 + 23 bits single, 1 + 52 bits double
  - always true:  $0 < \text{Significand} < 1$   
(for normalized numbers)
- ◆ Note: 0 has no leading 1, so reserve exponent value 0 just for number 0

# IEEE 754 Standard (2/4)

## ◆ Exponent:

- Need to represent positive and negative exponents
- Also want to compare FP numbers as if they were integers, to help in value comparisons
- If use 2's complement to represent?  
e.g.,  $1.0 \times 2^{-1}$  versus  $1.0 \times 2^{+1}$  (1/2 versus 2)

1/2

0	1111 1111	000 0000 0000 0000 0000 0000
---	-----------	------------------------------

2

0	0000 0001	000 0000 0000 0000 0000 0000
---	-----------	------------------------------

*If we use integer comparison for these two words,  
we will conclude that  $1/2 > 2$ !!!*

# Biased (Excess) Notation

## ◆ Biased 7

0000	-7
0001	-6
0010	-5
0011	-4
0100	-3
0101	-2
0110	-1
0111	0
1000	1
1001	2
1010	3
1011	4
1100	5
1101	6
1110	7
1111	8




# IEEE 754 Standard (3/4)

- ◆ Instead, let notation 0000 0000 be most negative, and 1111 1111 most positive
- ◆ Called **biased notation**, where bias is the number subtracted to get the real number
  - IEEE 754 uses bias of 127 for single precision:  
Subtract 127 from Exponent field to get actual value for exponent
  - 1023 is bias for double precision

1/2	0	0111 1110	000 0000 0000 0000 0000 0000
2	0	1000 0000	000 0000 0000 0000 0000 0000

# IEEE 754 Standard (4/4)

- ◆ Summary (single precision):



$$(-1)^S \times (1.\text{Significand}) \times 2^{(\text{Exponent}-127)}$$

- ◆ Double precision identical, except with exponent bias of 1023

# Example: FP to Decimal

0	0110 1000	101 0101 0100 0011 0100 0010
---	-----------	------------------------------

- ◆ Sign: 0 => positive
- ◆ Exponent:
  - $0110\ 1000_{\text{two}} = 104_{\text{ten}}$
  - Bias adjustment:  $104 - 127 = -23$
- ◆ Significand:
  - $1 + 2^{-1} + 2^{-3} + 2^{-5} + 2^{-7} + 2^{-9} + 2^{-14} + 2^{-15} + 2^{-17} + 2^{-22}$   
 $= 1.0 + 0.666115$
- ◆ Represents:  $1.666115_{\text{ten}} \times 2^{-23} \approx 1.986 \times 10^{-7}$

# Example 1: Decimal to FP

- ◆ Number = - 0.75
  - = -  $0.11_{\text{two}} \times 2^0$  (scientific notation)
  - = -  $1.1_{\text{two}} \times 2^{-1}$  (normalized scientific notation)
- ◆ Sign: negative => 1
- ◆ Exponent:
  - Bias adjustment:  $-1 + 127 = 126$
  - $126_{\text{ten}} = 0111\ 1110_{\text{two}}$

1	0111 1110	100 0000 0000 0000 0000
---	-----------	-------------------------

## Example 2: Decimal to FP

◆ A more difficult case: representing  $1/3$ ?

$$= 0.33333..._{10} = 0.0101010101..._2 \times 2^0$$

$$= 1.0101010101..._2 \times 2^{-2}$$

- Sign: 0
- Exponent =  $-2 + 127 = 125_{10} = 01111101_2$
- Significand = 0101010101...

0	0111 1101	0101 0101 0101 0101 0101 010
---	-----------	------------------------------

# Single-Precision Range

- ◆ Exponents 00000000 and 11111111 reserved
- ◆ Smallest value
  - Exponent: 00000001  
 $\Rightarrow$  actual exponent =  $1 - 127 = -126$
  - Fraction: 000...00  $\Rightarrow$  significand = 1.0
  - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- ◆ Largest value
  - exponent: 11111110  
 $\Rightarrow$  actual exponent =  $254 - 127 = +127$
  - Fraction: 111...11  $\Rightarrow$  significand  $\approx 2.0$
  - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

# Double-Precision Range

- ◆ Exponents 0000...00 and 1111...11 reserved
- ◆ Smallest value
  - Exponent: 000000000001  
 $\Rightarrow$  actual exponent =  $1 - 1023 = -1022$
  - Fraction: 000...00  $\Rightarrow$  significand = 1.0
  - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- ◆ Largest value
  - Exponent: 111111111110  
 $\Rightarrow$  actual exponent =  $2046 - 1023 = +1023$
  - Fraction: 111...11  $\Rightarrow$  significand  $\approx 2.0$
  - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$

# Floating-Point Precision

## ◆ Relative precision

- all fraction bits are significant
- Single: approx  $2^{-23}$ 
  - Equivalent to  $23 \times \log_{10} 2 \approx 23 \times 0.3 \approx 6$  decimal digits of precision
- Double: approx  $2^{-52}$ 
  - Equivalent to  $52 \times \log_{10} 2 \approx 52 \times 0.3 \approx 16$  decimal digits of precision



# Zero and Special Numbers

- ◆ What have we defined so far? (single precision)

<u>Exponent</u>	<u>Significand</u>	<u>Object</u>
0	0	<u>???</u>
0	nonzero	<u>???</u>
1-254	anything	+/- floating-point
255	0	<u>???</u>
255	nonzero	<u>???</u>

# Representation for 0

## ◆ Represent 0?

- exponent all zeroes
- significand all zeroes too
- What about sign?
- +0: 0 00000000 000000000000000000000000
- -0: 1 00000000 000000000000000000000000

## ◆ Why two zeroes?

- Helps in some limit comparisons

# Special Numbers

## ◆ What have we defined so far? (single precision)

<u>Exponent</u>	<u>Significand</u>	<u>Object</u>
0	0	0
0	nonzero	<u>???</u>
1-254	anything	+/- floating-point
255	0	<u>???</u>
255	nonzero	<u>???</u>

## ◆ Range:

$$1.0 \times 2^{-126} \approx 1.8 \times 10^{-38}$$

What if result too small? ( $>0$ ,  $< 1.8 \times 10^{-38} \Rightarrow$  Underflow!)

$$(2 - 2^{-23}) \times 2^{127} \approx 3.4 \times 10^{38}$$

What if result too large? ( $> 3.4 \times 10^{38} \Rightarrow$  Overflow!)

# Gradual Underflow

- ◆ Represent denormalized numbers (denorms)
  - Exponent : all zeroes
  - Significand : non-zeroes
  - Allow a number to degrade in significance until it become 0 (gradual underflow)
- The smallest normalized number
  - $1.0000\ 0000\ 0000\ 0000\ 0000\ 0000 \times 2^{-126}$
- The smallest de-normalized number
  - $0.0000\ 0000\ 0000\ 0000\ 0000\ 0001 \times 2^{-126}$

# Special Numbers

- ◆ What have we defined so far? (single precision)

Exponent

0

0

1-254

255

255

Significand

0

nonzero

anything

0

nonzero

Object

0

denorm

+/- floating-point

???

???

# Representation for +/- Infinity

- ◆ In FP, divide by zero should produce +/- infinity, not overflow
- ◆ Why?
  - OK to do further computations with infinity, e.g.,  $X/0 > Y$  may be a valid comparison
- ◆ IEEE 754 represents +/- infinity
  - Most positive exponent reserved for infinity
  - Significands all zeroes

S	1111 1111	0000 0000 0000 0000 0000 000
---	-----------	------------------------------

# Special Numbers (cont'd)

- ◆ What have we defined so far? (single-precision)

Exponent

0

0

1-254

255

255

Significand

0

nonzero

anything

0

nonzero

Object

0

denom

+/- fl. pt. #

+/- infinity

???

# Representation for Not a Number

- ◆ What do I get if I calculate  $\text{sqrt}(-4.0)$  or  $0/0$ ?
  - If infinity is not an error, these should not be either
  - They are called *Not a Number* (NaN)
  - Exponent = 255, Significand nonzero
- ◆ Why is this useful?
  - Hope NaNs help with debugging?
  - They contaminate:  $\text{op}(\text{NaN}, X) = \text{NaN}$
  - OK if calculate but don't use it



# Special Numbers (cont'd)

- ◆ What have we defined so far? (single-precision)

Exponent

0

0

1-254

255

255

Significand

0

nonzero

anything

0

nonzero

Object

0

denom

+/- fl. pt. #

+/- infinity

NaN

# Floating-Point Addition

## Basic addition algorithm:

(1) Align binary point :compute  $Y_e - X_e$

- ◆ right shift the smaller number, say  $X_m$ , that many positions to form  $X_m \times 2^{X_e - Y_e}$

(2) Add mantissa: compute  $X_m \times 2^{X_e - Y_e} + Y_m$

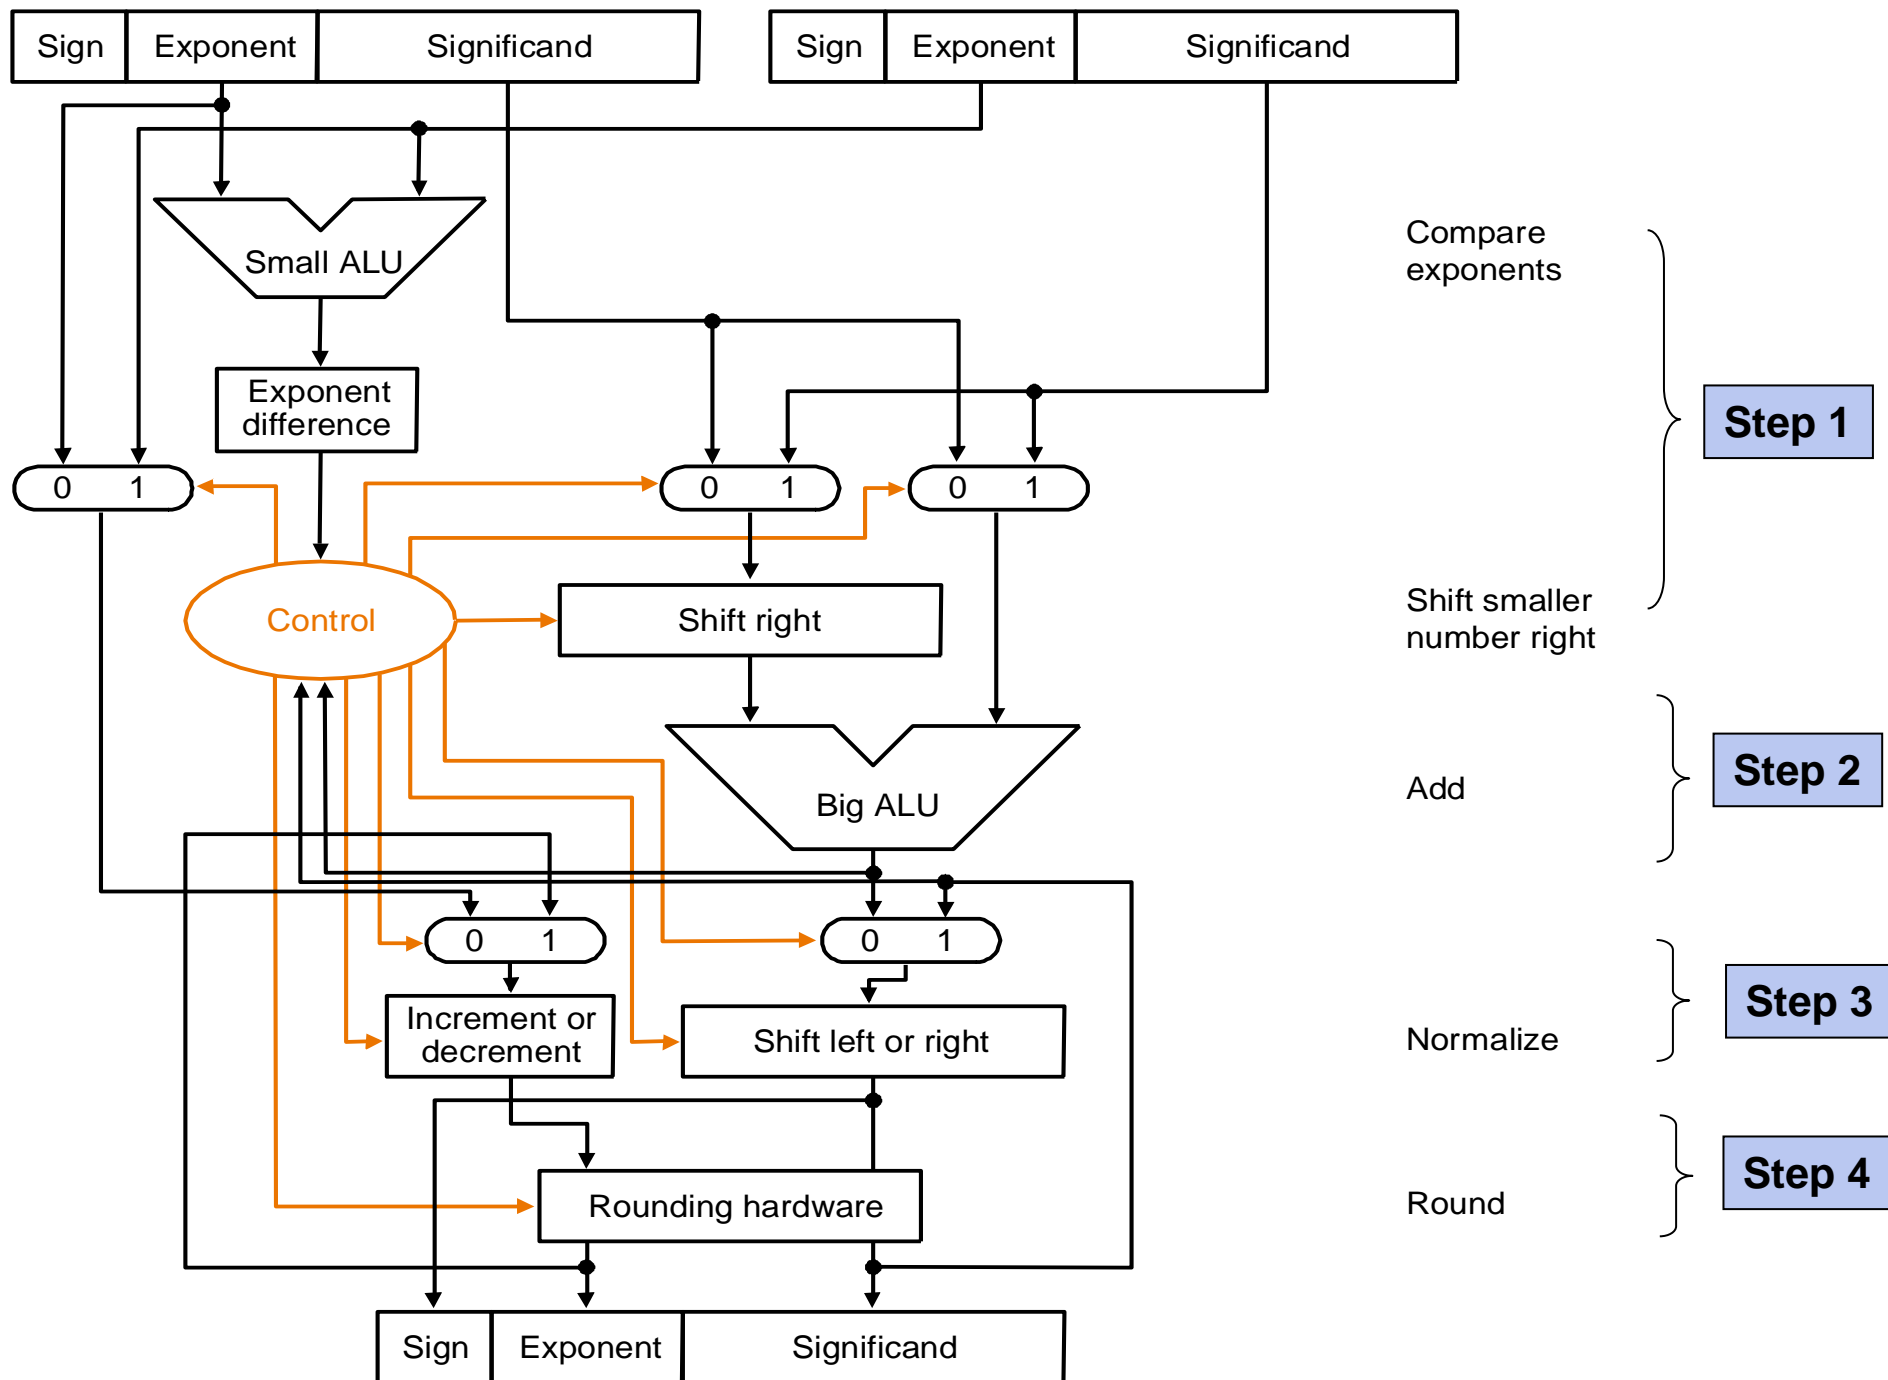
(3) Normalization & check for over/underflow if necessary:

- left shift result, decrement result exponent
- right shift result, increment result exponent
- check overflow or underflow during the shift

(4) Round the mantissa and renormalize if necessary

# Floating-Point Addition Example

- ◆ Now consider a 4-digit binary example
  - $1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2}$  (0.5 + -0.4375)
- ◆ 1. Align binary points
  - Shift number with smaller exponent
  - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$
- ◆ 2. Add mantissa
  - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$
- ◆ 3. Normalize result & check for over/underflow
  - $1.000_2 \times 2^{-4}$ , with no over/underflow
- ◆ 4. Round and renormalize if necessary
  - $1.000_2 \times 2^{-4}$  (no change) = 0.0625



# FP Adder Hardware

- ◆ Much more complex than integer adder
- ◆ Doing it in one clock cycle would take too long
  - Much longer than integer operations
  - Slower clock would penalize all instructions
- ◆ FP adder usually takes several cycles
  - Can be pipelined

# Floating-Point Multiplication

## Basic multiplication algorithm

- (1) Add exponents of operands to get exponent of product  
doubly biased exponent must be corrected:

$$X_e = 7$$

$$Y_e = -3 \quad X_e = 1111 \quad = 15 \quad = 7 + 8$$

$$\text{Excess } 8 \quad Y_e = 0101 \quad = 5 \quad = -3 + 8$$

$$10100 \quad 20 \quad 4 + 8 + 8$$

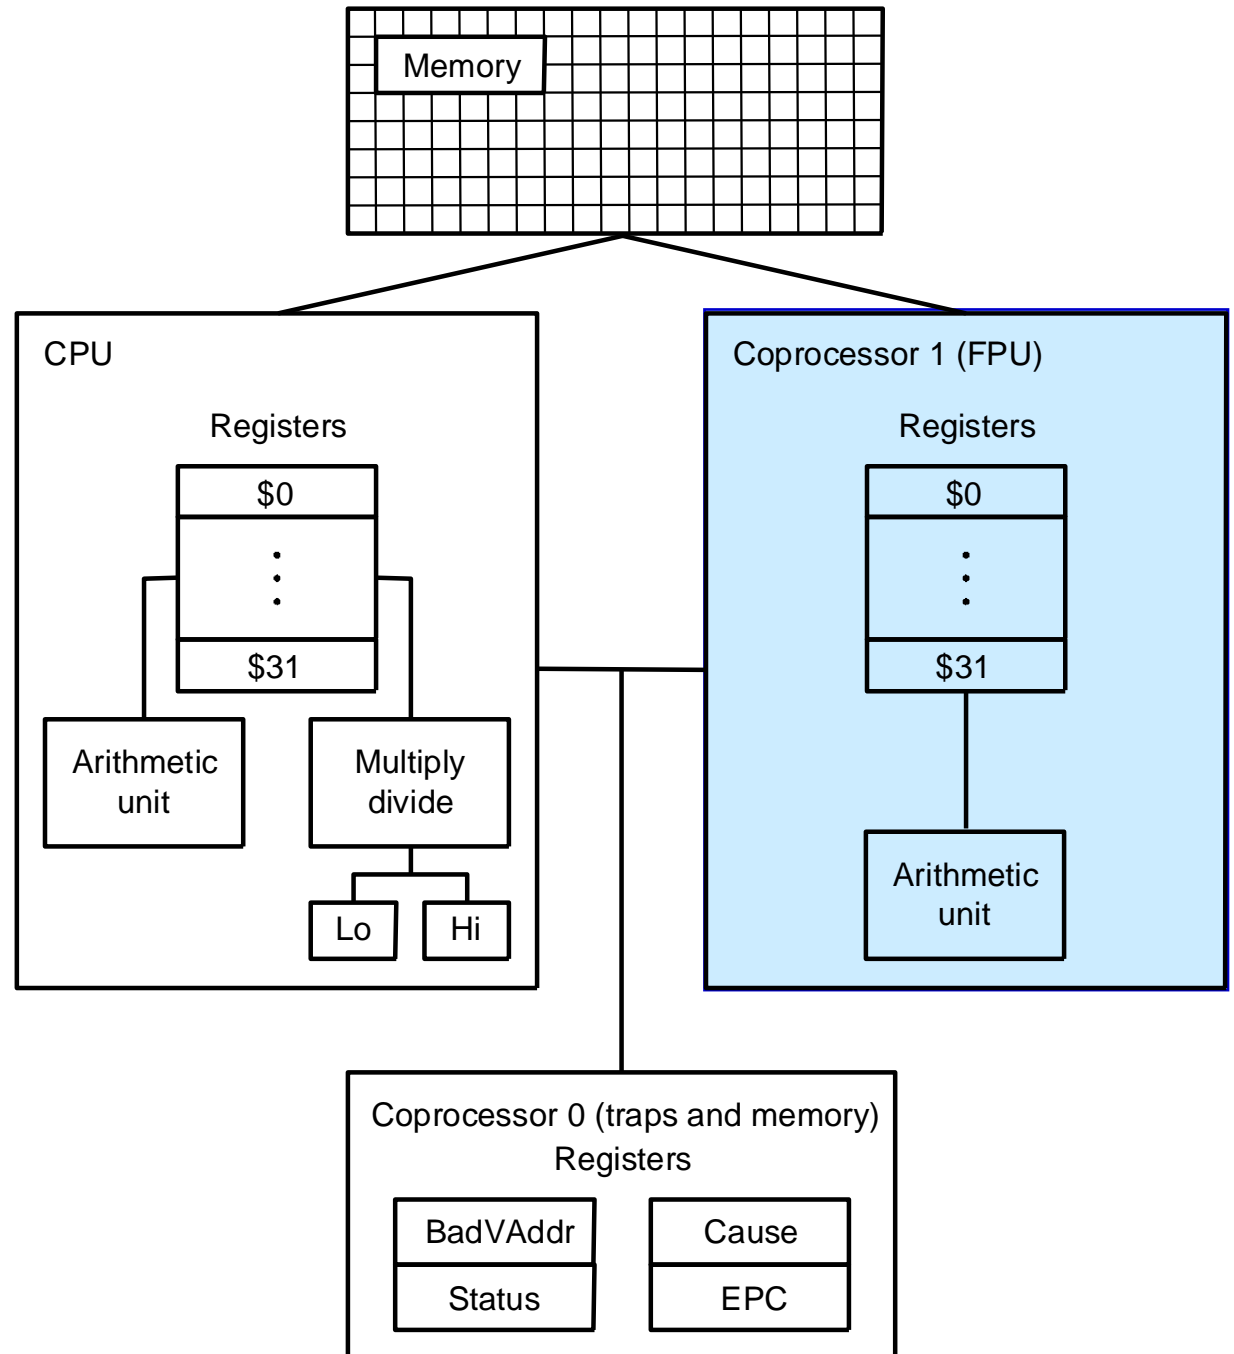
need extra subtraction step of the bias amount

- (2) Multiplication of operand mantissa  
(3) Normalize the product & check overflow or underflow during the shift  
(4) Round the mantissa and renormalize if necessary  
(5) Set the sign of product

# Floating-Point Multiplication Example

- ◆ Now consider a 4-digit binary example
  - $1.000_2 \times 2^{-1} \times -1.110_2 \times 2^{-2}$  ( $0.5 \times -0.4375$ )
- ◆ 1. Add exponents
  - Unbiased:  $-1 + -2 = -3$
  - Biased:  $(-1 + 127) + (-2 + 127) = -3 + 254 - 127 = -3 + 127$
- ◆ 2. Multiply operand mantissa
  - $1.000_2 \times 1.110_2 = 1.1102 \Rightarrow 1.110_2 \times 2^{-3}$
- ◆ 3. Normalize result & check for over/underflow
  - $1.110_2 \times 2^{-3}$  (no change) with no over/underflow
- ◆ 4. Round and renormalize if necessary
  - $1.110_2 \times 2^{-3}$  (no change)
- ◆ 5. Determine sign:
  - $-1.110_2 \times 2^{-3} = -0.21875$

# MIPS R2000 Organization





# MIPS Floating Point

- ◆ **Separate floating point instructions:**
  - Single precision: `add.s, sub.s, mul.s, div.s`
  - Double precision: `add.d, sub.d, mul.d, div.d`
- ◆ **FP part of the processor:**
  - contains 32 32-bit registers: `$f0, $f1, ...`
  - most registers specified in `.s` and `.d` instruction refer to this set
  - Double precision: by convention, even/odd pair contain one DP FP number: `$f0/$f1, $f2/$f3`
  - separate load and store: `lwc1` and `swc1`
  - Instructions to move data between main processor and coprocessors:
    - `mfc0, mtc0, mfc1, mtc1, etc.`

# Interpretation of Data

## The BIG Picture

- ◆ Bits have no inherent meaning
  - Interpretation depends on the instructions applied
- ◆ Computer representations of numbers
  - Finite range and precision
  - Need to account for this in programs

# Associativity

## ◆ Floating Point add, subtract associative ?

		$(x+y)+z$	$x+(y+z)$
x	-1.50E+38	0.00E+00	-1.50E+38
y	1.50E+38		1.50E+38
z	1.0	1.0	
		1.00E+00	0.00E+00

- ◆ Therefore, Floating Point add, subtract are not associative!
  - Why? FP result approximates real result!
  - This example:  $1.5 \times 10^{38}$  is so much larger than 1.0 that  $1.5 \times 10^{38} + 1.0$  in floating point representation is still  $1.5 \times 10^{38}$

# Associativity in Parallel Programming

- ◆ Parallel programs may interleave operations in unexpected orders
  - Assumptions of associativity may fail
- ◆ Need to validate parallel programs under varying degrees of parallelism

# x86 FP Architecture

- ◆ Originally based on 8087 FP coprocessor
  - 8 × 80-bit extended-precision registers
  - Used as a push-down stack
  - Registers indexed from TOS: ST(0), ST(1), ...
- ◆ FP values are 32-bit or 64 in memory
  - Converted on load/store of memory operand
  - Integer operands can also be converted on load/store
- ◆ Very difficult to generate and optimize code
  - Result: poor FP performance

# x86 FP Instructions

Data transfer	Arithmetic	Compare	Transcendental
<b>F</b> <b>I</b> LD mem/ST(i) <b>F</b> <b>I</b> ST <b>P</b> mem/ST(i) FLDPI FLD1 FLDZ	<b>F</b> <b>I</b> ADD <b>P</b> mem/ST(i) <b>F</b> <b>I</b> SUB <b>R</b> <b>P</b> mem/ST(i) <b>F</b> <b>I</b> MUL <b>P</b> mem/ST(i) <b>F</b> <b>I</b> DIV <b>R</b> <b>P</b> mem/ST(i) FSQRT FABS FRNDINT	<b>F</b> <b>I</b> COMP <b>F</b> <b>I</b> UCOMP FSTSW AX/mem	FPATAN F2XMI FCOS FPTAN FPREM FPSIN FYL2X

## ◆ Optional variations

- **I**: integer operand
- **P**: pop operand from stack
- **R**: reverse operand order
- But not all combinations allowed

# Streaming SIMD Extension 2 (SSE2)

- ◆ Adds  $4 \times 128$ -bit registers
  - Extended to 8 registers in AMD64/EM64T
- ◆ Can be used for multiple FP operands
  - $2 \times 64$ -bit double precision
  - $4 \times 32$ -bit double precision
  - Instructions operate on them simultaneously
    - Single-Instruction Multiple-Data

# Right Shift and Division

- ◆ Left shift by  $i$  places multiplies an integer by  $2^i$
- ◆ Right shift divides by  $2^i$ ?
  - Only for unsigned integers
- ◆ For signed integers
  - Arithmetic right shift: replicate the sign bit
  - e.g.,  $-5 / 4$ 
    - $11111011_2 \gg 2 = 11111110_2 = -2$
    - Rounds toward  $-\infty$
  - c.f.  $11111011_2 \ggg 2 = 00111110_2 = +62$



# Who Cares About FP Accuracy?

- ◆ Important for scientific code
  - But for everyday consumer use?
    - “My bank balance is out by 0.0002¢!” ☹
- ◆ The Intel Pentium FDIV bug
  - The market expects accuracy
  - See Colwell, *The Pentium Chronicles*

# Concluding Remarks

- ◆ ISAs support arithmetic
  - Signed and unsigned integers
  - Floating-point approximation to reals
- ◆ Bounded range and precision
  - Operations can overflow and underflow
- ◆ MIPS ISA
  - Core instructions: 54 most frequently used
    - 100% of SPECINT, 97% of SPECFP
  - Other instructions: less frequent