

# CS4100: 計算機結構

## Memory Hierarchy

---

國立清華大學資訊工程學系  
一零零學年度第二學期

# Outline

- ◆ **Memory hierarchy**
- ◆ **The basics of caches**
- ◆ **Measuring and improving cache performance**
- ◆ **Virtual memory**
- ◆ **A common framework for memory hierarchy**
- ◆ **Using a Finite State Machine to Control a Simple Cache**
- ◆ **Parallelism and Memory Hierarchies: Cache Coherence**

# Memory Technology

## ◆ Random access:

- Access time same for all locations
- **SRAM**: *Static Random Access Memory*
  - Low density, high power, expensive, fast
  - Static: content will last (forever until lose power)
  - Address not divided
  - Use for caches
- **DRAM**: *Dynamic Random Access Memory*
  - High density, low power, cheap, slow
  - Dynamic: need to be refreshed regularly
  - Addresses in 2 halves (memory as a 2D matrix):
    - ✧ RAS/CAS (Row/Column Access Strobe)
  - Use for main memory

## ◆ Magnetic disk


# Comparisons of Various Technologies

Memory technology	Typical access time	\$ per GB in 2008
SRAM	0.5 – 2.5 ns	\$2000 – \$5,000
DRAM	50 – 70 ns	\$20 – \$75
Magnetic disk	5,000,000 – 20,000,000 ns	\$0.20 – \$2

**Ideal memory**

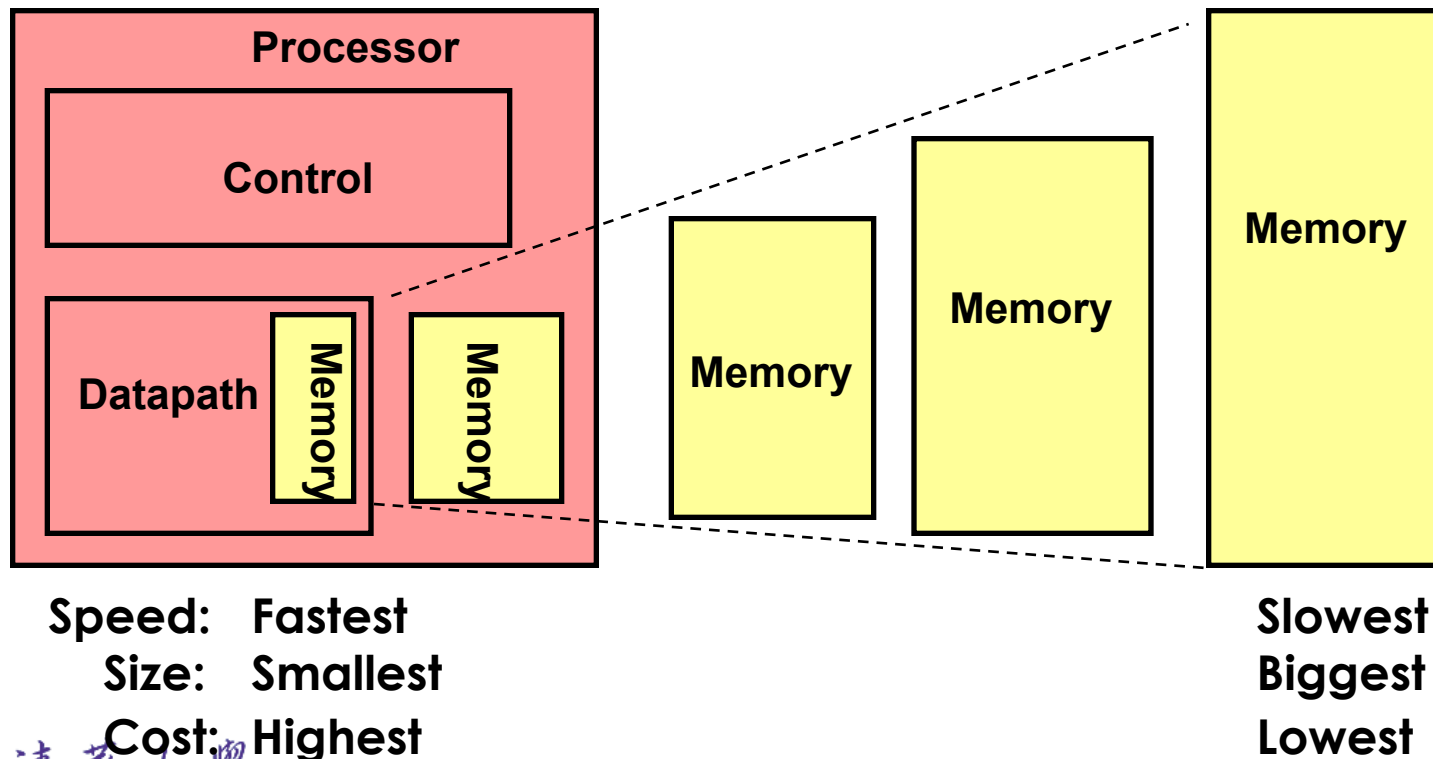
- ◆ Access time of SRAM
- ◆ Capacity and cost/GB of disk

# Processor Memory Latency Gap




# Solution: Memory Hierarchy

- ◆ An Illusion of a large, fast, cheap memory
  - Fact: Large memories slow, fast memories small
  - How to achieve: hierarchy, parallelism
- ◆ An expanded view of memory system:



# Memory Hierarchy: Principle

- ◆ At any given time, data is copied between only two adjacent levels:
  - **Upper level:** the one closer to the processor
    - Smaller, faster, uses more expensive technology
  - **Lower level:** the one away from the processor
    - Bigger, slower, uses less expensive technology
- ◆ **Block:** basic unit of information transfer
  - Minimum unit of information that can either be present or not present in a level of the hierarchy




# Why Hierarchy Works?

## ◆ Principle of Locality:

- Program access a relatively small portion of the address space at any instant of time
- 90/10 rule: 10% of code executed 90% of time

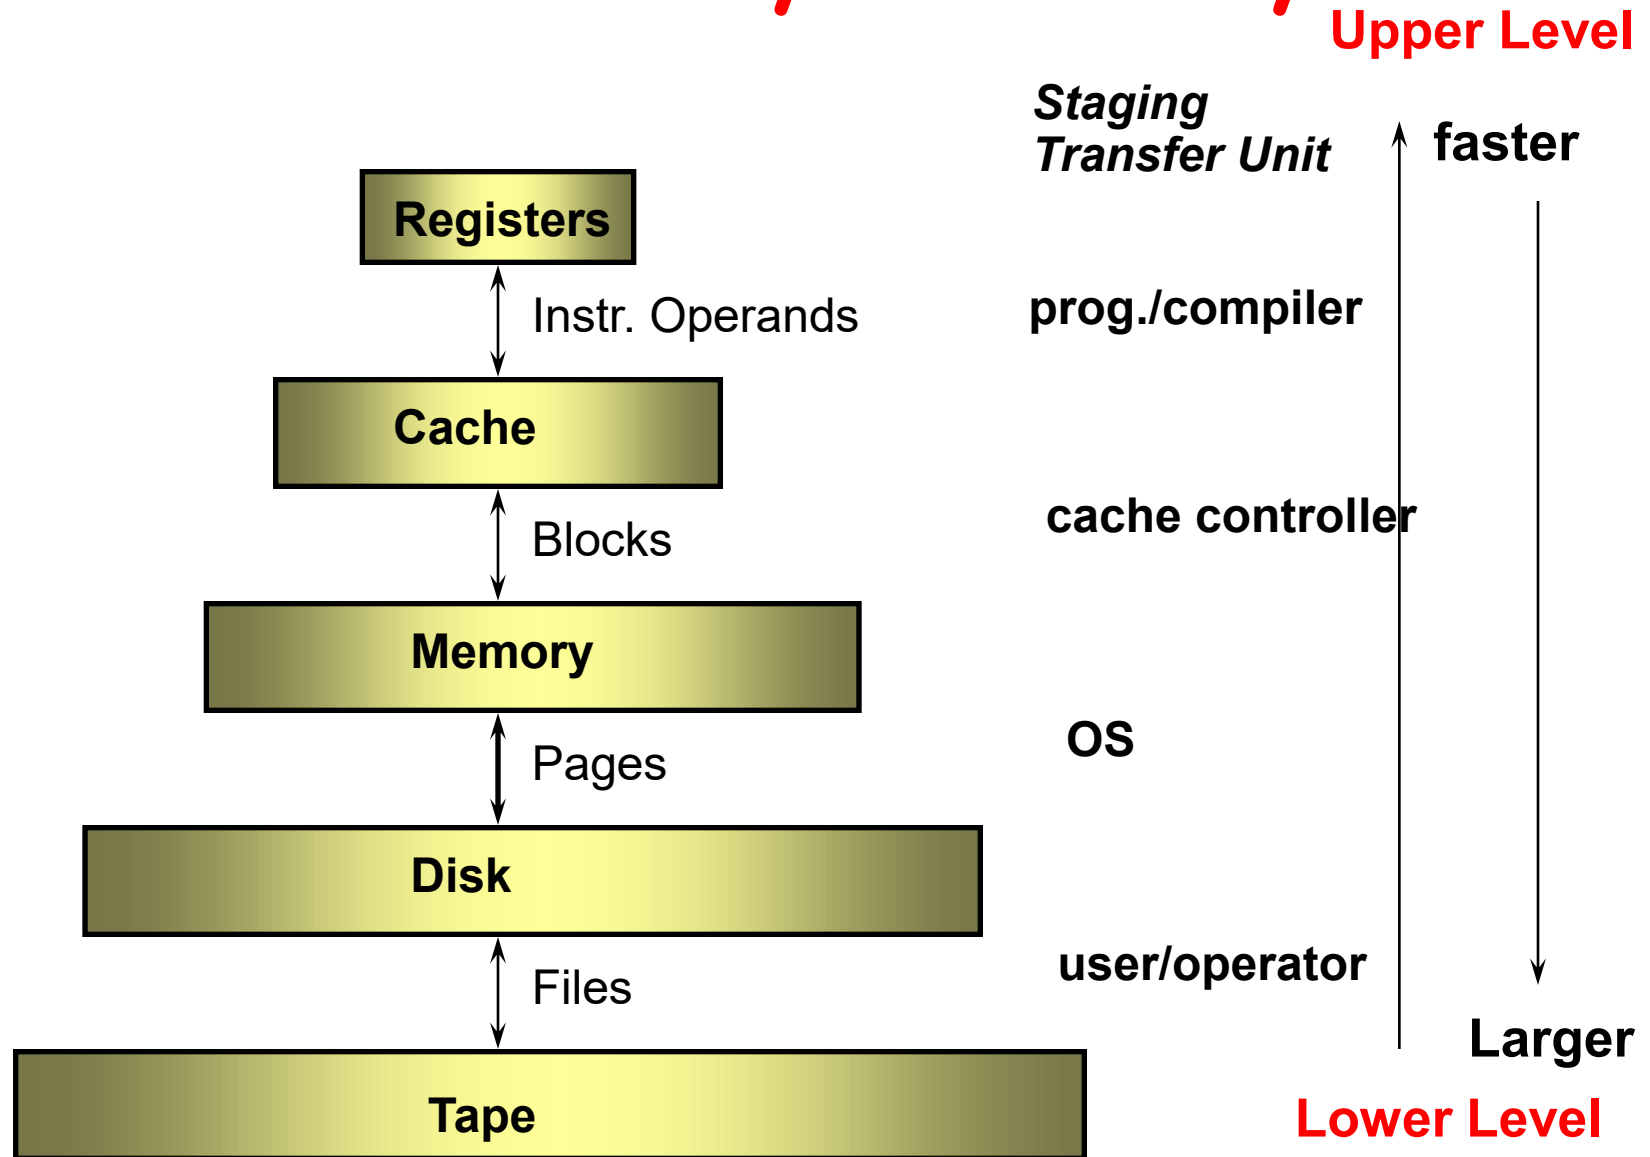
## ◆ Two types of locality:

- **Temporal locality**: if an item is referenced, it will tend to be referenced again soon
- **Spatial locality**: if an item is referenced, items whose addresses are close by tend to be referenced soon





# Levels of Memory Hierarchy




# How Is the Hierarchy Managed?

- ◆ Registers  $\leftrightarrow$  Memory
  - by compiler (programmer?)
- ◆ cache  $\leftrightarrow$  memory
  - by the hardware
- ◆ memory  $\leftrightarrow$  disks
  - by the hardware and operating system (virtual memory)
  - by the programmer (files)

# Memory Hierarchy: Terminology

- ◆ **Hit**: data appears in upper level (Block X)
  - **Hit rate**: fraction of memory access found in the upper level
  - **Hit time**: time to access the upper level
    - RAM access time + Time to determine hit/miss
- ◆ **Miss**: data needs to be retrieved from a block in the lower level (Block Y)
  - **Miss Rate** =  $1 - (\text{Hit Rate})$
  - **Miss Penalty**: time to replace a block in the upper level + time to deliver the block to the processor (latency + transmit time)
- ◆ **Hit Time**  $\ll$  **Miss Penalty**



# 4 Questions for Hierarchy Design

**Q1: Where can a block be placed in the upper level?**

**=> *block placement***

**Q2: How is a block found if it is in the upper level?**

**=> *block finding***

**Q3: Which block should be replaced on a miss?**

**=> *block replacement***

**Q4: What happens on a write?**

**=> *write strategy***

# Summary of Memory Hierarchy

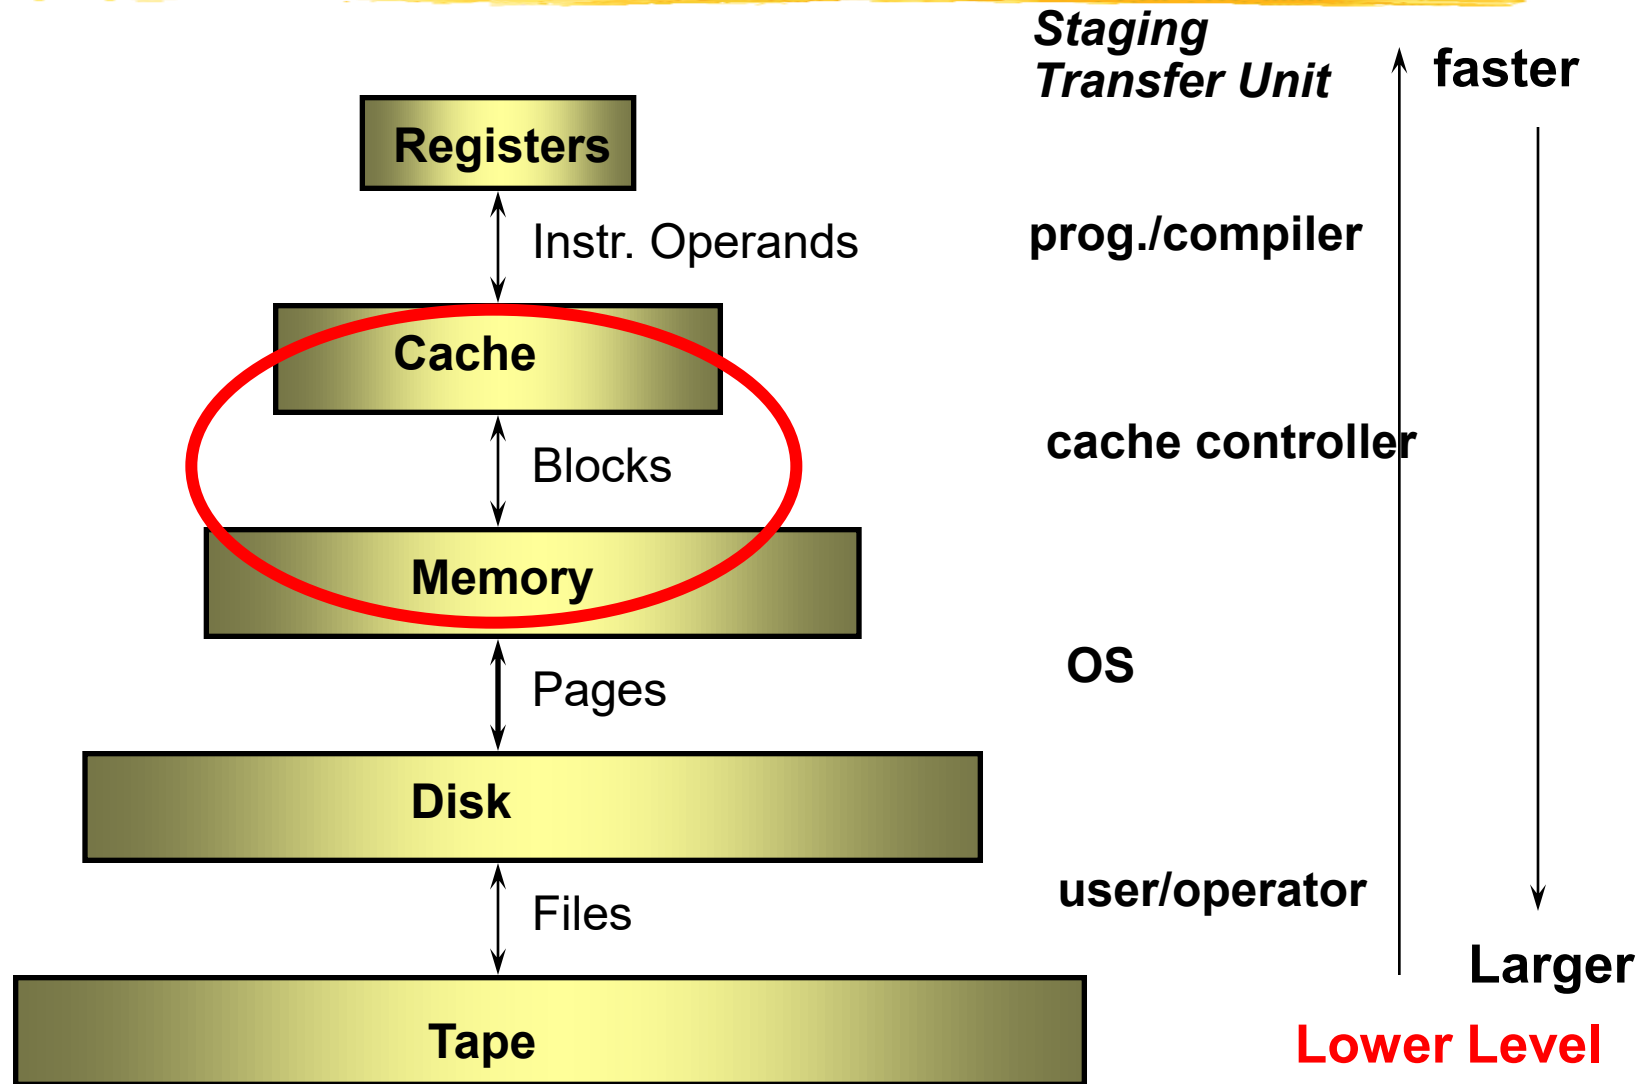
- ◆ Two different types of locality:
  - Temporal Locality (Locality in Time)
  - Spatial Locality (Locality in Space)
- ◆ Using the principle of locality:
  - Present the user with as much memory as is available in the cheapest technology.
  - Provide access at the speed offered by the fastest technology.
- ◆ DRAM is slow but cheap and dense:
  - Good for presenting users with a BIG memory system
- ◆ SRAM is fast but expensive, not very dense:
  - Good choice for providing users FAST accesses

# Outline

- ◆ Memory hierarchy
- ◆ The basics of caches
- ◆ Measuring and improving cache performance
- ◆ Virtual memory
- ◆ A common framework for memory hierarchy
- ◆ Using a Finite State Machine to Control a Simple Cache
- ◆ Parallelism and Memory Hierarchies: Cache Coherence

# Levels of Memory Hierarchy

Upper Level



Lower Level

Computer Architecture

# Cache Memory

- ◆ Cache memory
  - The level of the memory hierarchy closest to the CPU
- ◆ Given accesses  $X_1, \dots, X_{n-1}, X_n$

$X_4$
$X_1$
$X_{n-2}$
$X_{n-1}$
$X_2$
$X_3$

$X_4$
$X_1$
$X_{n-2}$
$X_{n-1}$
$X_2$
$X_n$
$X_3$

- ◆ How do we know if the data is present?
- ◆ Where do we look?


a. Before the reference to  $X_n$

b. After the reference to  $X_n$



# Basics of Cache

- ◆ Our first example: *direct-mapped cache*
- ◆ Block Placement :
  - For each item of data at the lower level, there is exactly one location in cache where it might be
  - Address mapping: modulo number of blocks



# Tags and Valid Bits: Block Finding

- ◆ How do we know which particular block is stored in a cache location?
  - Store block address as well as the data
  - Actually, only need the high-order bits
  - Called the **tag**
- ◆ What if there is no data in a location?
  - Valid bit: 1 = present, 0 = not present
  - Initially 0

# Cache Example

- ◆ 8-blocks, 1 word/block, direct mapped
- ◆ Initial state

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	N		
111	N		

# Cache Example

Word addr	Binary addr	Hit/miss	Cache block
22	10 110	Miss	110

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

# Cache Example

Word addr	Binary addr	Hit/miss	Cache block
26	11 010	Miss	010

Index	V	Tag	Data
000	N		
001	N		
010	Y	11	Mem[11010]
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

# Cache Example

Word addr	Binary addr	Hit/miss	Cache block
22	10 110	Hit	110
26	11 010	Hit	010

Index	V	Tag	Data
000	N		
001	N		
010	Y	11	Mem[11010]
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

# Cache Example

Word addr	Binary addr	Hit/miss	Cache block
16	10 000	Miss	000
3	00 011	Miss	011
16	10 000	Hit	000

Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	11	Mem[11010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

# Cache Example


Word addr	Binary addr	Hit/miss	Cache block
18	10 010	Miss	010

Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	11->10	Mem[10010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		




# Address Subdivision

- ◆ 1K words, 1-word block:
  - Cache index: lower 10 bits
  - Cache tag: upper 20 bits
  - Valid bit (When start up, valid is 0)



# Example: Larger Block Size

- ◆ 64 blocks, 16 bytes/block
  - To what block number does address 1200 map?
- ◆ Block address =  $\lfloor 1200/16 \rfloor = 75$
- ◆ Block number = 75 modulo 64 = 11
- ◆  $1200 = 10010110000_2 / 10000_2 \Rightarrow 1001011_2$
- ◆  $1001011_2 \Rightarrow 001011_2$




# Block Size Considerations

- ◆ Larger blocks should reduce miss rate
  - Due to spatial locality
- ◆ But in a fixed-sized cache
  - Larger blocks  $\Rightarrow$  fewer of them
    - More competition  $\Rightarrow$  increased miss rate
- ◆ Larger miss penalty
  - Larger blocks  $\Rightarrow$  pollution
  - Can override benefit of reduced miss rate
  - Early restart and critical-word-first can help

# Block Size on Performance

- ◆ Increase block size tends to decrease miss rate



# Cache Misses

## Read Hit and Miss:


- ◆ On cache hit, CPU proceeds normally
- ◆ On cache miss
  - Stall the CPU pipeline
  - Fetch block from next level of hierarchy
  - Instruction cache miss
    - Restart instruction fetch
  - Data cache miss
    - Complete data access

# Write-Through

## Write Hit:

- ◆ On data-write hit, could just update the block in cache
  - But then cache and memory would be inconsistent
- ◆ Write through: also update memory
- ◆ But makes writes take longer
  - e.g., if base CPI = 1, 10% of instructions are stores, write to memory takes 100 cycles
    - Effective CPI =  $1 + 0.1 \times 100 = 11$
- ◆ Solution: write buffer
  - Holds data waiting to be written to memory
  - CPU continues immediately
    - Only stalls on write if write buffer is already full

# Avoid Waiting for Memory in Write Through



- ◆ Use a *write buffer* (WB):
  - Processor: writes data into cache and WB
  - Memory controller: write WB data to memory
- ◆ Write buffer is just a FIFO:
  - Typical number of entries: 4
- ◆ Memory system designer's nightmare:
  - Store frequency  $> 1 / \text{DRAM write cycle}$
  - Write buffer saturation  $\Rightarrow$  CPU stalled

# Write-Back

- ◆ **Alternative: On data-write hit, just update the block in cache**
  - Keep track of whether each block is dirty
- ◆ **When a dirty block is replaced**
  - Write it back to memory
  - Can use a write buffer to allow replacing block to be read first



# Write Allocation

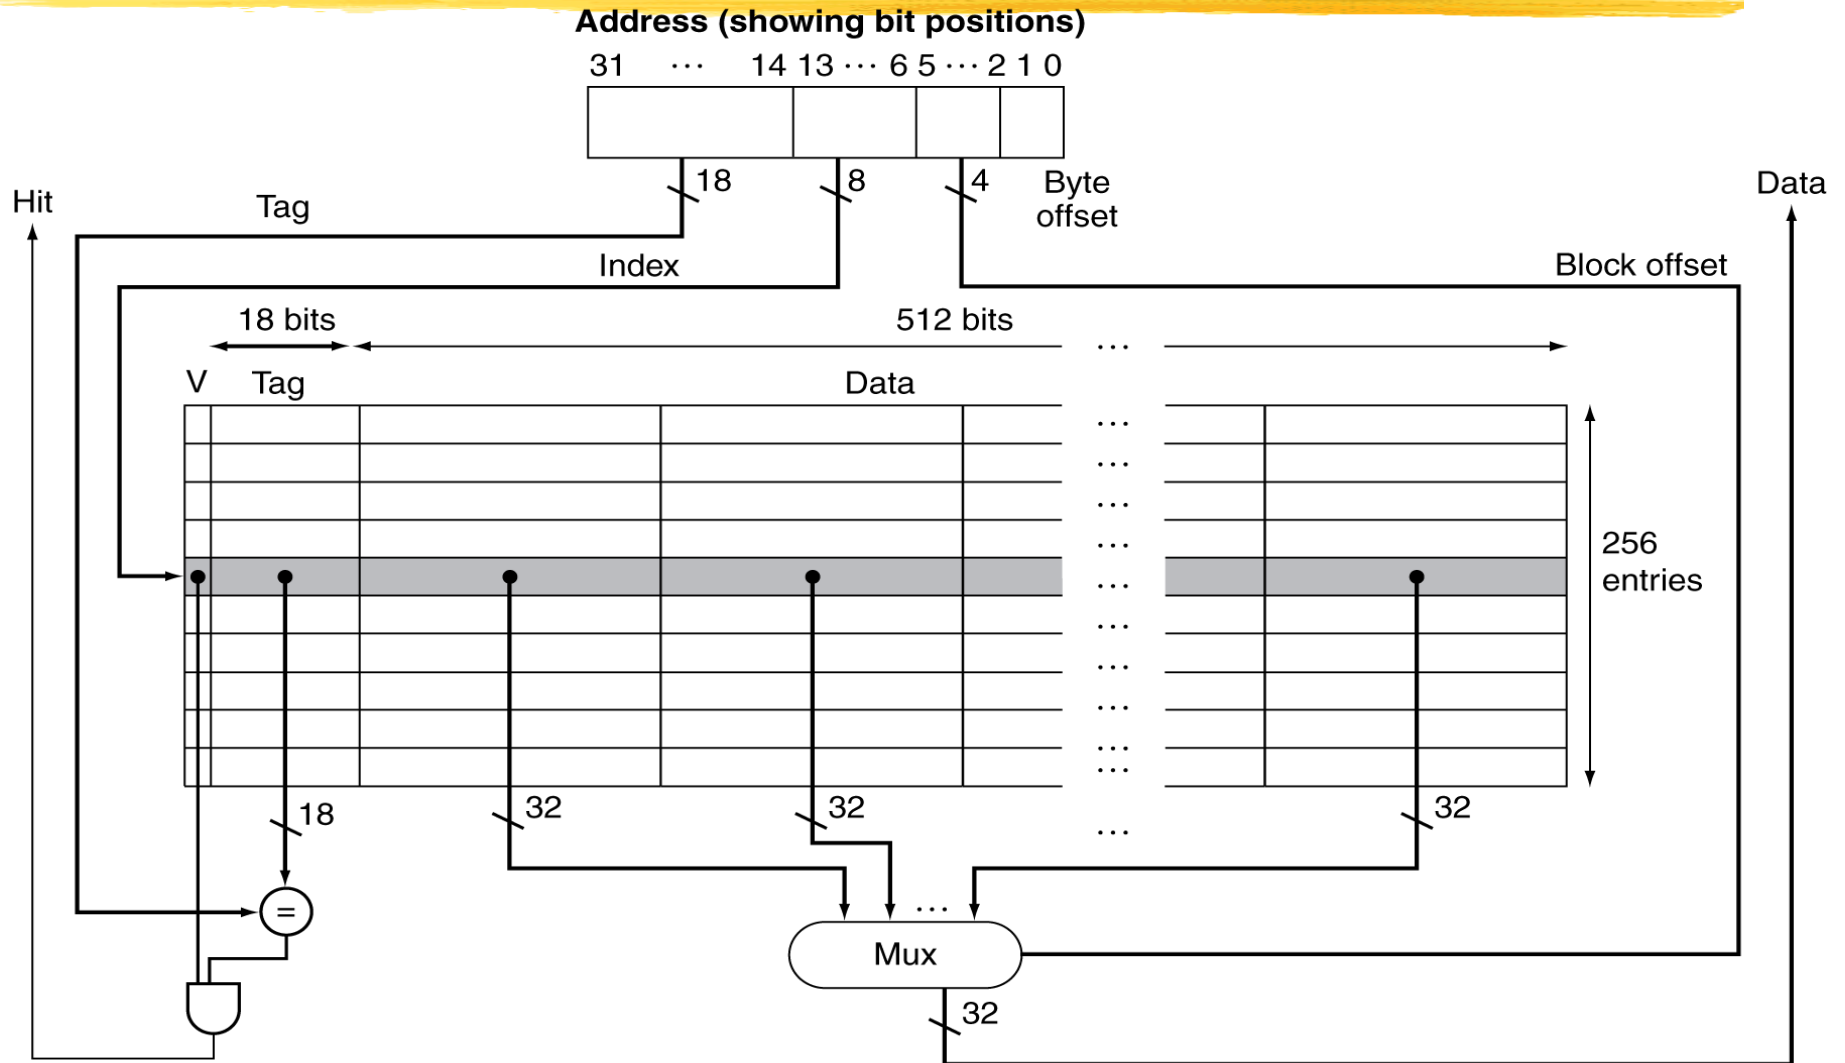
## Write Miss:

- ◆ What should happen on a write miss?
- ◆ Alternatives for write-through
  - Allocate on miss: fetch the block
  - Write around: don't fetch the block
    - Since programs often write a whole block before reading it (e.g., initialization)
- ◆ For write-back
  - Usually fetch the block

# Example: Intrinsity FastMATH


- ◆ **Embedded MIPS processor**
  - 12-stage pipeline
  - Instruction and data access on each cycle
- ◆ **Split cache: separate I-cache and D-cache**
  - Each 16KB: 256 blocks × 16 words/block
  - D-cache: write-through or write-back
- ◆ **SPEC2000 miss rates**
  - I-cache: 0.4%
  - D-cache: 11.4%
  - Weighted average: 3.2%

# Example: Intrinsity FastMATH




# Memory Design to Support Cache


- ◆ How to increase memory bandwidth to reduce miss penalty?



a. One-word-wide memory organization



b. Wide memory organization

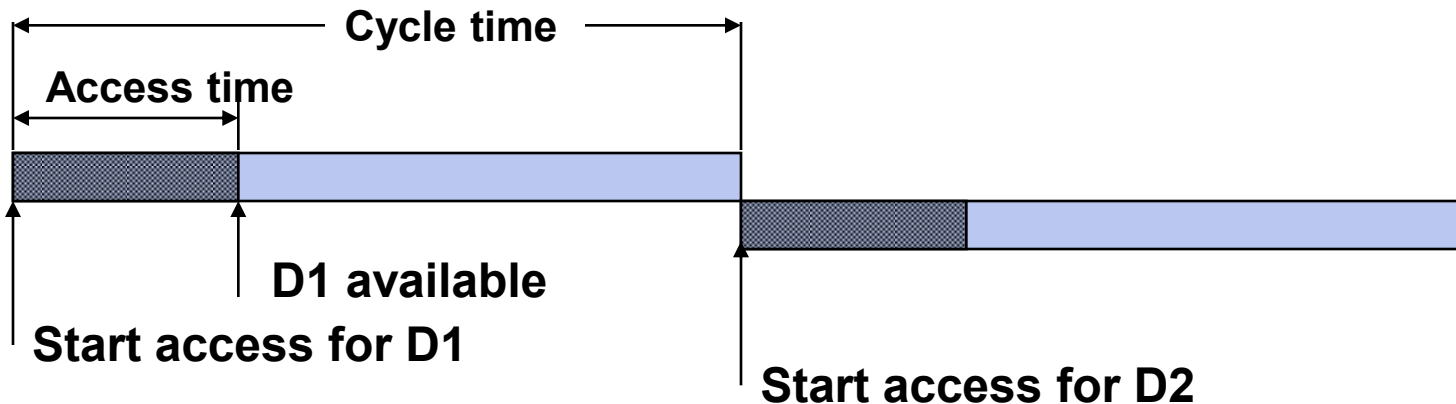


c. Interleaved memory organization

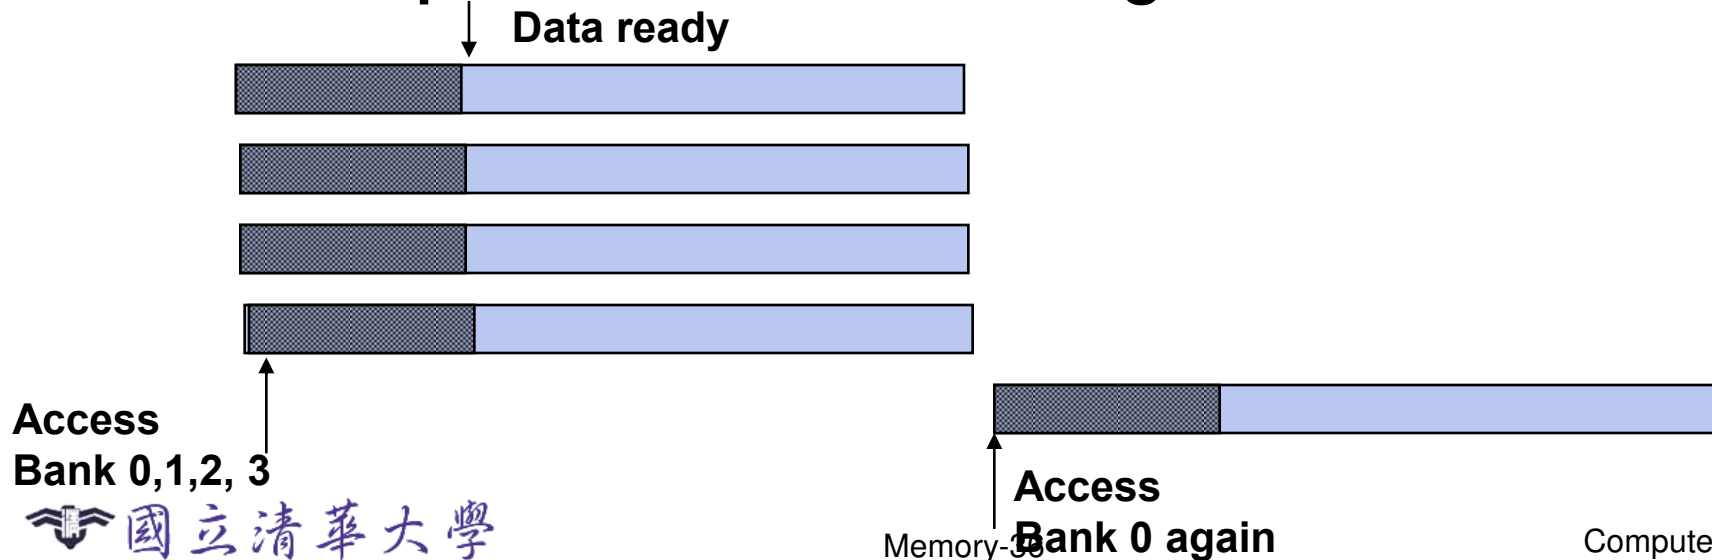
Fig. 5.11

# Interleaving for Bandwidth

## ◆ Access pattern without interleaving:



## ◆ Access pattern with interleaving



Memory-30

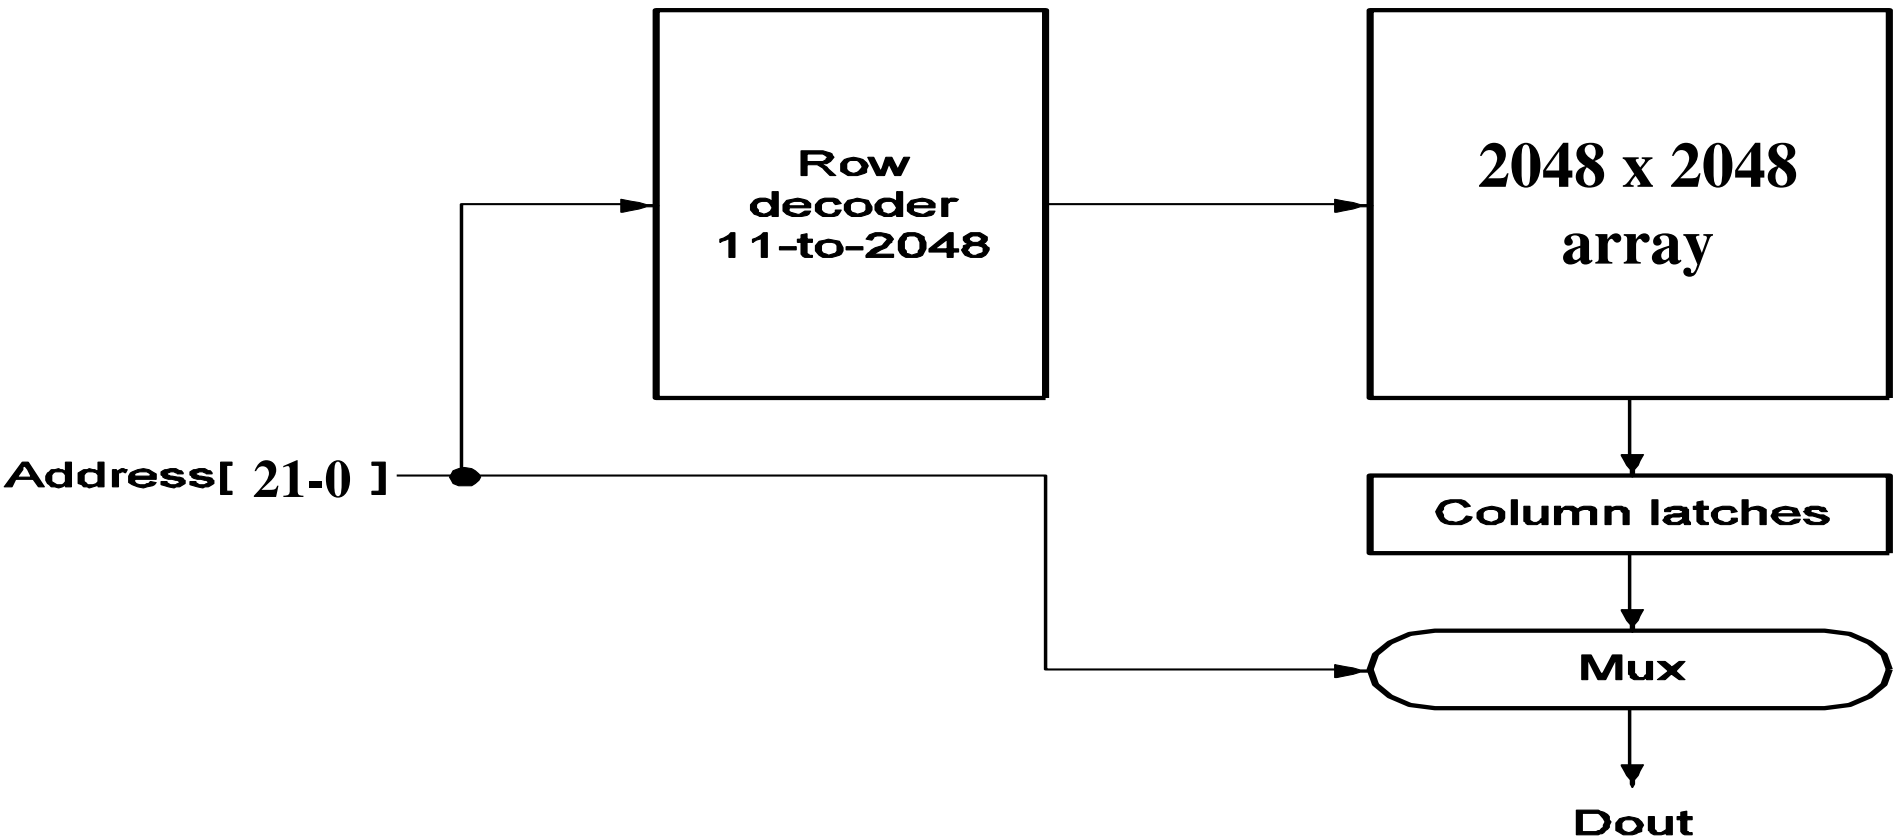
Computer Architecture

# Miss Penalty for Different Memory Organizations

## Assume

- ◆ 1 memory bus clock to send the address
- ◆ 15 memory bus clocks for each DRAM access initiated
- ◆ 1 memory bus clock to send a word of data
- ◆ A cache block = 4 words
- ◆ Three memory organizations :
  - A **one-word-wide bank** of DRAMs
  - Miss penalty =  $1 + 4 \times 15 + 4 \times 1 = 65$
  - A **four-word-wide bank** of DRAMs
  - Miss penalty =  $1 + 15 + 1 = 17$
  - A **four-bank, one-word-wide bus** of DRAMs
  - Miss penalty =  $1 + 1 \times 15 + 4 \times 1 = 20$

# Access of DRAM



# DDR SDRAM


## Double Data Rate Synchronous DRAMs

- ◆ Burst access from a sequential locations
- ◆ Starting address, burst length
- ◆ Data transferred under control of clock  
(300 MHz, 2004)
- ◆ Clock is used to eliminate the need of synchronization and the need of supplying successive address
- ◆ Data transfer on both leading and falling edge of clock



# DRAM Generations

Year	Capacity	\$/GB
1980	64Kbit	\$1500000
1983	256Kbit	\$500000
1985	1Mbit	\$200000
1989	4Mbit	\$50000
1992	16Mbit	\$15000
1996	64Mbit	\$10000
1998	128Mbit	\$4000
2000	256Mbit	\$1000
2004	512Mbit	\$250
2007	1Gbit	\$50



Trac:access time to a new row

Tcac:column access time to existing row

# Outline

- ◆ Memory hierarchy
- ◆ The basics of caches
- ◆ Measuring and improving cache performance
- ◆ Virtual memory
- ◆ A common framework for memory hierarchy
- ◆ Using a Finite State Machine to Control a Simple Cache
- ◆ Parallelism and Memory Hierarchies: Cache Coherence

# Measuring Cache Performance

- ◆ Components of CPU time
  - Program execution cycles
    - Includes cache hit time
  - Memory stall cycles
    - Mainly from cache misses
- ◆ With simplifying assumptions:

$$\begin{aligned} & \text{Memory stall cycles} \\ &= \frac{\text{Memory accesses}}{\text{Program}} \times \text{Miss rate} \times \text{Miss penalty} \\ &= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Misses}}{\text{Instruction}} \times \text{Miss penalty} \end{aligned}$$

# Cache Performance Example

## ◆ Given

- I-cache miss rate = 2%
- D-cache miss rate = 4%
- Miss penalty = 100 cycles
- Base CPI (ideal cache) = 2
- Load & stores are 36% of instructions

## ◆ Miss cycles per instruction

- I-cache:  $0.02 \times 100 = 2$
- D-cache:  $0.36 \times 0.04 \times 100 = 1.44$

## ◆ Actual CPI = $2 + 2 + 1.44 = 5.44$

- Ideal CPU is  $5.44/2 = 2.72$  times faster

# Average Access Time

- ◆ Hit time is also important for performance
- ◆ Average memory access time (AMAT)
  - $AMAT = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$
- ◆ Example
  - CPU with 1ns clock, hit time = 1 cycle, miss penalty = 20 cycles, l-cache miss rate = 5%
  - $AMAT = 1 + 0.05 \times 20 = 2\text{ns}$ 
    - 2 cycles per instruction

# Performance Summary

- ◆ When CPU performance increased
  - Miss penalty becomes more significant
- ◆ Decreasing base CPI
  - Greater proportion of time spent on memory stalls
- ◆ Increasing clock rate
  - Memory stalls account for more CPU cycles
- ◆ Can't neglect cache behavior when evaluating system performance


# Improving Cache Performance

- ◆ Reduce the time to hit in the cache
- ◆ Decreasing the miss ratio
- ◆ Decreasing the miss penalty

# Direct Mapped

## ◆ Block Placement :

- For each item of data at the lower level, there is exactly one location in cache where it might be
- Address mapping: modulo number of blocks





# Associative Caches

## ◆ Fully associative

- Allow a given block to go in any cache entry
- Requires all entries to be searched at once
- Comparator per entry (expensive)


## ◆ $n$ -way set associative

- Each set contains  $n$  entries
- Block number determines which set
  - (Block number) modulo (#Sets in cache)
- Search all entries in a given set at once
- $n$  comparators (less expensive)


# Associative Cache Example

## ◆ Placement of a block whose address is 12:


Direct mapped



Set associative



Fully associative



# Possible Associativity Structures

(direct mapped)

Block	Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

Two-way set associative

Set	Tag	Data	Tag	Data
0				
1				
2				
3				

## An 8-block cache

Four-way set associative

Set	Tag	Data	Tag	Data	Tag	Data	Tag	Data
0								
1								

Eight-way set associative (fully associative)

Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data

# Associativity Example

- ◆ Compare 4-block caches
  - Direct mapped, 2-way set associative, fully associative
  - Block access sequence: 0, 8, 0, 6, 8
- ◆ Direct mapped

time ↓

Block address	Cache index	Hit/miss	Cache content after access			
			0	1	2	3
0	0	miss	Mem[0]			
8	0	miss	Mem[8]			
0	0	miss	Mem[0]			
6	2	miss	Mem[0]		Mem[6]	
8	0	miss	Mem[8]		Mem[6]	

# Associativity Example

## ◆ 2-way set associative

Block address	Cache index	Hit/miss	Cache content after access			
			Set 0		Set 1	
0	0	miss	Mem[0]			
8	0	miss	Mem[0]	Mem[8]		
0	0	hit	Mem[0]	Mem[8]		
6	0	miss	Mem[0]	Mem[6]		
8	0	miss	Mem[8]	Mem[6]		

time ↓

## ◆ Fully associative

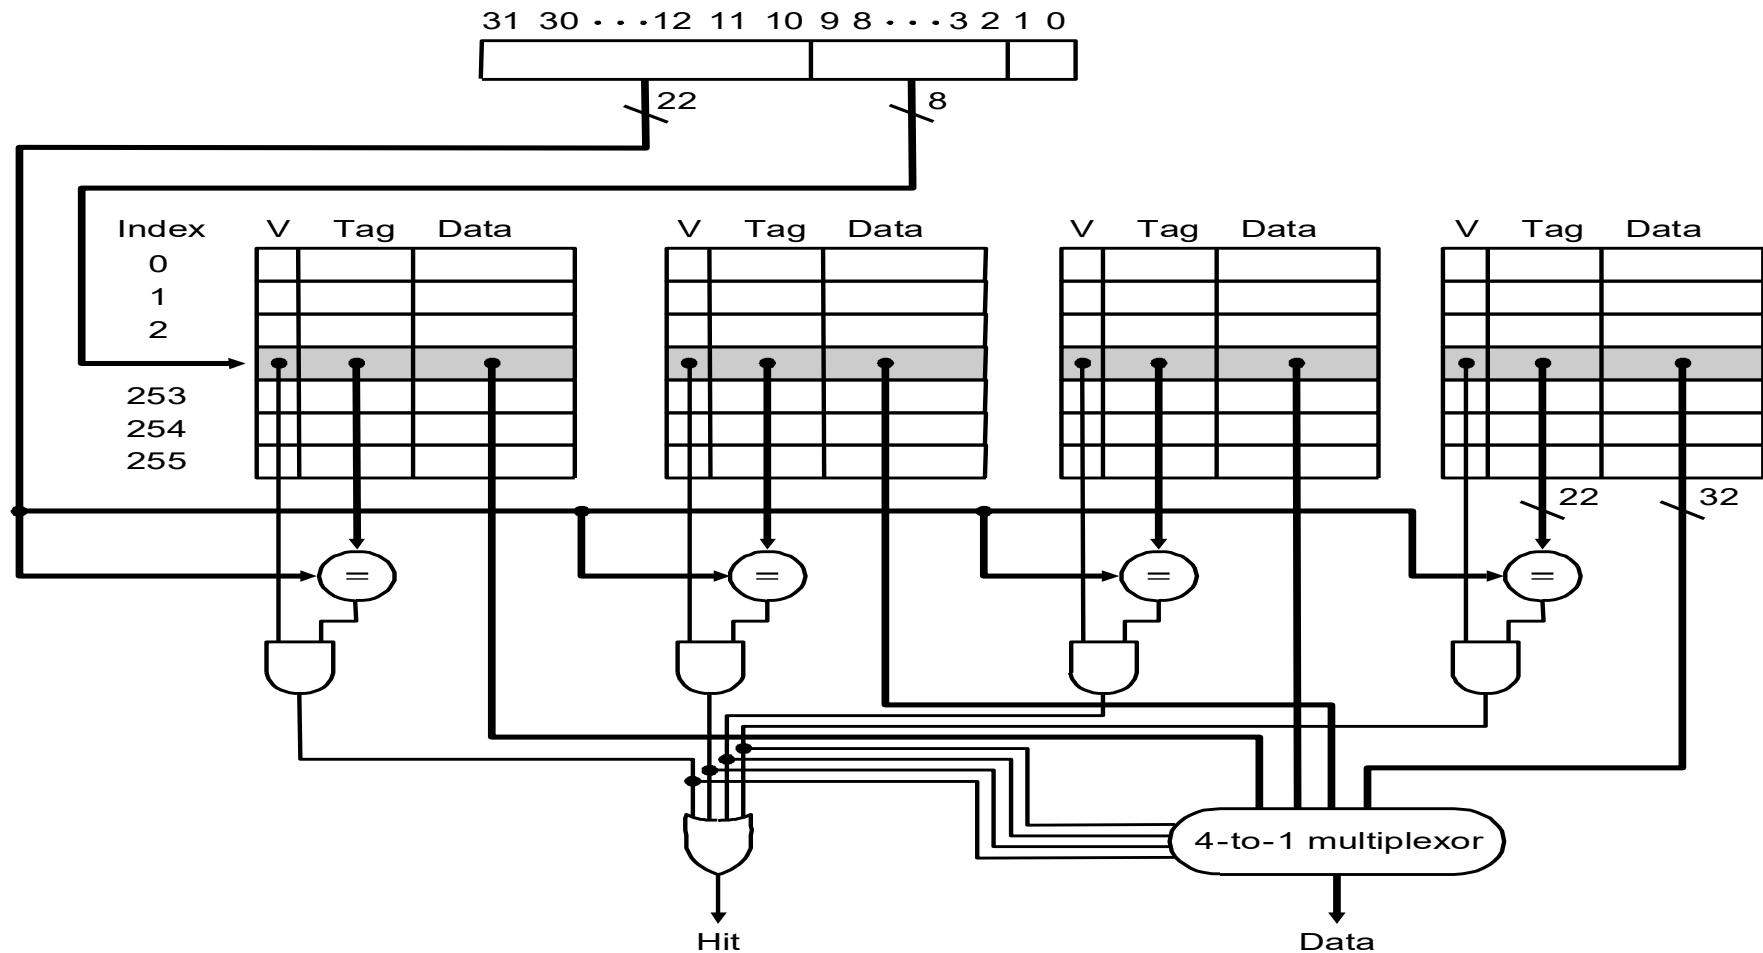
Block address		Hit/miss	Cache content after access			
0		miss	Mem[0]			
8		miss	Mem[0]	Mem[8]		
0		hit	Mem[0]	Mem[8]		
6		miss	Mem[0]	Mem[8]	Mem[6]	
8		hit	Mem[0]	Mem[8]	Mem[6]	

time ↓

# How Much Associativity

- ◆ Increased associativity decreases miss rate
  - But with diminishing returns
- ◆ Simulation of a system with 64KB D-cache, 16-word blocks, SPEC2000
  - 1-way: 10.3%
  - 2-way: 8.6%
  - 4-way: 8.3%
  - 8-way: 8.1%

# A 4-Way Set-Associative Cache



◆ Increasing associativity shrinks index, expands tag

# Data Placement Policy

- ◆ **Direct mapped cache:**
  - Each memory block mapped to one location
  - No need to make any decision
  - Current item replaces previous one in location
- ◆ **N-way set associative cache:**
  - Each memory block has choice of N locations
- ◆ **Fully associative cache:**
  - Each memory block can be placed in ANY cache location
- ◆ **Misses in N-way set-associative or fully associative cache:**
  - Bring in new block from memory
  - Throw out a block to make room for new block
  - Need to decide on which block to throw out



# Cache Block Replacement

- ◆ Direct mapped: no choice
- ◆ Set associative or fully associative:
  - Random
  - LRU (Least Recently Used):
    - Hardware keeps track of the access history and replace the block that has not been used for the longest time
  - An example of a pseudo LRU (for a two-way set associative) :
    - use a pointer pointing at each block in turn
    - whenever an access to the block the pointer is pointing at, move the pointer to the next block
    - when need to replace, replace the block currently pointed at

# Comparing the Structures

- ◆ **N-way set-associative cache**
  - N comparators vs. 1
  - Extra MUX delay for the data
  - Data comes **AFTER** Hit/Miss decision and set selection
- ◆ **Direct mapped cache**
  - Cache block is available **BEFORE** Hit/Miss:
  - Possible to assume a hit and continue, recover later if miss

# Multilevel Caches

- ◆ **Primary cache attached to CPU**
  - Small, but fast
- ◆ **Level-2 cache services misses from primary cache**
  - Larger, slower, but still faster than main memory
- ◆ **Main memory services L-2 cache misses**
- ◆ **Some high-end systems include L-3 cache**

# Multilevel Cache Example

## ◆ Given

- CPU base CPI = 1, clock rate = 4GHz
- Miss rate/instruction = 2%
- Main memory access time = 100ns

## ◆ With just primary cache

- Miss penalty =  $100\text{ns} / 0.25\text{ns} = 400$  cycles
- Effective CPI =  $1 + 0.02 \times 400 = 9$

## Example (cont.)

- ◆ Now add L-2 cache
  - Access time = 5ns (to M: 100ns)
  - Global miss rate to main memory = 0.5% (to M 2%)
- ◆ Primary miss with L-2 hit
  - Penalty =  $5\text{ns}/0.25\text{ns} = 20$  cycles
- ◆ Primary miss with L-2 miss (0.5%)
  - Extra penalty = 400 cycles
- ◆  $\text{CPI} = 1 + 0.02 \times 20 + 0.005 \times 400 = 3.4$
- ◆ Performance ratio =  $9/3.4 = 2.6$

# Multilevel Cache Considerations


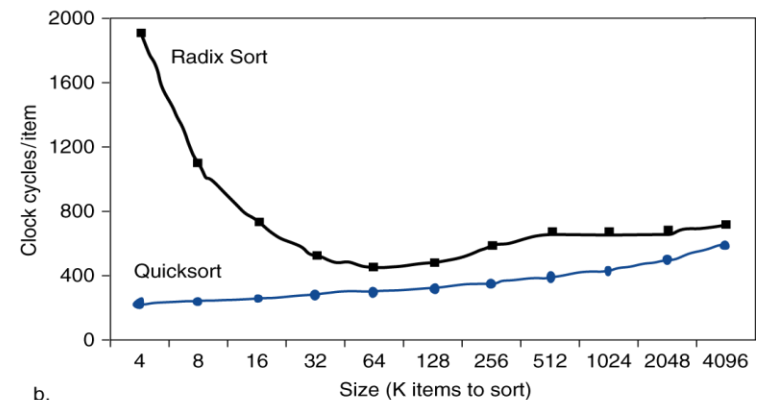

- ◆ **Primary cache**
  - Focus on minimal hit time
- ◆ **L-2 cache**
  - Focus on low miss rate to avoid main memory access
  - Hit time has less overall impact
- ◆ **Results**
  - L-1 cache usually smaller than a single-level cache
  - L-1 block size smaller than L-2 block size

# Interactions with Advanced CPUs

- ◆ **Out-of-order CPUs can execute instructions during cache miss**
  - Pending store stays in load/store unit
  - Dependent instructions wait in reservation stations
    - Independent instructions continue
- ◆ **Effect of miss depends on program data flow**
  - Much harder to analyze
  - Use system simulation

# Interactions with Software

- ◆ Misses depend on memory access patterns
  - Algorithm behavior
  - Compiler optimization for memory access





# Outline

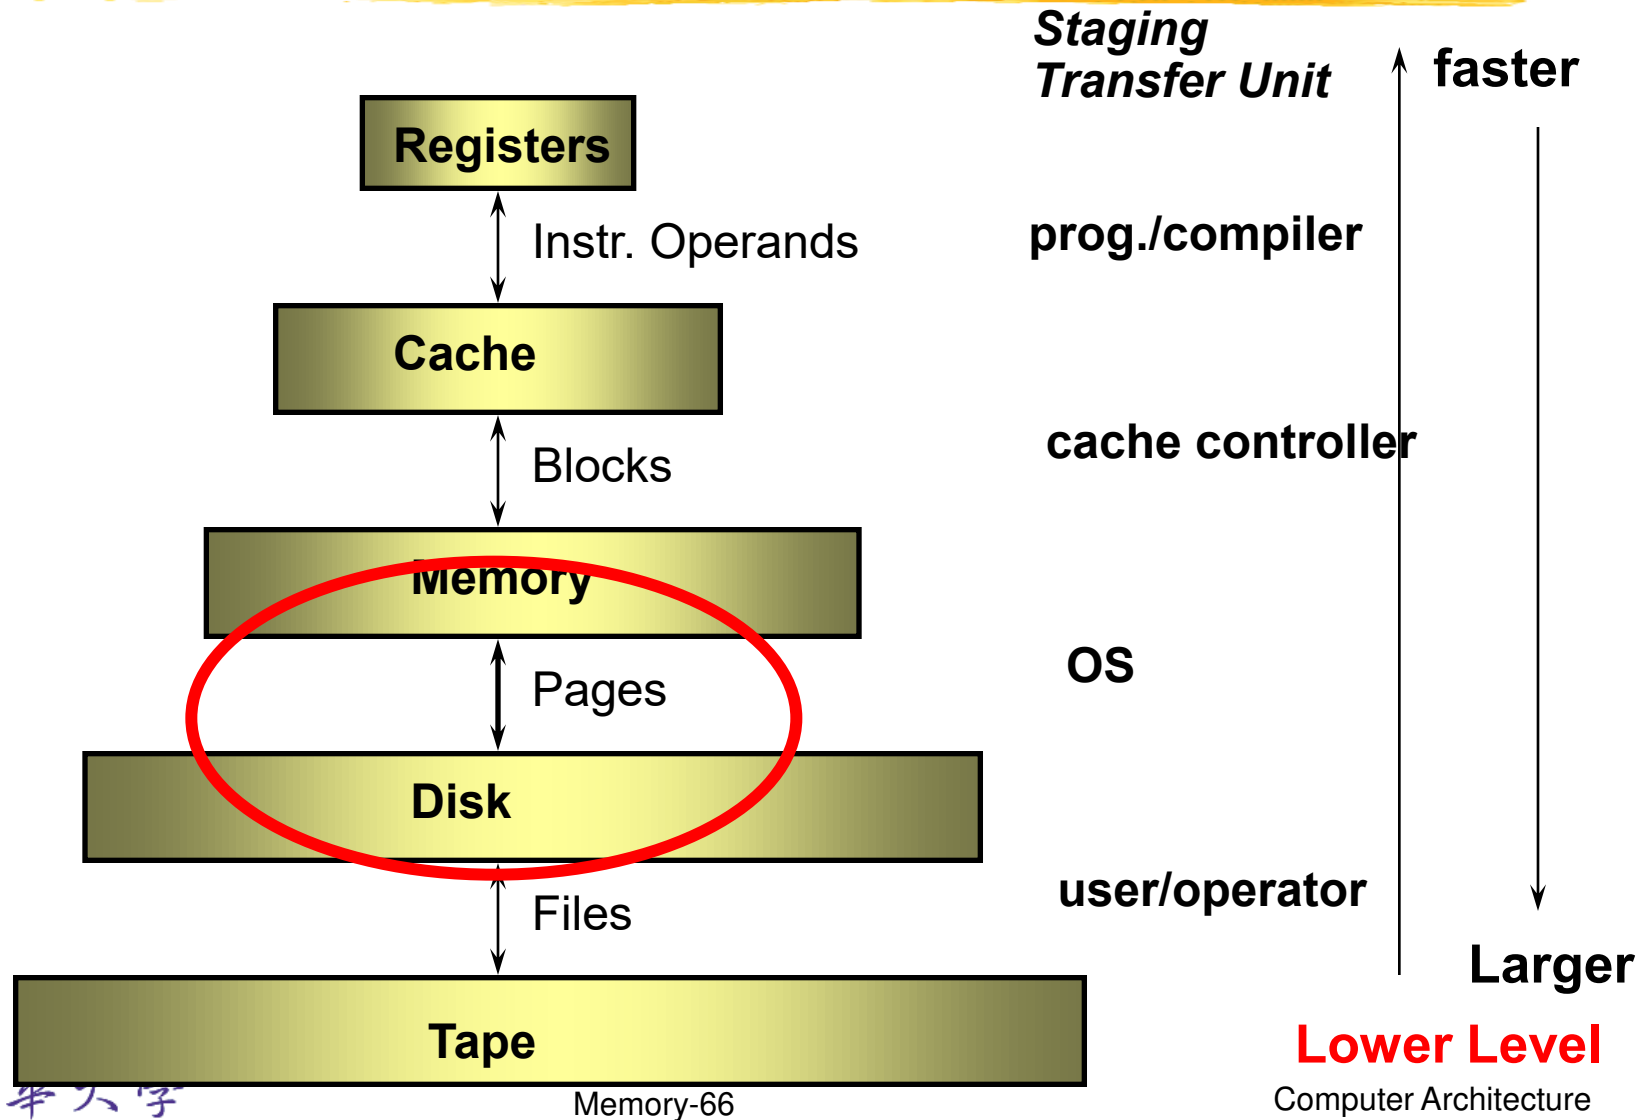
- ◆ Memory hierarchy
- ◆ The basics of caches
- ◆ Measuring and improving cache performance
- ◆ Virtual memory
- ◆ A common framework for memory hierarchy
- ◆ Using a Finite State Machine to control a simple cache
- ◆ Parallelism and memory hierarchies: cache coherence

# Outline

- ◆ Memory hierarchy
- ◆ The basics of caches
- ◆ Measuring and improving cache performance
- ◆ Virtual memory
  - Basics

# Levels of Memory Hierarchy

Upper Level




# Virtual Memory

- ◆ Use main memory as a “cache” for secondary (disk) storage
  - Managed jointly by **CPU hardware** and the **operating system** (OS)
- ◆ Programs share main memory
  - Each gets a private virtual address space holding its frequently used code and data, starting at address 0, only accessible to itself
    - yet, any can run anywhere in physical memory
    - executed in a name space (virtual address space) different from memory space (physical address space)
    - virtual memory implements the **translation** from virtual space to physical space
  - Protected from other programs
  - Every program has lots of memory (> physical memory)

# Virtual Memory - Continued


- ◆ CPU and OS translate virtual addresses to physical addresses
  - VM “block” is called **a page**
  - VM translation “miss” is called **a page fault**

# Virtual Memory



# Basic Issues in Virtual Memory

- ◆ Size of data blocks that are transferred from disk to main memory
- ◆ Which region of memory to hold new block  
=> placement policy
- ◆ When memory is full, then some region of memory must be released to make room for the new block => replacement policy
- ◆ When to fetch missing items from disk?
  - Fetch only on a fault => demand load policy



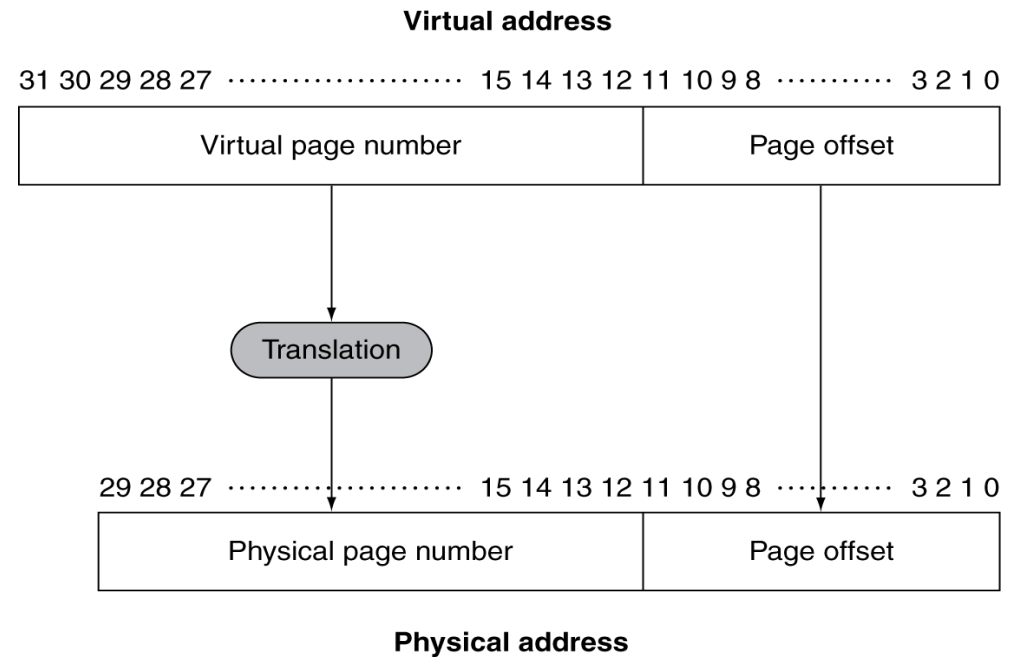

# Block Size and Placement Policy

- ◆ Huge miss penalty: a page fault may take millions of cycles to process
  - Pages should be fairly large (e.g., 4KB) to amortize the high access time
  - Reducing page faults is important
    - fully associative **placement**  
=> use page table (in memory) to **locate** pages



# Address Translation

## ◆ Fixed-size pages (e.g., 4K)



# Paging

- Virtual and physical address space

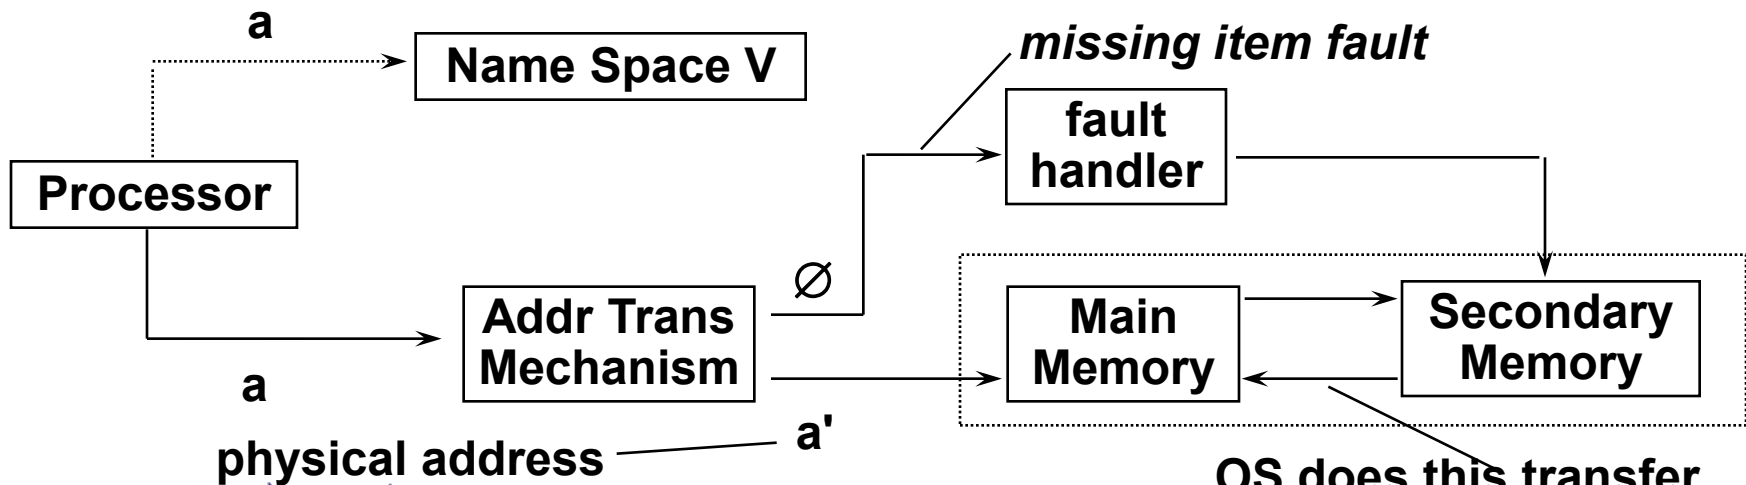
*pages* *page frames*  
partitioned into blocks of equal size

- Key operation: address mapping

MAP:  $V \rightarrow M \cup \{\emptyset\}$  address mapping function

$\text{MAP}(a) = a'$  if data at virtual address  $a$  is present in  
physical address  $a'$  and  $a'$  in  $M$

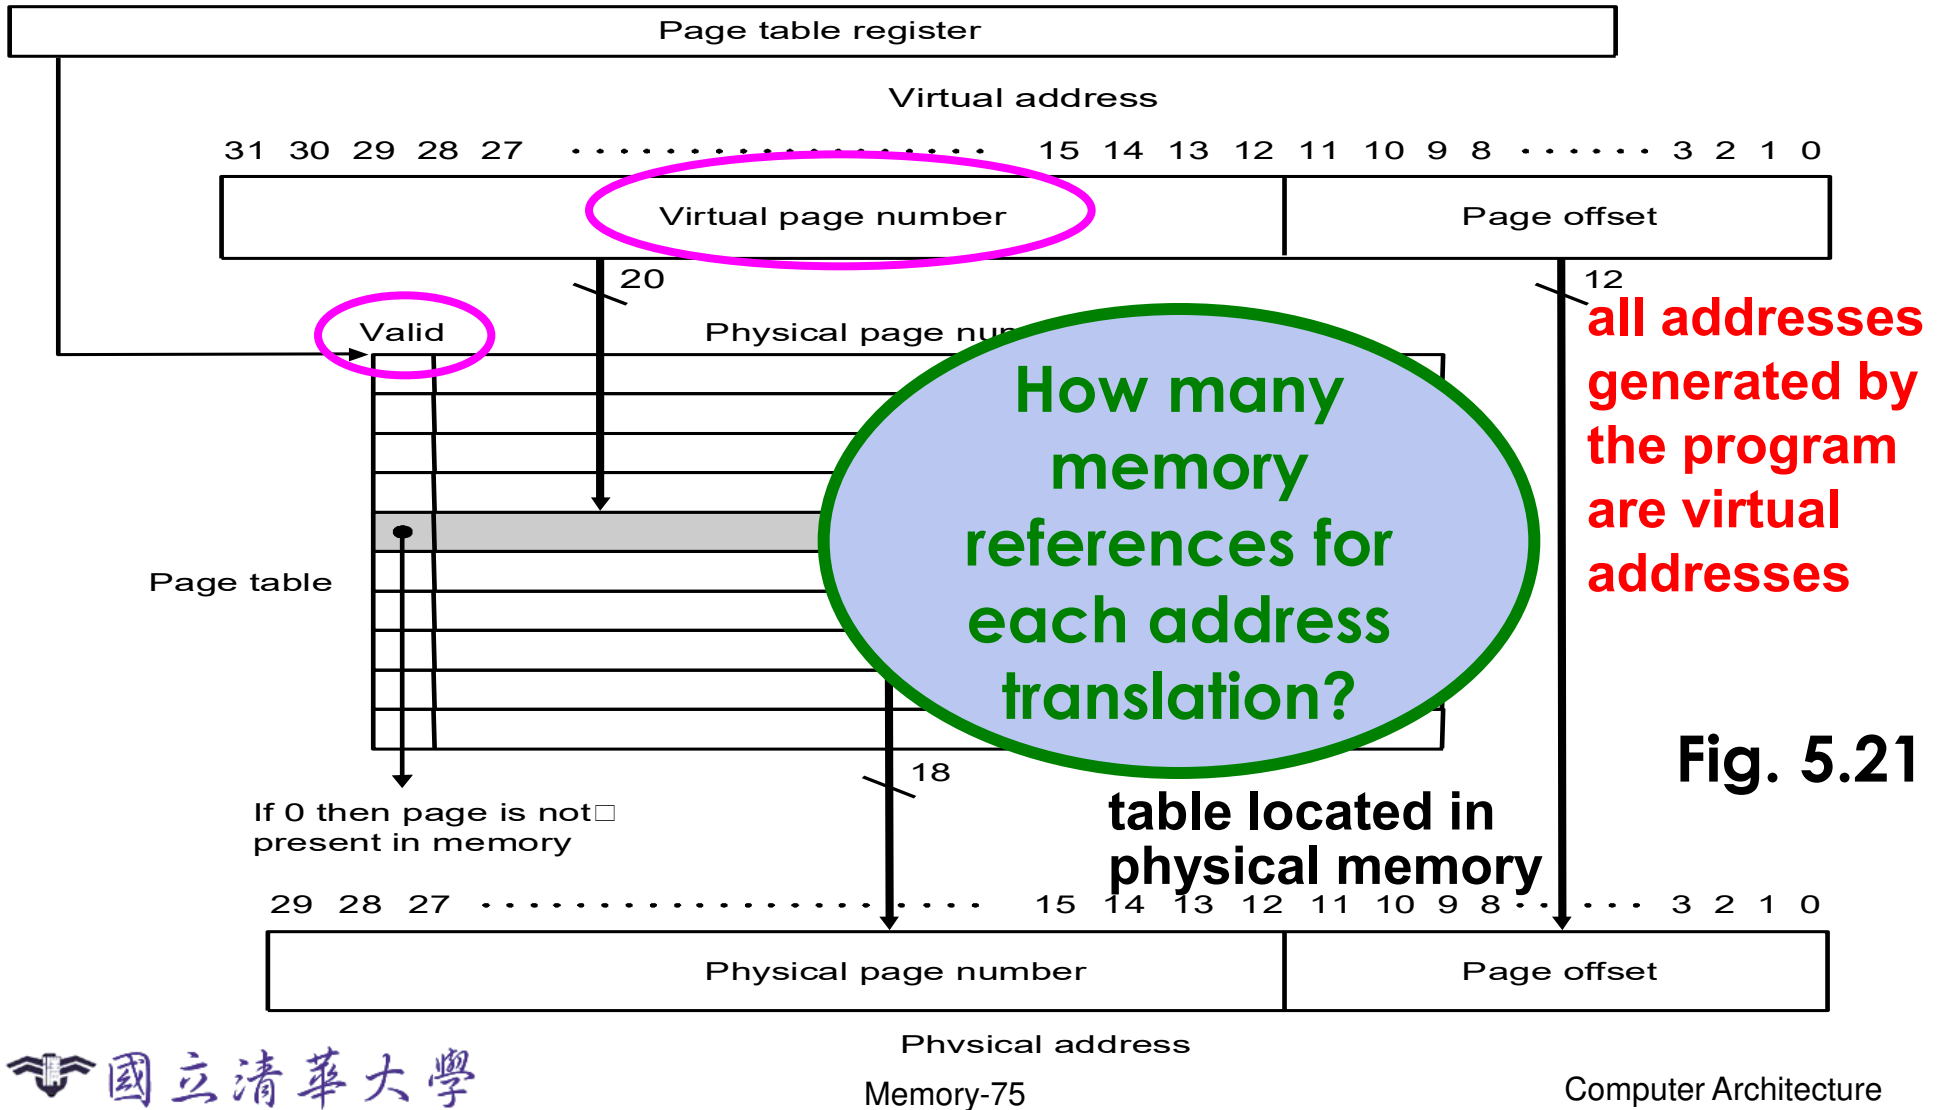
$= \emptyset$  if data at virtual address  $a$  is not present in  $M$



# Page Tables

- ◆ Stores placement information
  - Array of page table entries, indexed by virtual page number
  - Page table register in CPU points to page table in physical memory
- ◆ If page is present in memory
  - PTE stores the physical page number
  - Plus other status bits (referenced, dirty, ...)
- ◆ If page is not present
  - PTE can refer to location in swap space on disk

# Page Tables



# Page Fault: What Happens When You Miss?

- ◆ Page fault means that page is not resident in memory
- ◆ Huge miss penalty: a page fault may take millions of cycles to process
- ◆ Hardware must detect situation but it cannot remedy the situation
- ◆ Can handle the faults in software instead of hardware, because handling time is small compared to disk access
  - the software can be very smart or complex
  - the faulting process can be context-switched

# Handling Page Faults

- ◆ **Hardware must trap to the operating system so that it can remedy the situation**
  - Pick a page to discard (may write it to disk)
  - Load the page in from disk
  - Update the page table
  - Resume to program so HW will retry and succeed!
  
- ◆ **In addition, OS must know where to find the page**
  - Create space on disk for all pages of process (swap space)
  - Use a data structure to record where each valid page is on disk (may be part of page table)
  - Use another data structure to track which process and virtual addresses use each physical page  
=> for replacement purpose

# Page Replacement and Writes

- ◆ To reduce page fault rate, prefer least-recently used (LRU) replacement
  - Reference bit (aka **use bit**) in PTE set to 1 on access to page
  - Periodically cleared to 0 by OS
  - A page with reference bit = 0 has not been used recently
- ◆ Disk writes take millions of cycles
  - Block at once, not individual locations
  - Write through is impractical
  - Use write-back
  - Dirty bit in PTE set when page is written

# Problems of Page Table

- ◆ Page table is too big
- ◆ Access to page table is too slow (needs one memory read)



# Outline


- ◆ Memory hierarchy
- ◆ The basics of caches
- ◆ Measuring and improving cache performance
- ◆ Virtual memory
  - Basics
  - Handling huge page table: page table is too big

# Impact of Paging: Huge Page Table

- ◆ Page table occupies storage  
32-bit VA, 4KB page, 4bytes/entry  
=>  $2^{20}$  PTE, 4MB table
- ◆ Possible solutions:
  - Use bounds register to limit table size; add more if exceed
  - Let pages to grow in both directions  
=> 2 tables, 2 limit registers, one for hash, one for stack
  - Use hashing => page table same size as physical pages
  - Multiple levels of page tables
  - Paged page table (page table resides in virtual space)

# Hashing: Inverted Page Tables

- ◆ 28-bit virtual address, 4 KB per page, and 4 bytes per page-table entry
- ◆ The number of pages: 64K
- ◆ The number of physical frames: 16K
  - Page table size :  $64 \text{ K (pages \#)} \times 4 = 256 \text{ KB}$
  - Inverted page table :  $16 \text{ K (frame \#)} \times (4+?) = 64 \text{ KB}$



=> TLBs or virtually addressed caches are critical

# Two-level Page Tables

32-bit address:



○ 4 GB virtual address space


○ 4 KB of PTE1

(Each entry indicate if any page in the segment is allocated)

○ 4 MB of PTE2

- paged, holes

*What about a 48-64 bit address space?*



# Outline

- ◆ Memory hierarchy
- ◆ The basics of caches
- ◆ Measuring and improving cache performance
- ◆ Virtual memory
  - Basics
  - Handling huge page table
  - TLB (Translation Lookaside Buffer): access to page table is too slow (needs one memory read)

# Impact of Paging: More Memory Access !

- ◆ Each memory operation (instruction fetch, load, store) requires a page-table access!
  - Basically double number of memory operations
  - One to access the PTE
  - Then the actual memory access
- ◆ access to page tables has good locality

# Fast Translation Using a TLB

- ◆ Access to page tables has good locality
  - Fast cache of PTEs within the CPU
  - Called a Translation Look-aside Buffer (TLB)

# Fast Translation Using TLB (Translation Lookaside Buffer)

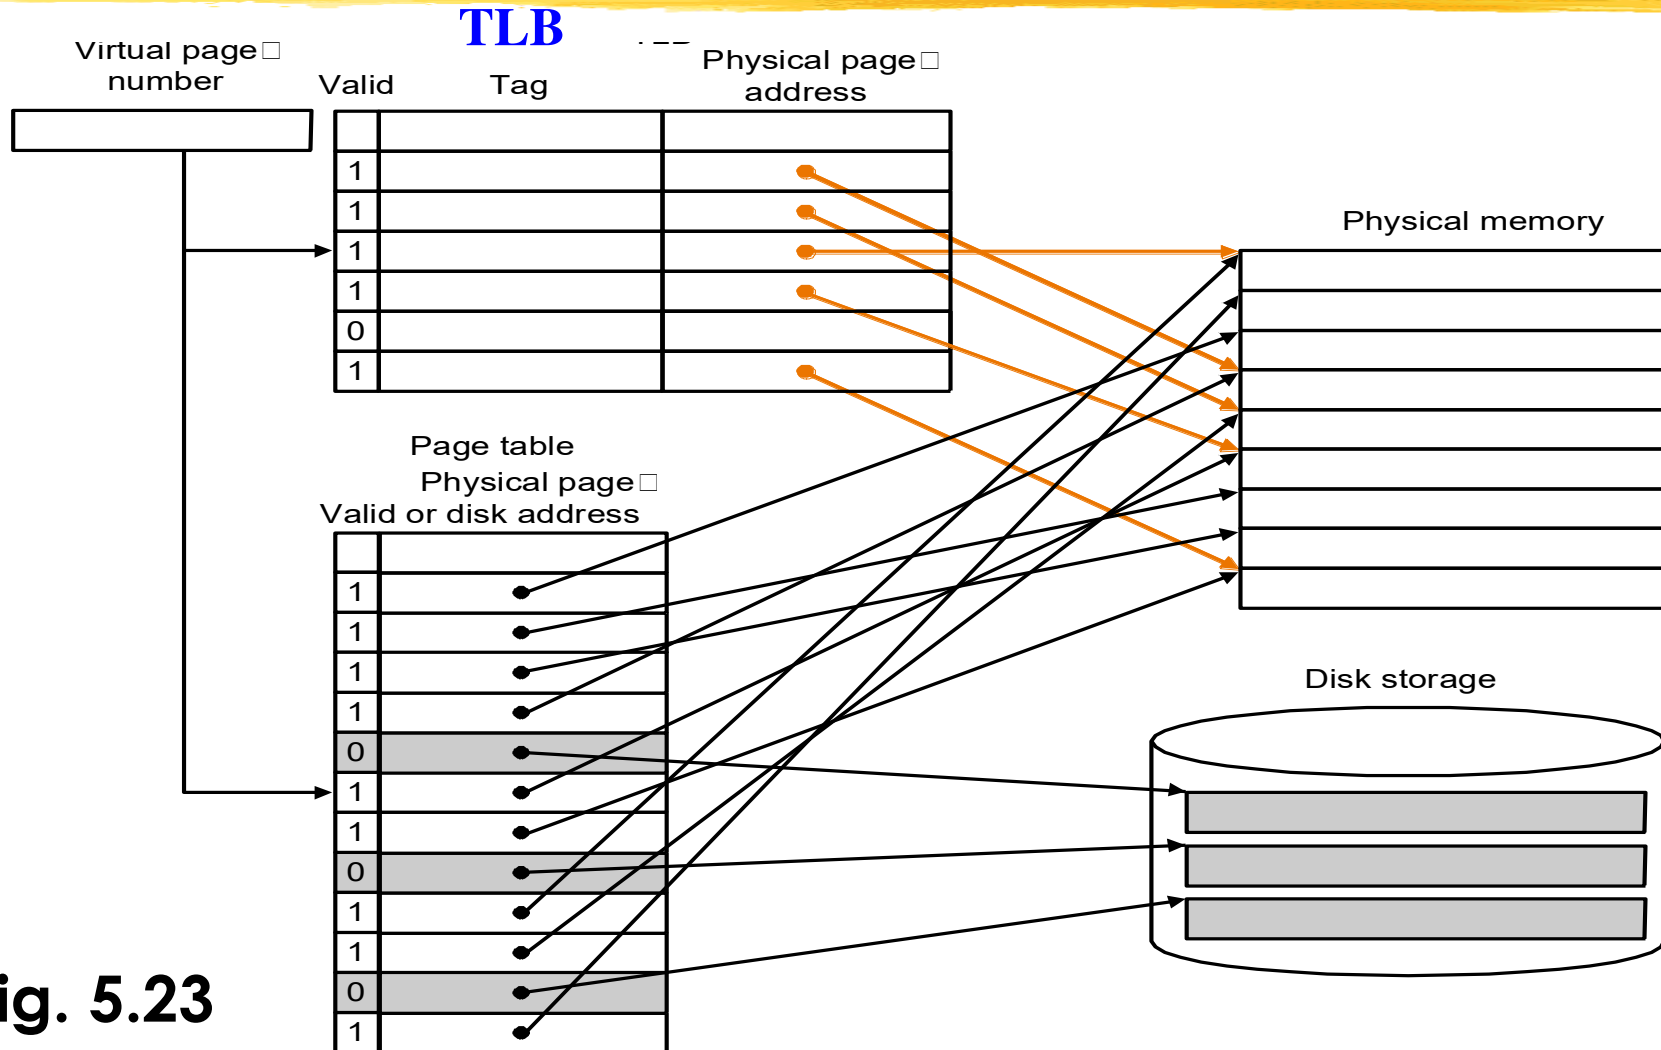


Fig. 5.23



# Translation Lookaside Buffer

- ◆ Typical RISC processors have *memory management unit* (MMU) which includes TLB and does page table lookup
  - TLB can be organized as fully associative, set associative, or direct mapped
  - TLBs are small, typical: 16–512 PTEs, 0.5–1 cycle for hit, 10–100 cycles for miss, 0.01%–1% miss rate
  - Misses could be handled by hardware or software

# TLB Hit

- ◆ TLB hit on read
- ◆ TLB hit on write:
  - Toggle dirty bit (write back to page table on replacement)

# TLB Miss

## ◆ If page is in memory

- Load the PTE from memory and retry
- Could be handled in hardware
  - Can get complex for more complicated page table structures
- Or in software
  - Raise a special exception, with optimized handler

## ◆ If page is not in memory (page fault)


- OS handles fetching the page and updating the page table (software)
- Then restart the faulting instruction

# Outline

- ◆ Memory hierarchy
- ◆ The basics of caches
- ◆ Measuring and improving cache performance
- ◆ Virtual memory
  - Basics
  - Handling huge page table
  - TLB (Translation Lookaside Buffer)
  - TLB and cache

# Making Address Translation Practical

- ◆ In VM, memory acts like a cache for disk
  - Page table maps virtual page numbers to physical frames
  - Use a page table cache for recent translation
    - => **Translation Lookaside Buffer (TLB)**



# Integrating TLB and Cache

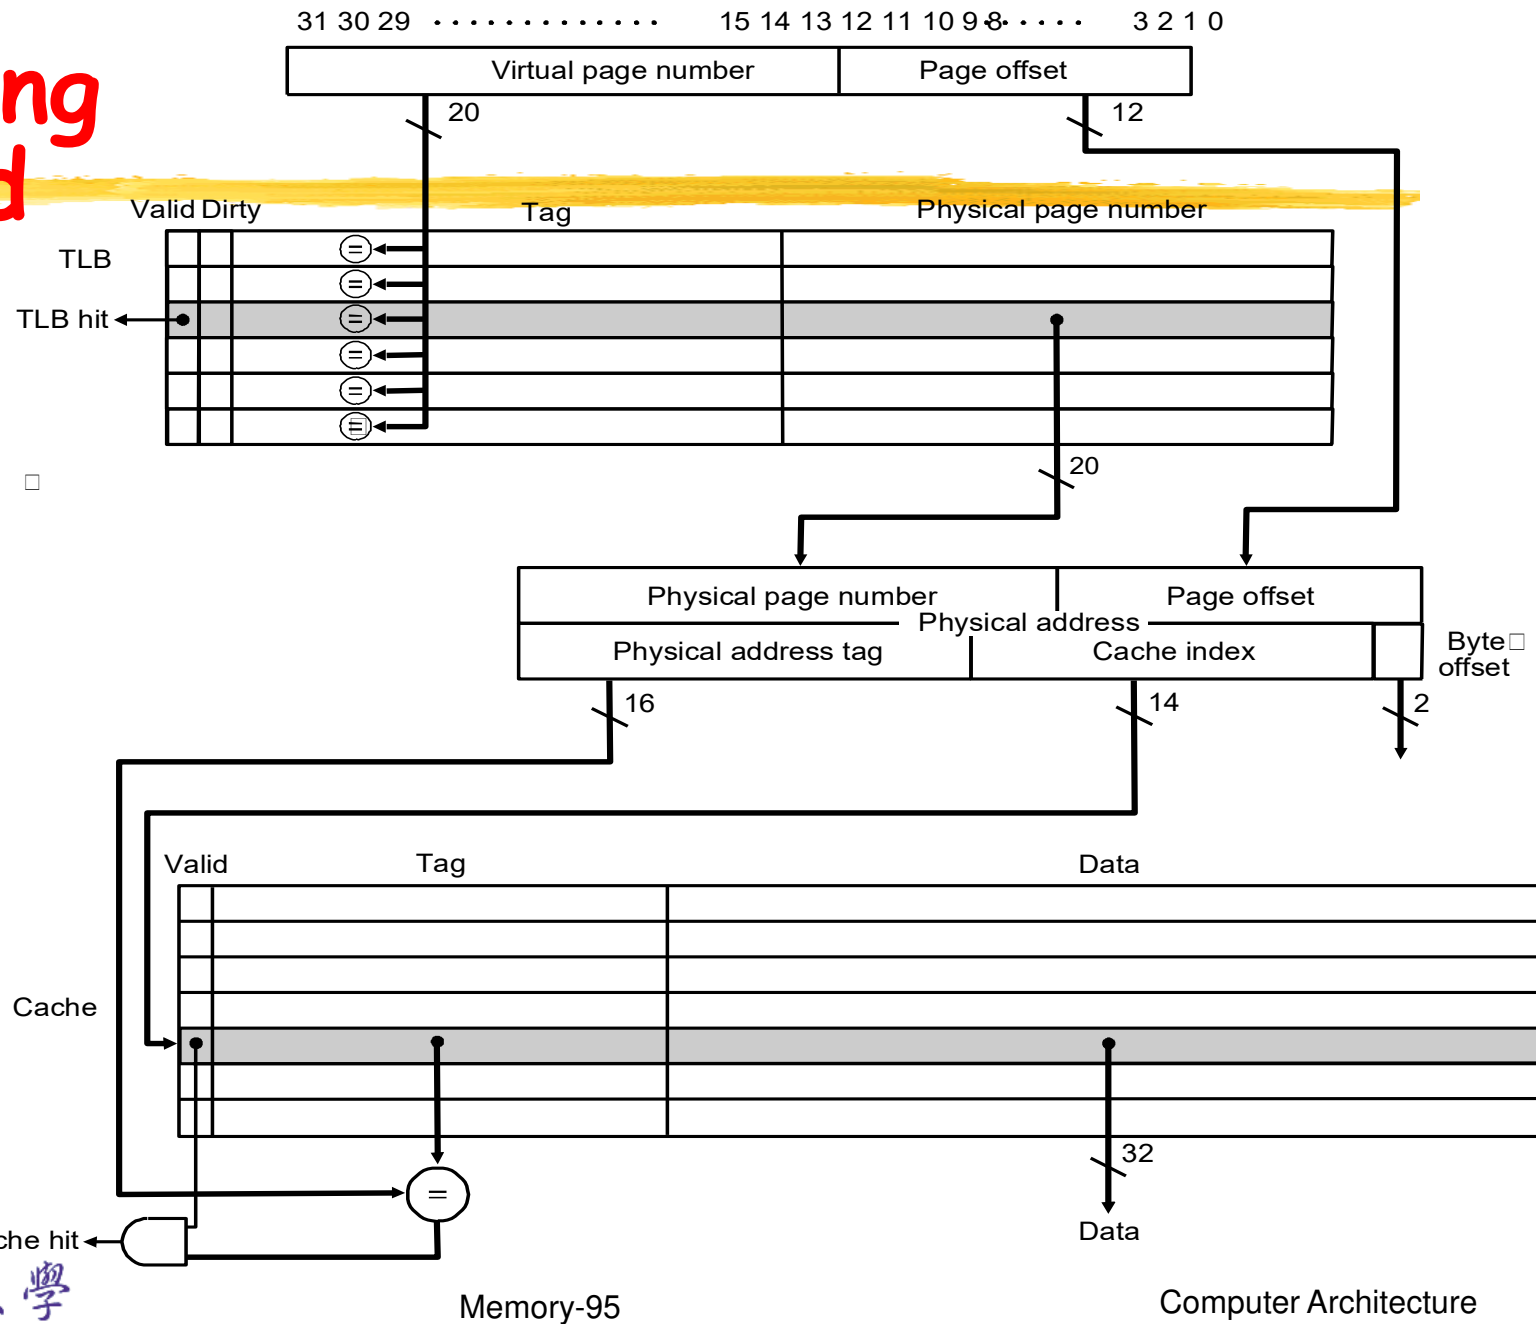
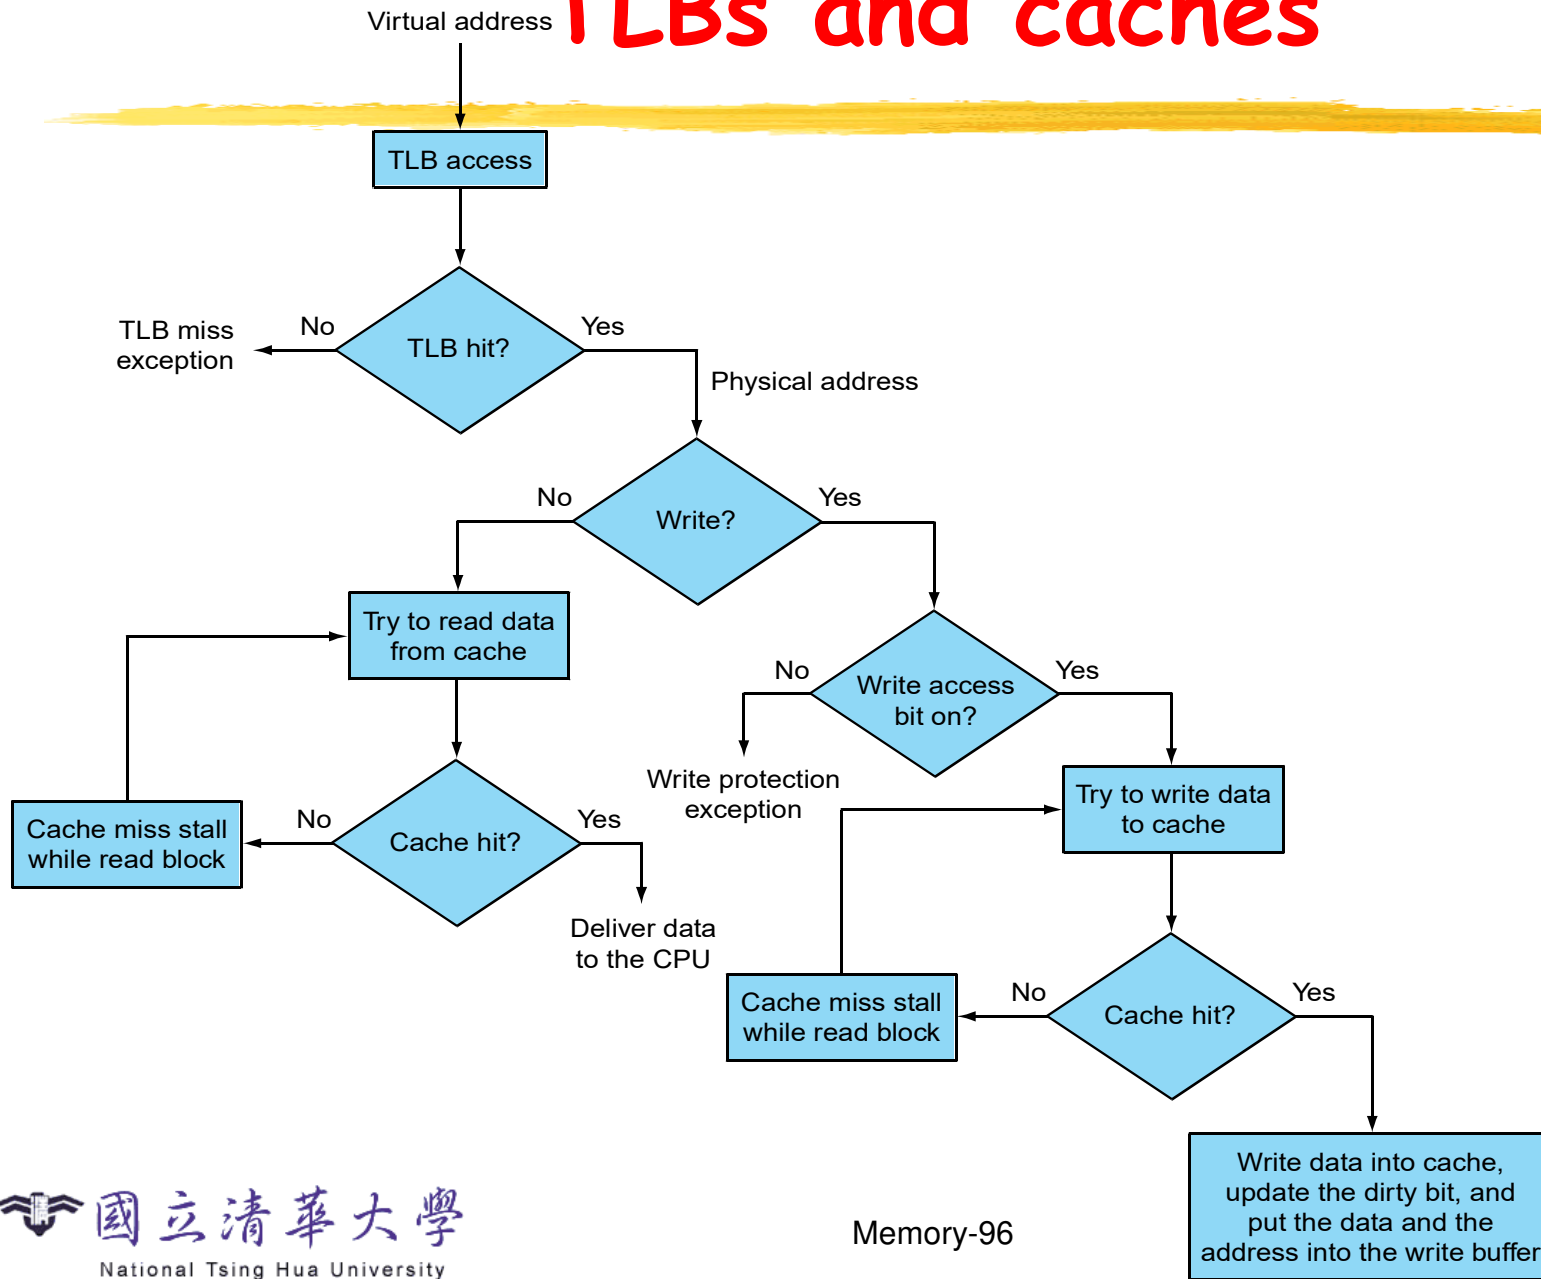


Fig. 5.24

# TLBs and caches




# Possible Combinations of Events

Cache	TLB	Page table	Possible? Conditions?
Miss	Hit	Hit	Yes; but page table never checked if TLB hits
Hit	Miss	Hit	TLB miss, but entry found in page table; after retry, data in cache
Miss	Miss	Hit	TLB miss, but entry found in page table; after retry, data miss in cache
Miss	Miss	Miss	TLB miss and is followed by a page fault; after retry, data miss in cache
Miss	Hit	Miss	impossible; not in TLB if page not in memory
Hit	Hit	Miss	impossible; not in TLB if page not in memory
Hit	Miss	Miss	impossible; not in cache if page not in memory




# Virtual Address and Cache

- ◆ TLB access is serial with cache access
  - Cache is physically indexed and tagged



- ◆ Alternative: *virtually addressed cache*
  - Cache is *virtually indexed and virtually tagged*



# Virtually Addressed Cache

- ◆ Require address translation only on miss!
- ◆ Problem:
  - Same virtual addresses (different processes) map to different physical addresses: tag + process id
  - *Synonym/alias problem*: two different virtual addresses map to same physical address
    - Two different cache entries holding data for the same physical address!
    - For update: must update all cache entries with same physical address or memory becomes inconsistent
    - Determining this requires significant hardware, essentially an associative lookup on the physical address tags to see if you have multiple hits;
    - Or software enforced alias boundary: same least-significant bits of VA & PA > cache size

# Integrating TLB and Cache

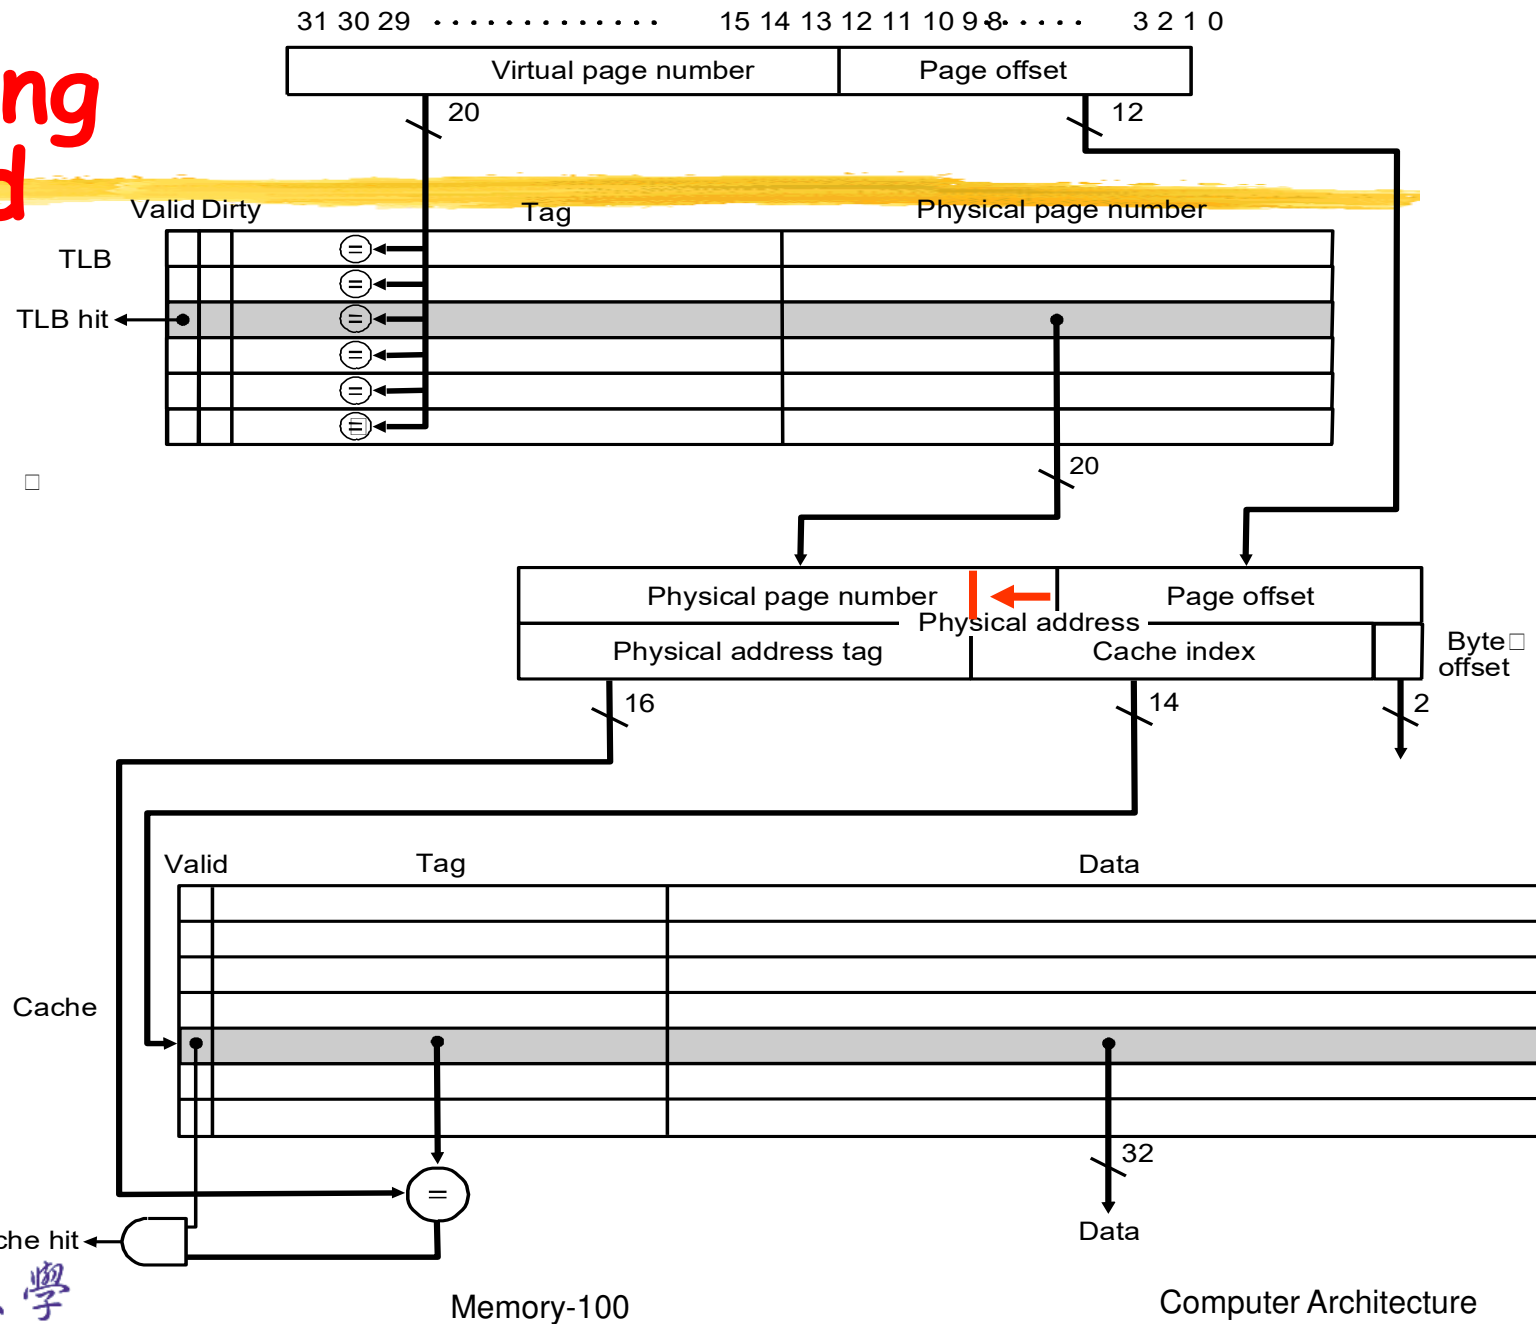
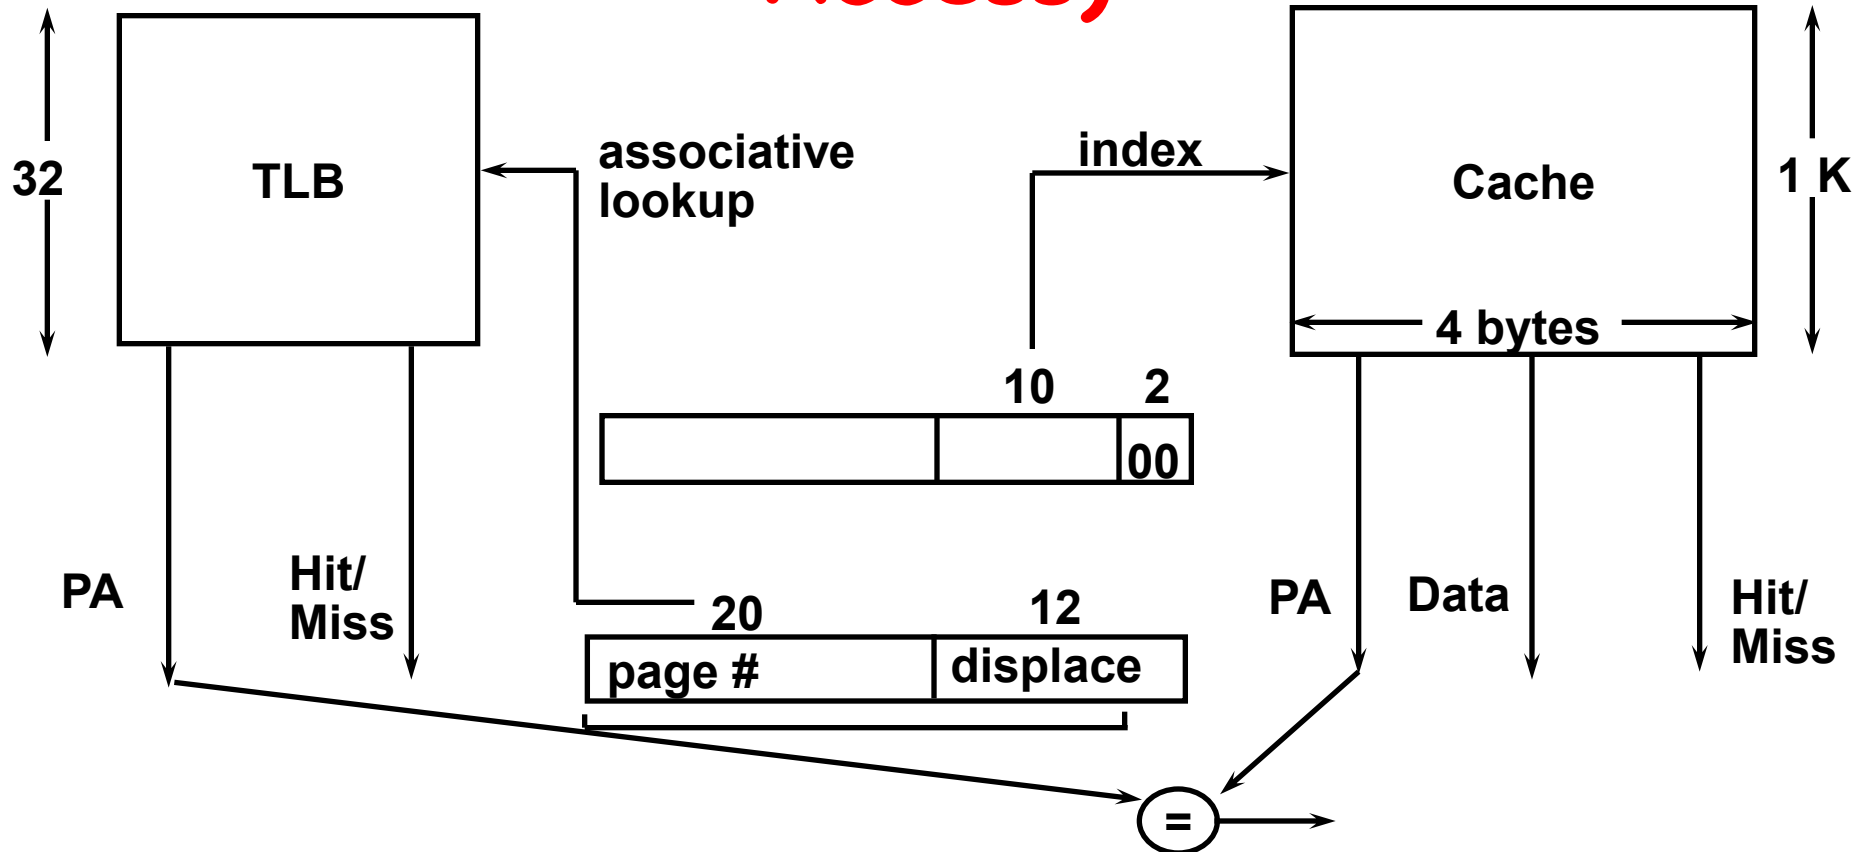


Fig. 5.24

# An Alternative: Virtually Indexed but Physically Tagged (Overlapped Access)



IF cache hit AND (cache tag = PA) then deliver data to CPU  
ELSE IF [cache invalid OR (cache tag  $\neq$  PA)] and TLB hit THEN  
access memory with the PA from the TLB  
ELSE do standard VA translation

# Memory Protection

- ◆ Different tasks can share parts of their virtual address spaces
  - But need to protect against errant access
  - Requires OS assistance
- ◆ Hardware support for OS protection
  - 2 modes: kernel, user
  - Privileged supervisor mode (aka kernel mode)
  - Privileged instructions
  - Page tables and other state information only accessible in supervisor mode
  - System call exception (e.g., syscall in MIPS) : CPU from user to kernel

# Outline

- ◆ Memory hierarchy
- ◆ The basics of caches
- ◆ Measuring and improving cache performance
- ◆ Virtual memory
- ◆ A common framework for memory hierarchy
- ◆ Using a Finite State Machine to Control a Simple Cache
- ◆ Parallelism and Memory Hierarchies: Cache Coherence

# The Memory Hierarchy

## The BIG Picture

- ◆ Common principles apply at all levels of the memory hierarchy
  - Based on notions of caching
- ◆ At each level in the hierarchy
  - Block placement
  - Finding a block
  - Replacement on a miss
  - Write policy

# Block Placement

- ◆ **Determined by associativity**
  - **Direct mapped (1-way associative)**
    - One choice for placement
  - **n-way set associative**
    - n choices within a set
  - **Fully associative**
    - Any location
- ◆ **Higher associativity reduces miss rate**
  - Increases complexity, cost, and access time



# Finding a Block

Associativity	Location method	Tag comparisons
Direct mapped	Index	1
n-way set associative	Set index, then search entries within the set	n
Fully associative	Search all entries	#entries
	Full lookup table	0

- ◆ **Hardware caches**
  - Reduce comparisons to reduce cost
- ◆ **Virtual memory**
  - Full table lookup makes full associativity feasible
  - Benefit in reduced miss rate

# Replacement

- ◆ Choice of entry to replace on a miss
  - Least recently used (LRU)
    - Complex and costly hardware for high associativity
  - Random
    - Close to LRU, easier to implement
- ◆ Virtual memory
  - LRU approximation with hardware support

# Write Policy

- ◆ **Write-through**
  - Update both upper and lower levels
  - Simplifies replacement, but may require write buffer
- ◆ **Write-back**
  - Update upper level only
  - Update lower level when block is replaced
  - Need to keep more state
- ◆ **Virtual memory**
  - Only write-back is feasible, given disk write latency

# Sources of Misses

- ◆ **Compulsory misses (aka cold start misses)**
  - First access to a block
- ◆ **Capacity misses**
  - Due to finite cache size
  - A replaced block is later accessed again
- ◆ **Conflict misses (aka collision misses)**
  - In a non-fully associative cache
  - Due to competition for entries in a set
  - Would not occur in a fully associative cache of the same total size

# Challenge in Memory Hierarchy

- ◆ Every change that potentially improves miss rate can negatively affect overall performance

Design change	Effects on miss rate	Possible effects
size ↑	capacity miss ↓	access time ↑
associativity ↑	conflict miss ↓	access time ↑
block size ↑	spatial locality ↑	miss penalty ↑

- ◆ Trends:


- Synchronous SRAMs (provide a burst of data)
- Redesign DRAM chips to provide higher bandwidth or processing
- Restructure code to increase locality
- Use prefetching (make cache visible to ISA)

# Outline

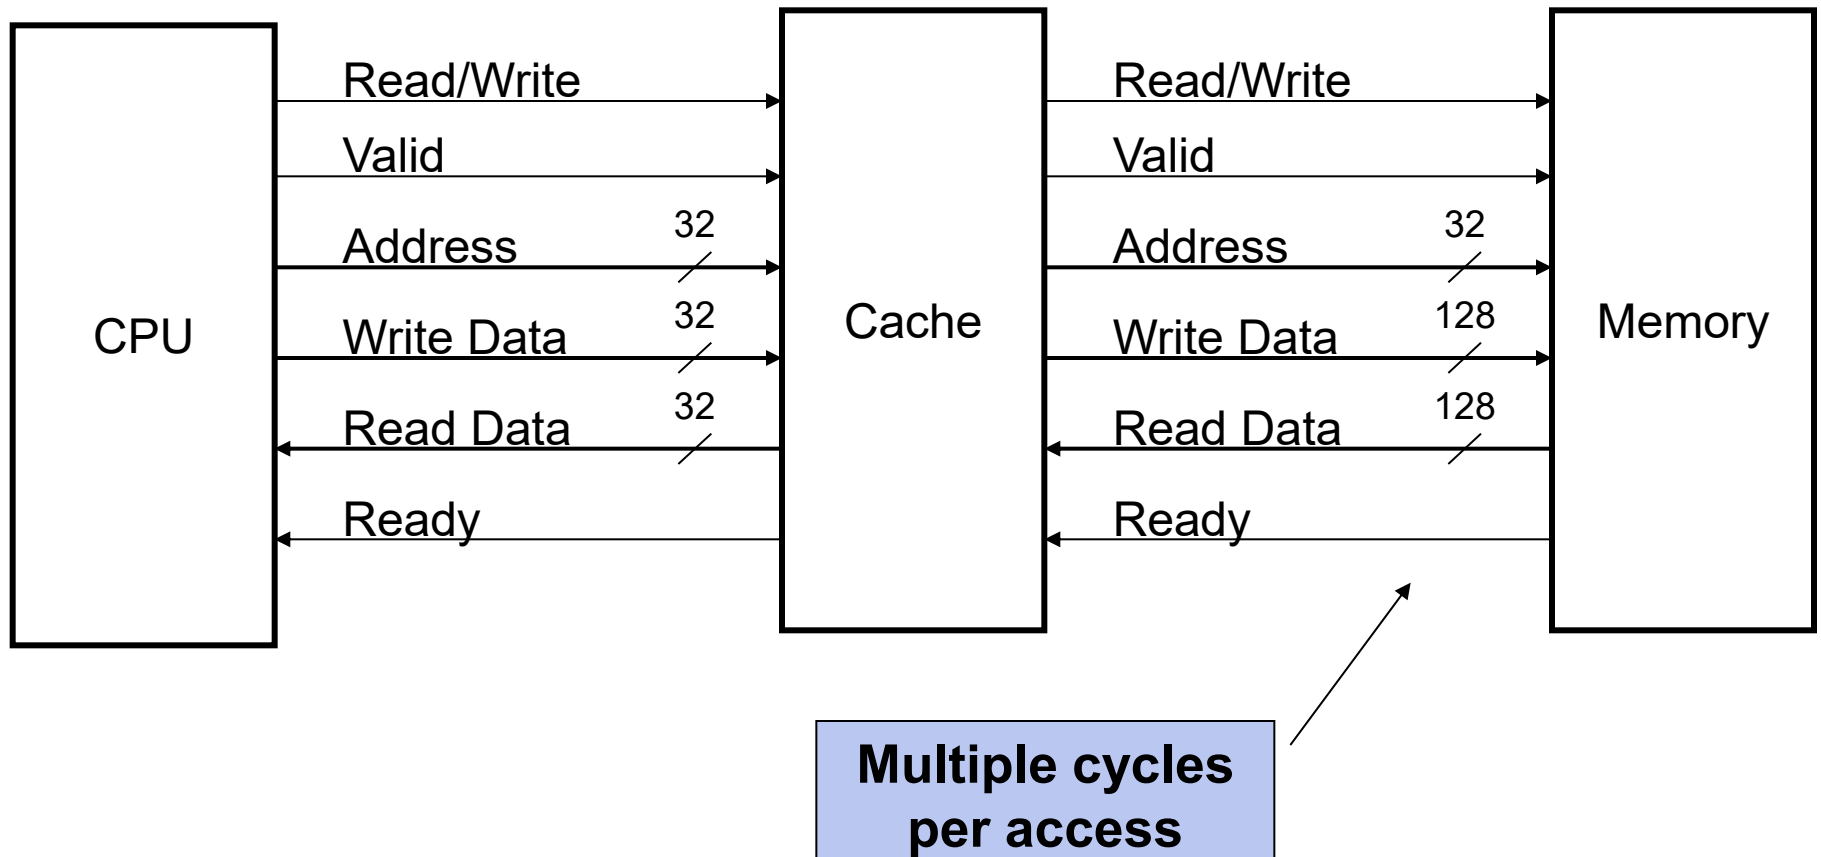
- ◆ Memory hierarchy
- ◆ The basics of caches
- ◆ Measuring and improving cache performance
- ◆ Virtual memory
- ◆ A common framework for memory hierarchy
- ◆ Using a Finite State Machine to Control a Simple Cache
- ◆ Parallelism and Memory Hierarchies: Cache Coherence

# Cache Control

- ◆ Example cache characteristics
  - Direct-mapped, write-back, write allocate
  - Block size: 4 words (16 bytes)
  - Cache size: 16 KB (1024 blocks)
  - 32-bit byte addresses
  - Valid bit and dirty bit per block
  - Blocking cache
    - CPU waits until access is complete

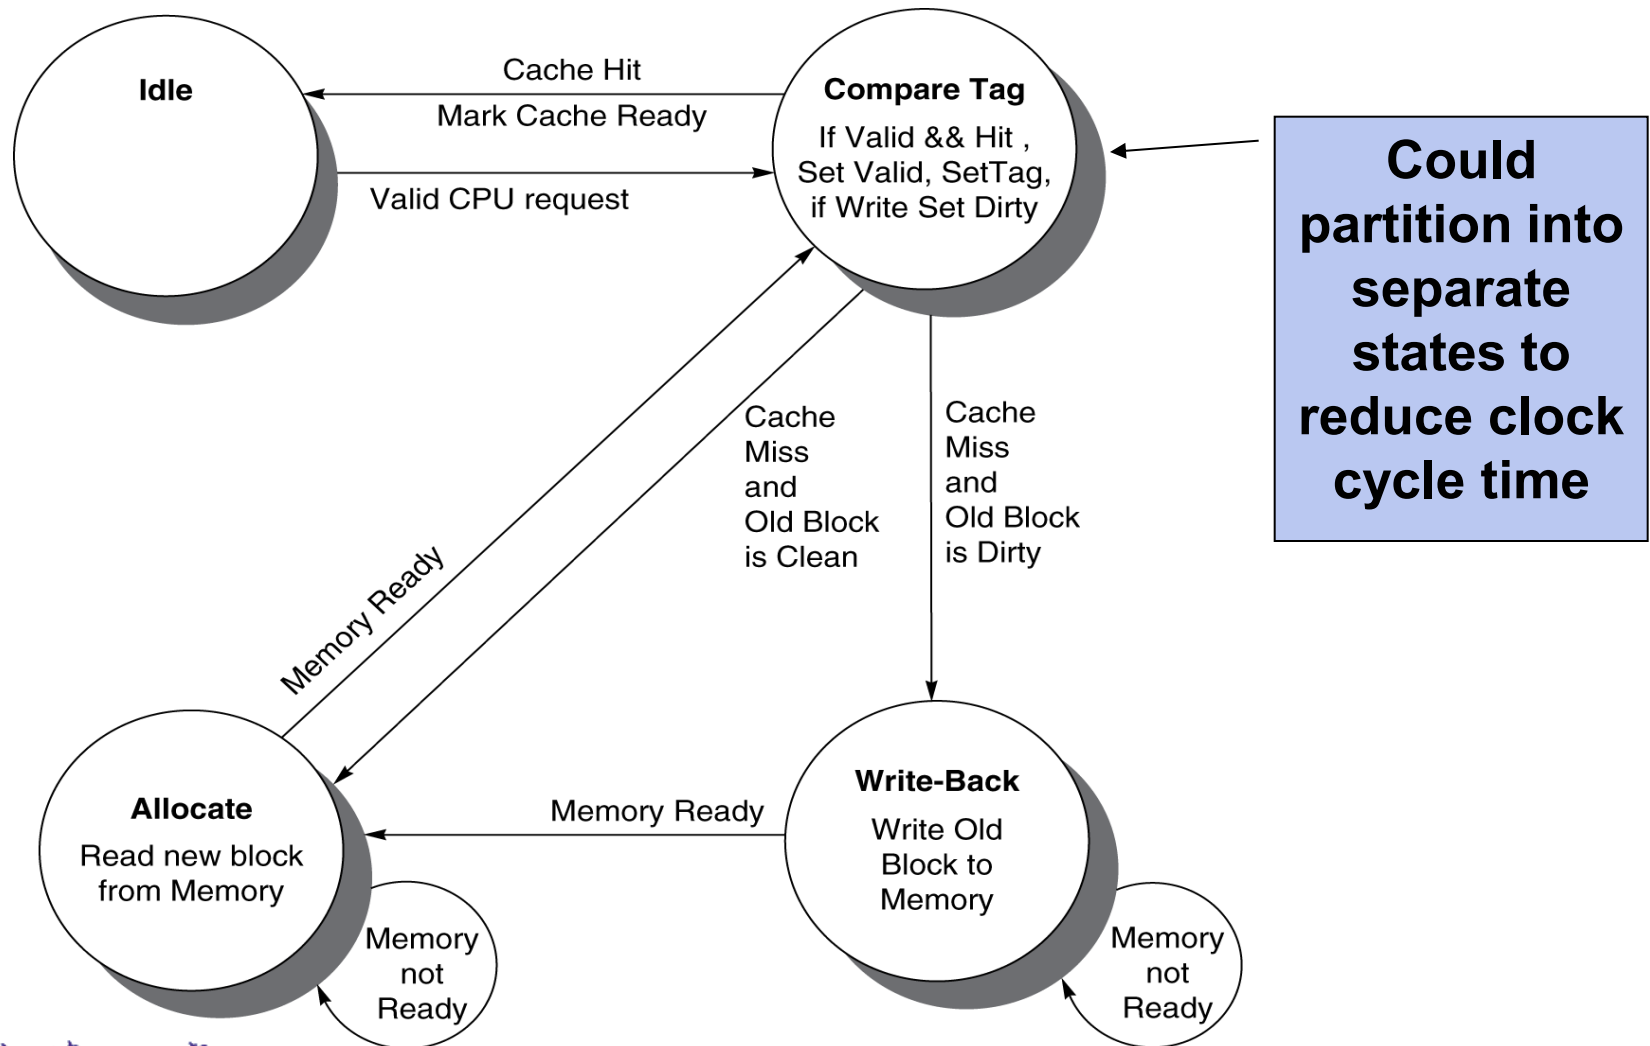


# Interface Signals



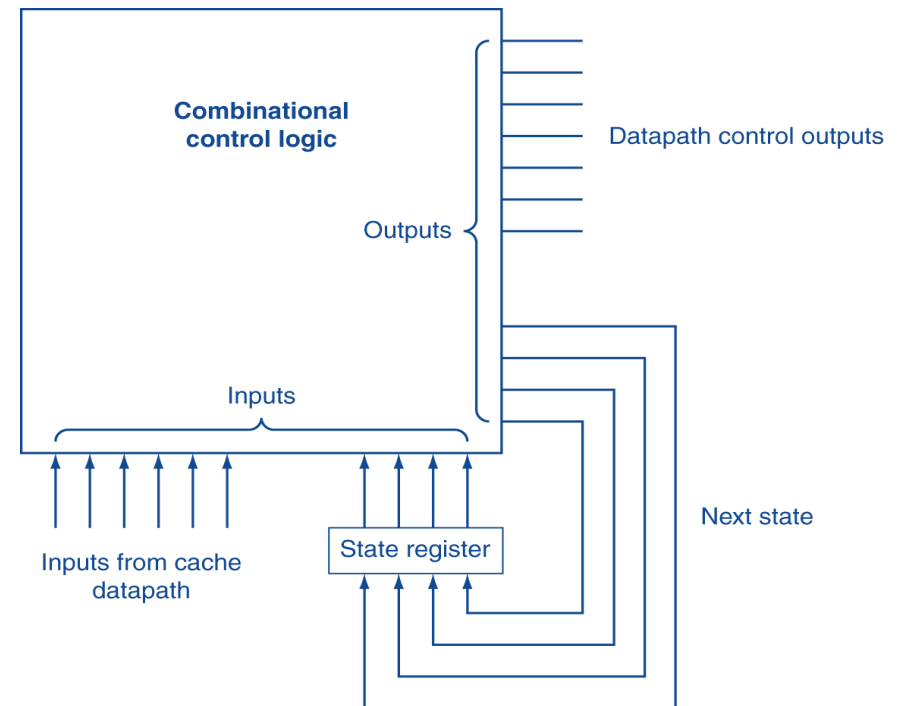


# Cache Controller FSM



# Finite State Machines

- ◆ Use an FSM to sequence control steps
- ◆ Set of states, transition on each clock edge
  - State values are binary encoded
  - Current state stored in a register
  - Next state =  $f_n$  (current state, current inputs)
- ◆ Control output signals =  $f_o$  (current state)



# Outline

- ◆ Memory hierarchy
- ◆ The basics of caches
- ◆ Measuring and improving cache performance
- ◆ Virtual memory
- ◆ A common framework for memory hierarchy
- ◆ Using a Finite State Machine to Control a Simple Cache
- ◆ Parallelism and Memory Hierarchies: Cache Coherence

# Cache Coherence Problem

- ◆ Suppose two CPU cores share a physical address space
  - Write-through caches

Time step	Event	CPU A's cache	CPU B's cache	Memory
0				0
1	CPU A reads X	0		0
2	CPU B reads X	0	0	0
3	CPU A writes 1 to X	1	0	1

- ◆ Two different processors have two different values for the same location

# Coherence Defined

- ◆ Informally: Reads return most recently written value
- ◆ Formally:
  - P writes X; P reads X (no intervening writes)  
⇒ read returns written value
  - $P_1$  writes X;  $P_2$  reads X (sufficiently later)  
⇒ read returns written value
    - c.f. CPU B reading X after step 3 in example
  - $P_1$  writes X,  $P_2$  writes X  
⇒ all processors see writes in the same order
    - End up with the same final value for X

# Cache Coherence Protocols

- ◆ **Operations performed by caches in multiprocessors**
  - **Migration of data to local caches**
    - Reduces bandwidth for shared memory
  - **Replication of read-shared data**
    - Reduces contention for access
- ◆ **Snooping protocols**
  - Each cache monitors bus reads/writes
- ◆ **Directory-based protocols**
  - Caches and memory record sharing status of blocks in a directory

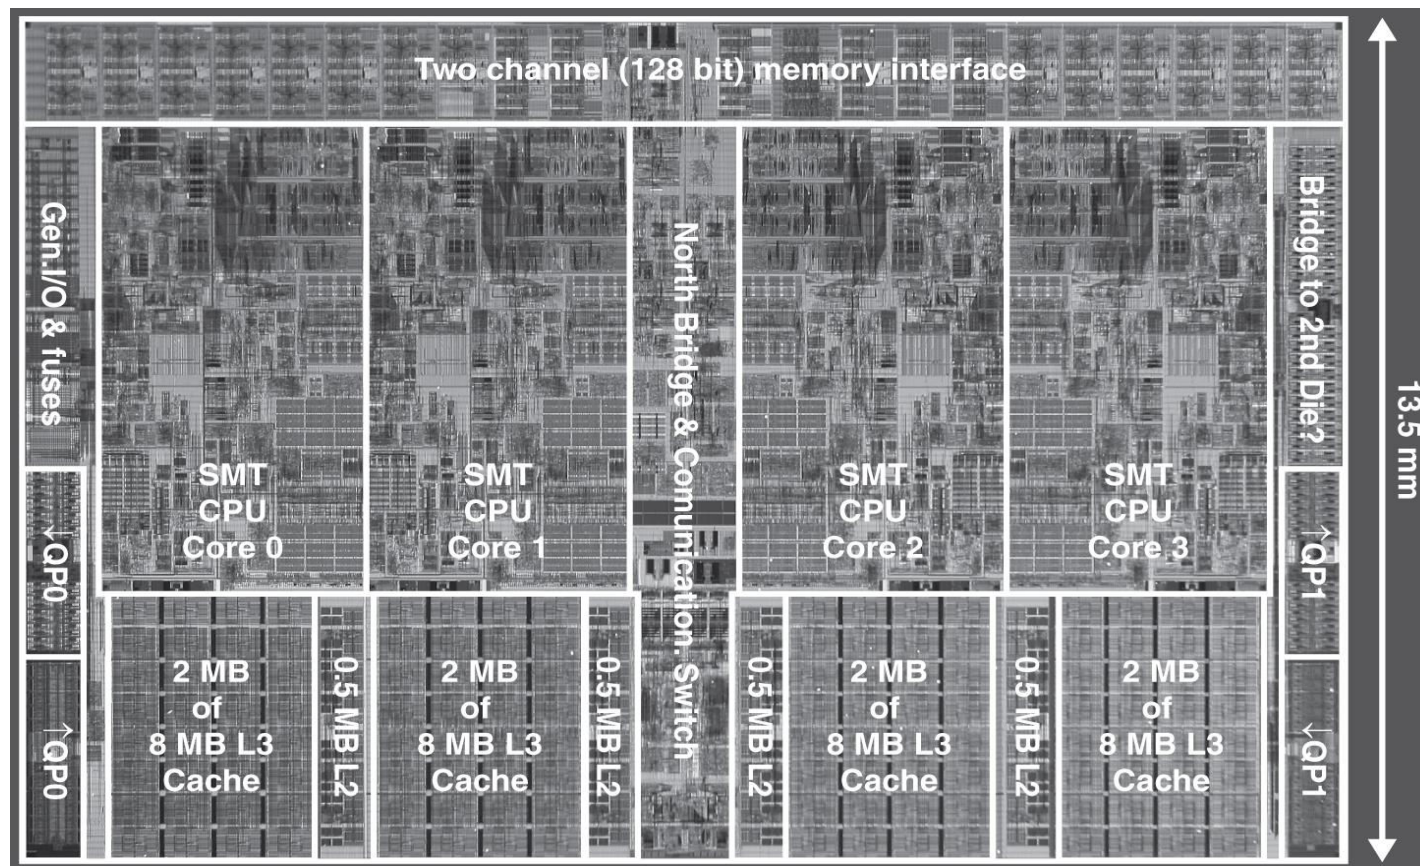
# Invalidating Snooping Protocols

- ◆ Cache gets exclusive access to a block when it is to be written
  - Broadcasts an invalidate message on the bus
  - Subsequent read in another cache misses
    - Owning cache supplies updated value

CPU activity	Bus activity	CPU A's cache	CPU B's cache	Memory
				0
CPU A reads X	Cache miss for X	0		0
CPU B reads X	Cache miss for X	0	0	0
CPU A writes 1 to X	Invalidate for X	1		0
CPU B read X	Cache miss for X	1	1	1

# Multilevel On-Chip Caches

Intel Nehalem 4-core processor



Per core: 32KB L1 I-cache, 32KB L1 D-cache, 512KB L2 cache



# 2-Level TLB Organization

	Intel Nehalem	AMD Opteron X4
Virtual addr	48 bits	48 bits
Physical addr	44 bits	48 bits
Page size	4KB, 2/4MB	4KB, 2/4MB
L1 TLB (per core)	L1 I-TLB: 128 entries for small pages, 7 per thread (2×) for large pages L1 D-TLB: 64 entries for small pages, 32 for large pages Both 4-way, LRU replacement	L1 I-TLB: 48 entries L1 D-TLB: 48 entries Both fully associative, LRU replacement
L2 TLB (per core)	Single L2 TLB: 512 entries 4-way, LRU replacement	L2 I-TLB: 512 entries L2 D-TLB: 512 entries Both 4-way, round-robin LRU
TLB misses	Handled in hardware	Handled in hardware

# 3-Level Cache Organization

	Intel Nehalem	AMD Opteron X4
L1 caches (per core)	L1 I-cache: 32KB, 64-byte blocks, 4-way, approx LRU replacement, hit time n/a L1 D-cache: 32KB, 64-byte blocks, 8-way, approx LRU replacement, write-back/allocate, hit time n/a	L1 I-cache: 32KB, 64-byte blocks, 2-way, LRU replacement, hit time 3 cycles L1 D-cache: 32KB, 64-byte blocks, 2-way, LRU replacement, write-back/allocate, hit time 9 cycles
L2 unified cache (per core)	256KB, 64-byte blocks, 8-way, approx LRU replacement, write-back/allocate, hit time n/a	512KB, 64-byte blocks, 16-way, approx LRU replacement, write-back/allocate, hit time n/a
L3 unified cache (shared)	8MB, 64-byte blocks, 16-way, replacement n/a, write-back/allocate, hit time n/a	2MB, 64-byte blocks, 32-way, replace block shared by fewest cores, write-back/allocate, hit time 32 cycles

n/a: data not available

# Pitfalls

- ◆ **Byte vs. word addressing**
  - **Example: 32-byte direct-mapped cache, 4-byte blocks**
    - Byte 36 maps to block 1
    - Word 36 maps to block 4
- ◆ **Ignoring memory system effects when writing or generating code**
  - **Example: iterating over rows vs. columns of arrays**
  - **Large strides result in poor locality**

# Pitfalls

- ◆ In multiprocessor with shared L2 or L3 cache
  - Less associativity than cores results in conflict misses
  - More cores  $\Rightarrow$  need to increase associativity
- ◆ Using AMAT (Average Memory Access Time) to evaluate performance of out-of-order processors
  - Ignores effect of non-blocked accesses
  - Instead, evaluate performance by simulation

# Concluding Remarks

- ◆ **Fast memories are small, large memories are slow**
  - We really want fast, large memories ☹
  - Caching gives this illusion ☺
- ◆ **Principle of locality**
  - Programs use a small part of their memory space frequently
- ◆ **Memory hierarchy**
  - L1 cache ↔ L2 cache ↔ ... ↔ DRAM memory  
↔ disk
- ◆ **Memory system design is critical for multiprocessors**